G.H.O.S.T. Capabilities

COMBINED CAPABILITY PROFILES

7 capabilities that sit on top of every G.H.O.S.T. agent.

01	AI Council	Convene more than one expert
02	M.A.C.S.	Wargame the plan before the world does
03	Agentic Loop	One prompt, many rounds
04	Smart Router	One box. The router decides.
05	Memory	Every session starts where the last one ended
06	Hunter	It doesn't wait for an alert
07	G.H.O.S.T. Mode	One request. The right agents.

CAPABILITY 01

AI Council

STANDING TEAM
MULTI-MODEL SEATS
ONE ACCOUNTABLE VERDICT

Convene more than one expert | Multi-agent deliberation

Real problems cross domains - and single-agent AI forces a staff to break the problem into pieces, ask separate questions, and stitch the answers together by hand. The stitching is where the mistakes live. The AI Council convenes several specialist agents on the same problem: state your tasking, and the platform recommends a chair agent to organize the effort and specialist seats that each bring their own persona, workflows, and document corpus to the same thread. The council deliberates in front of the operator and returns one verdict - then stands by for the next tasking.

HOW IT WORKS

1 TASKING
State the main goal

2 RECOMMEND
GHOST proposes chair + seats

3 THE COUNCIL
Chair and seats. Each seat: its own persona, workflows, corpus - and its own model, chosen per role

4 DELIBERATE
Challenge and reconcile - dissent stays visible

5 VERDICT
One synthesis, all reasoning kept

A STANDING TEAM - TASK IT AGAIN

ONE SEAT, ONE MODEL

Council seats are not clones of a single model. Each seat can run on a different underlying model - frontier or local - so the deliberation carries genuinely different perspectives, not one model agreeing with itself five times.

A STANDING TEAM

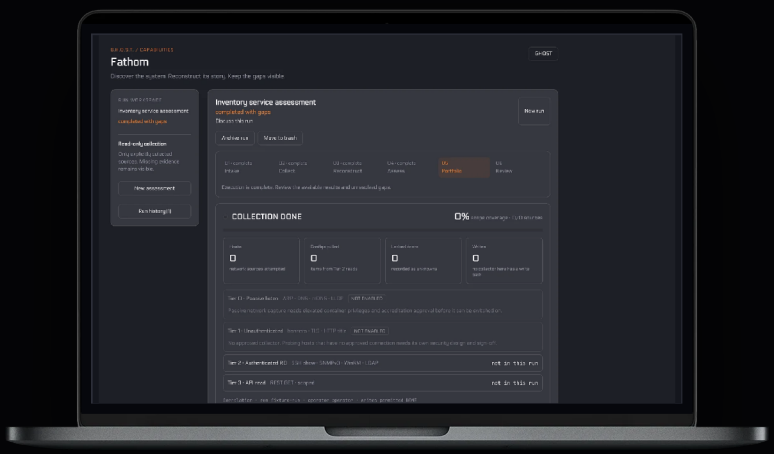
Convene once, then keep tasking it. The council is a long-running capability that carries context forward as the mission evolves - the team you assembled yesterday takes today's follow-on tasking without re- briefing.

ADVISORY-ONLY BY DESIGN

The operator keeps the decision. The council explains, assesses, debates, and recommends; the verdict arrives with each seat's reasoning traceable, and dissent recorded rather than hidden.

DEPLOYS YOUR WAY

Air-gapped on-prem, cloud-connected, or hybrid, the same platform - and the same council - deployed to match your security posture and mission need, built from the base agents you already field.



The AI Council convenes the G.H.O.S.T. agent roster - one chair, many perspectives, one accountable verdict.

CAPABILITY 02

M.A.C.S.

AN OPPONENT WITH A DOCTRINE
ROUNDS YOU CONTROL
A WRITTEN RECORD

Wargame the plan before the world does | Multi-agent competitive simulation

Every plan survives until it meets the other side. M.A.C.S. lets you meet the other side early - on your own terms. Stand up a friendly agent that knows your unit: your doctrine, your OPORD, your constraints. Stand up the adversary and arm it with everything you hold on them - a foreign tactics manual, intelligence, observed behavior. Then let them fight. Every hole the opposition finds is one you close before the real world finds it. Same engine in the boardroom: your company on one side, your three biggest competitors on the other - and the product launch or price change wargamed before you make the move.

HOW IT WORKS

1 YOUR SIDE

The friendly agent, built from your material. Doctrine, OPORD, unit details - or strategy, pricing, the plan you intend to execute

2 THE OPPOSITION

The adversary or your three biggest competitors. Foreign tactics manual, intel, observed behavior, market posture - everything you hold

3 THE ENGAGEMENT

Probe, counter, exploit - round after round, in front of the operator

Set the rounds: quick test or deep campaign

4 THE RECORD

Exposed assumptions, holes found, counters that worked - in writing

5 REVISE AND RE-RUN

Fix the plan, keep the sides standing, and fight it again until it holds

AN OPPONENT WITH A DOCTRINE

Not a devil's advocate - an adversary. The opposition argues from the tactics manuals, intelligence, and competitor context you load, so its moves are theirs, not a generic critique of your plan.

ROUNDS YOU CONTROL

One control sets how many rounds the sides trade. Dial it down for a fast pressure test before a decision brief; dial it up for a sustained campaign against your hardest problem.

ONE ENGINE, TWO ARENAS

Pit your unit against a foreign adversary - or your company against the competitors most likely to answer your next move. New product, price change, market entry: wargame it first.

ADVISORY-ONLY BY DESIGN

The commander - or the executive - keeps the decision. A simulation is an assessment, not an order; every finding in the record traces to the round and the material that produced it.

> M.A.C.S. puts a red team on retainer - wargame the plan before the world does.

CAPABILITY 03

Agentic Loop

MULTI-ROUND RETRIEVAL
SOURCE-CITED OUTPUT
GAPS STATED, NOT FILLED

One prompt, many rounds | Retrieve, verify, assemble

Ask most AI for a work product and it answers from one pass at your files - then writes confidently over whatever it did not find. The gap is where the mistakes live. The Agentic Loop is how a G.H.O.S.T. agent works a tasking the way a trained analyst would: it plans its own lines of inquiry, searches your document corpus and your trusted sources of truth round after round - regulation, doctrine, vendor documentation, standard operating procedures, whatever governs your work - opens and reads the passages that matter, checks what it still lacks, and only then assembles the product, delivered with its sources cited and its assumptions declared.

HOW IT WORKS

1 TASKING
One plain-language request

2 STAGE
Workflow, files, model and context selected

3 THE AGENTIC LOOP
Plan, search, read, verify, re-query. Round n - it keeps going until the evidence holds. Your corpus and your trusted sources of truth, searched and re-searched - every hit scored and traceable

4 ASSEMBLE
Built section by section from what it read

5 DELIVER
Product, sources, and stated gaps

REFINE IT IN PLACE - THE NEXT PROMPT PICKS UP THE WHOLE THREAD

IT SHOWS ITS WORK

Every round is on screen as it happens - the queries it chose, the files and reference pages it hit, the relevance score of each, and the passages it opened. Nothing about how the answer was reached is hidden from the operator.

GROUNDED, NOT GUESSED

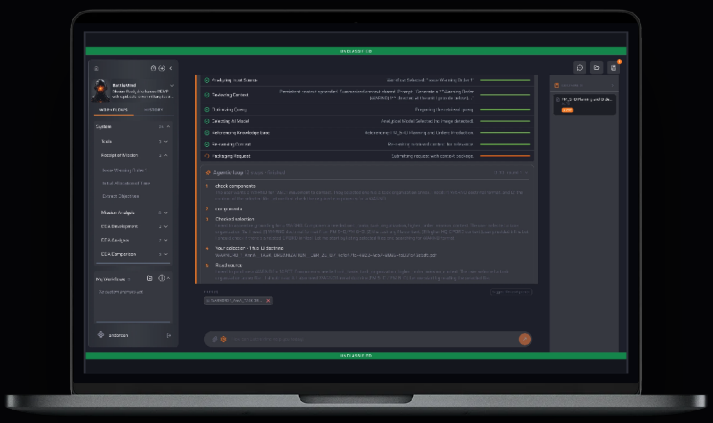
Products arrive with a grounded-on list naming the source documents and authoritative references behind them, so a reviewer can walk any line of the output back to the page it came from.

GAPS STATED, NOT FILLED

What the corpus does not contain is declared as an assumption or gap for your people to resolve - intent, timings, decisions still owed - rather than quietly invented to complete the page.

DEPLOYS YOUR WAY

Air-gapped on-prem, cloud-connected, or hybrid, the same platform - and the same loop - deployed to match your security posture and mission need, built from the agents you already field.



> The Agentic Loop is how a G.H.O.S.T. agent earns its answer - retrieve, verify, assemble, and cite your trusted sources of truth.

CAPABILITY 04

Smart Router

ALWAYS-ON CLASSIFIER
POLICY-SET THRESHOLD
REVERSIBLE REDACTION

One box. The router decides. | Complexity-aware model routing

Local-only AI gives up the hardest reasoning; frontier-only AI gives up control of the data - and of the bill. The work is not uniform, so the answer should not be either. The Smart Router scores every request for reasoning depth before anything moves, applies your policy and spend window, and sends the task to the AI model that actually fits - a small model on your own server for high-volume pattern work, a large model on your own server for full-document analysis, and an off-site frontier model only when the reasoning requires it. When frontier is selected, sensitive values are tokenized at the boundary and rehydrated on return: the reasoning leaves, the data does not.

HOW IT WORKS

1 REQUEST

Operator states the task

2 CLASSIFY

Custom-trained classifier scores reasoning depth - always on

Policy, threshold, spend window: sovereignty-weighted, balanced, quality-first, or frontier off - you control where the tokens go

3 ROUTING PROXY

Three AI models, one interface - the operator never picks

4 ONE MODEL SELECTED

Small AI model on your server: high volume, strict schema, pattern work. Large AI model on your server: full-document context, multi-step comparison. Frontier AI model off-site: long reasoning chains, maximum context

5 SANITIZE, REHYDRATE

PII becomes tokens before the call; the mapping stays in an on-prem vault and is never transmitted

ONE ANSWER RETURNS COMPLETE - NOTHING SENSITIVE LEFT

THE RIGHT MODEL, AUTOMATICALLY

Nobody picks a model. The classifier runs on every request and the proxy makes the call, so routine work is absorbed locally and frontier reasoning is reserved for the problems that actually need it.

REVERSIBLE REDACTION

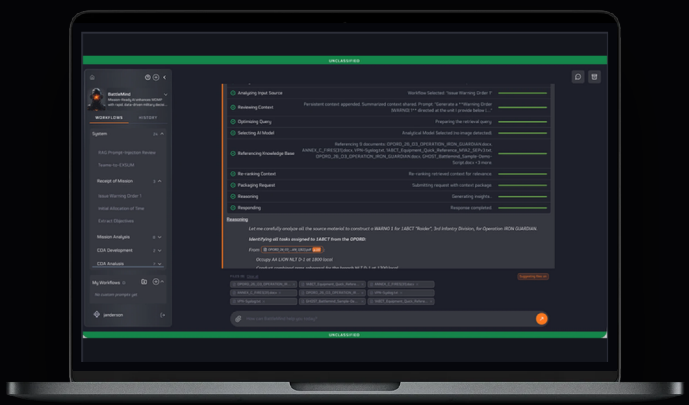
Sanitization is not deletion. Names, dates of birth, record numbers, and party names become tokens at the boundary and are restored from a local vault on return - a full-fidelity answer with no exposure.

COST STAYS UNDER YOUR CONTROL

Token spend is where most organizations quietly bleed money on AI. Frontier use is bounded by an envelope you set; when consumption runs past it, the router falls back to local tiers until the window resets - the work continues, the invoice does not climb.

ADVISORY-ONLY BY DESIGN

The router decides where reasoning happens, never what the answer should be. Routing decisions, policy trips, and sanitization steps are visible to the operator, and the human keeps authority over the work.



The Smart Router is the routing layer beneath the G.H.O.S.T. agent roster - every agent inherits it.

CAPABILITY 05

Memory

LEARNED, NOT UPLOADED
NOMINATED, THEN PROMOTED
PINNED, NEVER FORGOTTEN

Every session starts where the last one ended | Personal + team recall

Your files answer what you know; the Authority Library answers what the rules are. Neither answers the question every returning team asks: what have we already learned? MEMORY does. As GHOST works, it captures corrections, decisions and their rationale, conventions, and constraints - and users can pin a fact themselves with "remember this." Personal memory files itself and stays yours to edit. Nothing enters Team Memory silently: a curator promotes every entry, so the shared store stays small, current, and trusted. The pinned core rides in every session at a fixed, capped cost; what retires stays searchable.

HOW IT WORKS

1 CAPTURED

The session yields candidates. GHOST proposes what's worth keeping - or the user says "remember this"

2 CURATED

Personal memory files itself and stays user-editable; team entries wait in the queue until the curator promotes them

3 THE PINNED CORE

Rides in every session. Hard-capped, so the cost is fixed and known - retired entries stay searchable

4 IT STAYS CURRENT

Provenance and a last-confirmed date on every entry - stale ones retire to the archive

The memory budget: lean core, past this it's RAG

AUTHORITY OUTRANKS MEMORY: DOCTRINE, POLICY, AND SYSTEM FILES ALWAYS WIN; A CONFLICTING MEMORY IS FLAGGED, NOT OBEYED

LEARNED, NOT UPLOADED

RAG holds what you gave it; MEMORY holds what working together produced - the correction that ended a bad habit, the decision and why it was made, the convention nobody wrote down. Corrections are captured first because re-teaching is the most expensive mistake.

NOTHING ENTERS TEAM MEMORY SILENTLY

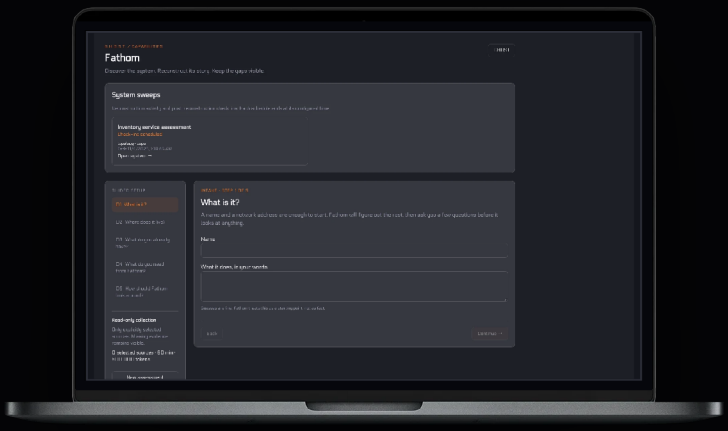
GHOST and users can only nominate a shared entry; a named curator promotes it from the queue. The whole store is visible to every member, and every entry carries its provenance - who, when, and the session it came from.

YOURS TO EDIT, YOURS TO SILENCE

Personal memory is fully user-editable - view, change, or delete any entry, any time - and any session can run with memory off. Team memory is off by default and enabled per team, deliberately.

MEMORY NEVER OUTRANKS THE RULES

Task instructions, the Authority Library, and System Files all take precedence. When a memory drifts out of step with doctrine or policy, GHOST raises the conflict for a human to resolve - memory doubles as a drift detector.



MEMORY is the G.H.O.S.T. platform's institutional recall - it learns as your team works, so no session starts from zero.

CAPABILITY 06

Hunter

SWEEP EVERYTHING, CHEAPLY
TRIAGE ONLY WHAT SURVIVES
CHASE WHAT IT CONFIRMS

It doesn't wait for an alert | Autonomous 24/7 hunt

Detection waits for a rule to fire, which means it only ever finds what somebody already knew to look for. HUNTER goes looking instead. Point it at a log repository, a SIEM export, a database, or a published API; it maps the terrain deterministically, takes its mission in plain language, then sweeps every record with detectors that cost nothing to run - and shows a model only the handful of candidates that survive. Full coverage is what makes it cheap, and cheap is what lets it run around the clock. Confirmations pivot into their own sub-hunts, and the loop keeps going inside the GPU duty cycle you set.

HOW IT WORKS

1 THE TERRAIN

Logs, SIEM, database, API. Enumerated and baselined first: inventory, entities, what normal looks like - 0 tokens

2 THE HUNT

Detectors sweep the whole corpus for free; the model triages only survivors - and confirmations pivot

The mission, in plain words: you state the intent and arm the behaviors; the plan comes back for approval, and that intent anchors every hunt that follows it

3 THE FINDINGS

Severity, confidence, evidence, next step. Incident case, timeline, coverage and assurance, rolling shift summary

Set the GPU duty cycle: share the card, or run flat out

4 IT DOESN'T STOP

Each finding becomes a hypothesis and runs down in its own sub-hunt; what that turns up sets the next - always against your intent

COVERAGE YOU CAN PROVE

Deterministic detectors read every record, so the assurance report can state exactly which detector examined what - including the hypotheses eliminated and the work deferred to the next period.

CANDIDATES ARE NOT FINDINGS

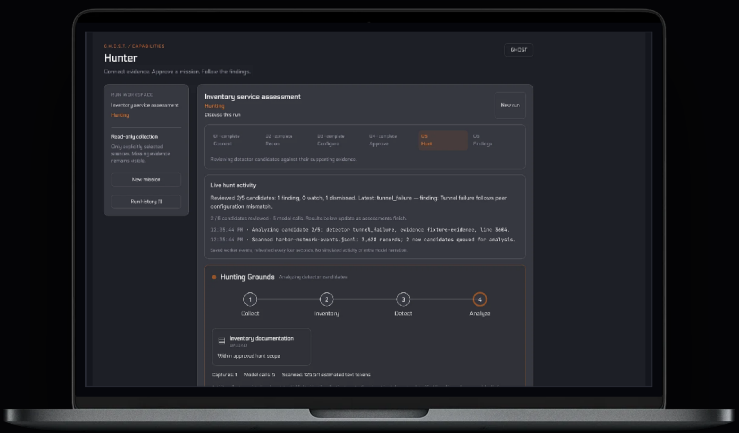
Triage weighs every candidate against the terrain baseline and reports the rejections alongside the escalations, with the reasoning that decided each.

EVERY FIND OPENS THE NEXT HUNT

A confirmation becomes a hypothesis and runs down in its own sub-hunt - foothold to traversal to staging - and whatever that turns up seeds the hunt after it, always against the intent you set.

ADVISORY-ONLY BY DESIGN

HUNTER detects, correlates, and recommends. Analysts and incident commanders keep the decision; every finding carries its detector, its triage reasoning, and its evidence on the face of the report.



HUNTER maps the terrain, works the mission, and never stops looking - the hunt runs while the queue sleeps.

CAPABILITY 07

G.H.O.S.T. Mode

DOMAIN DETECTED
AGENTS AUTO-SELECTED
ONE SYNTHESIZED PRODUCT

One request. The right agents. | Full-agentic orchestration

Real taskings cross specialties - a warning order carries foreign-language threat reporting and another service's fires; a protest response turns on acquisition procedure, regulatory reference, and cloud architecture at once. Nobody should have to know which agent to open first. In G.H.O.S.T. Mode the user states the task in plain language and attaches what they have. G.H.O.S.T. detects the mission domain, assembles the specialist team, binds each agent to the Automated Workflow that fits its part of the job, and runs them in sequence - returning one deliverable with grounding citations and an explicit assumptions-and-gaps list.

HOW IT WORKS

1 REQUEST
Plain language plus your documents

2 DOMAIN LOCK
Mission lexicons matched - ambiguity is reported, not guessed

3 ORCHESTRATION
Lead + specialists + bound workflows. Each agent is selected with a stated rationale and bound to the Automated Workflow for its part of the job

4 EXECUTION
Sequenced, logged - pauses to ask rather than assume

5 DELIVERABLE
One synthesis, cited - staged for human review

IT ASKS INSTEAD OF ASSUMING

When a step needs an input the request never supplied, execution pauses and asks one targeted question. Completed context is held in place - answering resumes the run rather than restarting it, and the answer is carried through the finished product.

A RUN BECOMES A WORKFLOW

Any successful run can be saved as a named Automated Workflow with its steps, triggers, and input parameters captured. Questions asked mid-run become parameters asked up front, and the workflow can be shared to the unit or office catalog.

ADVISORY-ONLY BY DESIGN

G.H.O.S.T. Mode assembles, analyzes, drafts, and recommends. Commanders, Contracting Officers, and counsel remain the sole authorities; every product is staged for human review with grounding citations, assumptions, gaps, and flagged conflicts on its face.

DEPLOYS YOUR WAY

Air-gapped on-prem, cloud-connected, or hybrid, the same platform deployed to match your security posture and mission need - orchestrating the base agents and workflows you already field.



> G.H.O.S.T. Mode is the orchestration layer over your G.H.O.S.T. agent roster - one request in, one reviewed product out.