G.H.O.S.T.

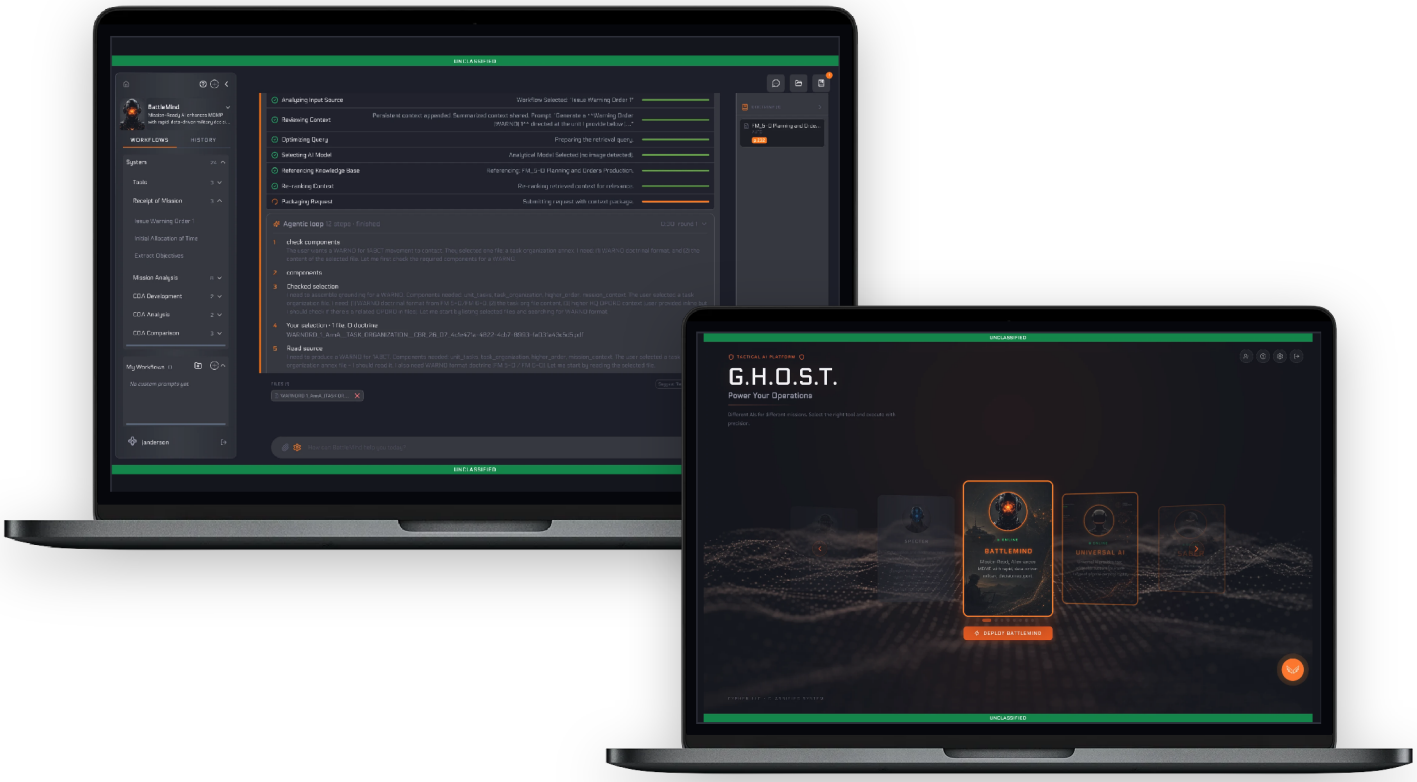
PLATFORM HANDBOOK

The Handbook

GUIDED HEURISTIC ON-PREM SUPPORT & TROUBLESHOOTING

The platform, its capabilities, its 34 agents, and how it is used in the field and in business. One place for everything G.H.O.S.T. does.

MISSION AI | INSIDE YOUR OWN NETWORK | ADVISORY ONLY



CONTENTS

The G.H.O.S.T. handbook

About this handbook 04

01 THE PLATFORM 05

Overview 06

Platform architecture 07

How it works 08

How G.H.O.S.T. is different 09

Featured agents 10

Using G.H.O.S.T. on screen 11

Inside the platform 12

Government alignment and compliance 13

Mission use cases 14

Security and governance 15

Unique advantages 16

Plug-and-play, ROI and licensing 17

02 CAPABILITIES 18

03 THE AGENT ROSTER 26

The roster by mission area 27

04 IN THE FIELD AND IN BUSINESS 62

Why G.H.O.S.T. 63

Specialized agents 64

Real-world business impact 65

Contracts and acquisition intelligence 66

Agents across the enterprise 67

05 DEPLOY AND GET STARTED 68

Deployment models 69

Get started 70

CONTENTS

Capabilities and agents

CAPABILITIES

AI Council 19
M.A.C.S. 20
Agentic Loop 21
Smart Router 22
Memory 23
Hunter 24
G.H.O.S.T. Mode 25

THE 34 AGENT CARDS

AEROCENT Air Component Planning Intelligence 28
AMBASSADOR Translation & Linguistic Intelligence 29
ARIA Post-Award Acquisition Intelligence 30
BASTION Security Assessment & Authorization 31
BATTLEMIND Mission-Ready Military Decision Support 32
CADENCE Requirements, Deliverables & Program Rhythm 33
CASCADE Semantic Change-Impact Analysis 34
DataMind Datacenter Facility & Critical Infrastructure Analysis 35
FATHOM Legacy System Reconstruction from Evidence 36
GAUNTLET AI Evaluation & Assurance 37
GridLock Microgrid Engineering, Operations & Test-and-Evaluation 38
HAVEN Employee Support & Workplace Guidance 39
InWard Insider Threat & Counterintelligence Analytics 40
J-ARC Joint Analysis & Recommendation Center 41
JURIS Legal Evaluation & Regulatory Analysis 42
LATTICE Database, ETL & Data-Quality Analysis 43
LORE Tacit Knowledge Capture & Expert Transfer 44
LYNX Linux & FRR Cloud-Router Engineering 45
NODE Network Operations & Infrastructure Analysis 46
PERISCOPE OSINT Collection Planning & Source Vetting 47
PRYME Pre-Award Acquisition Intelligence 48
PULSE Telemetry Assurance & Blind-Spot Analysis 49
QUARRY Mission Data Exploitation & Value Scoring 50
REDACT FOIA Processing & Declassification Review 51
SABER Marine Corps Mission Planning Intelligence 52
SIFT Digital Forensics & Incident Response Analysis 53
SUB-ZERO Naval Mission Planning & Information Warfare 54
SYNQ Staff Work Integration & Decision Support 55
SysTower Server Operations & Infrastructure Analysis 56
TALON Cyber Threat Intelligence 57
TRIAD Multi-Cloud Infrastructure & Engineering 58
UNIVERSAL AI General-Purpose Assistant & Concierge 59
VANTAGE Analytic Tradecraft & Finished Intelligence 60
VELLUM Corpus Engineering & Knowledge Governance 61

ABOUT THIS HANDBOOK

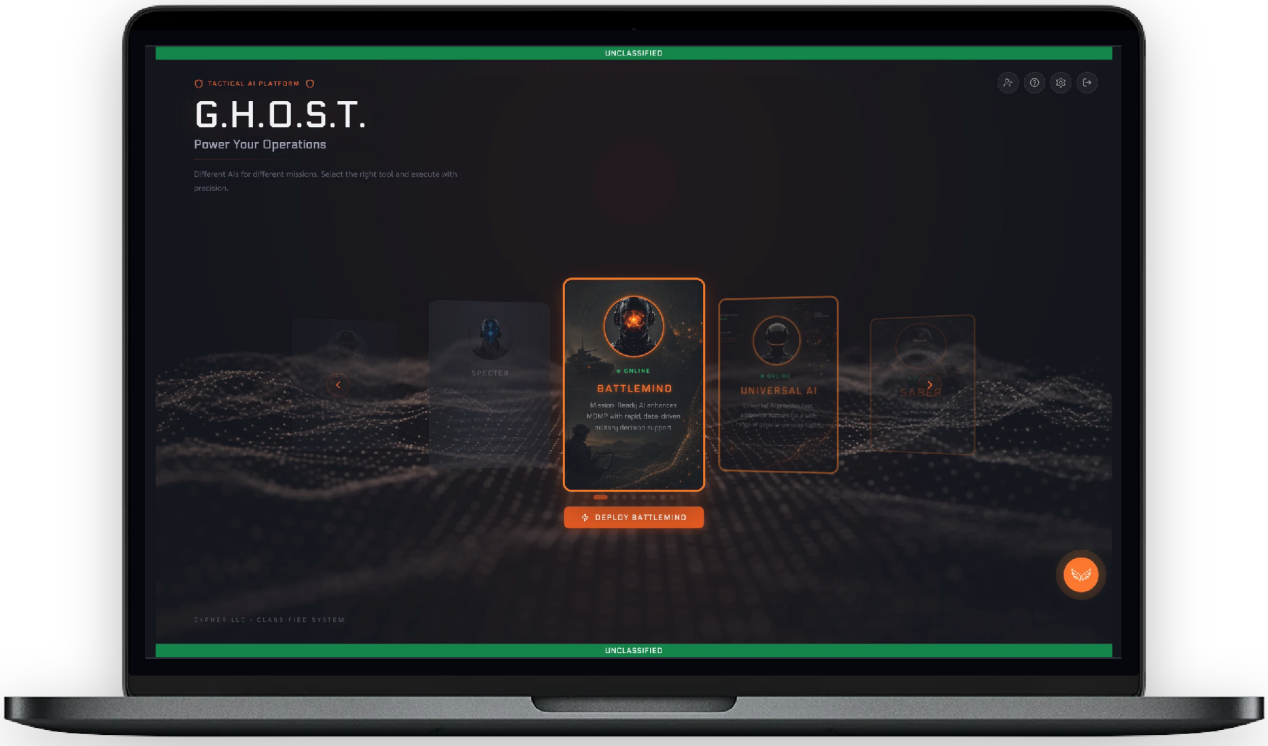
Everything G.H.O.S.T., in one document.

G.H.O.S.T. (Guided Heuristic On-Prem Support & Troubleshooting) is a commercial, modular agentic AI platform. It hosts, orchestrates and extends a library of mission-specific AI agents, and it runs on your own hardware, in a private cloud, or fully disconnected.

This handbook brings together the platform capability brochure, the commercial brochure, the extended capability profiles and all 34 agent cards, so a reader can find what the platform is, what sits on top of it, which agent does what, and how it is deployed and bought.

HOW IT IS ORGANIZED

1 THE PLATFORM	What G.H.O.S.T. is, how it is built, how it works on screen, and how it is secured and governed
2 CAPABILITIES	AI Council, M.A.C.S., Agentic Loop, Smart Router, Memory, Hunter and G.H.O.S.T. Mode
3 THE AGENT ROSTER	The 34 agents, grouped by mission area, then one card for each
4 IN THE FIELD AND IN BUSINESS	Real-world results, acquisition intelligence and mission planning
5 DEPLOY AND GET STARTED	Deployment models, security and how to request a demo





The platform

What G.H.O.S.T. is, how it is built, how it works, and how it is secured and governed.

Overview and why G.H.O.S.T. is different

Architecture and how it works

Government alignment and compliance

Security and governance

OVERVIEW

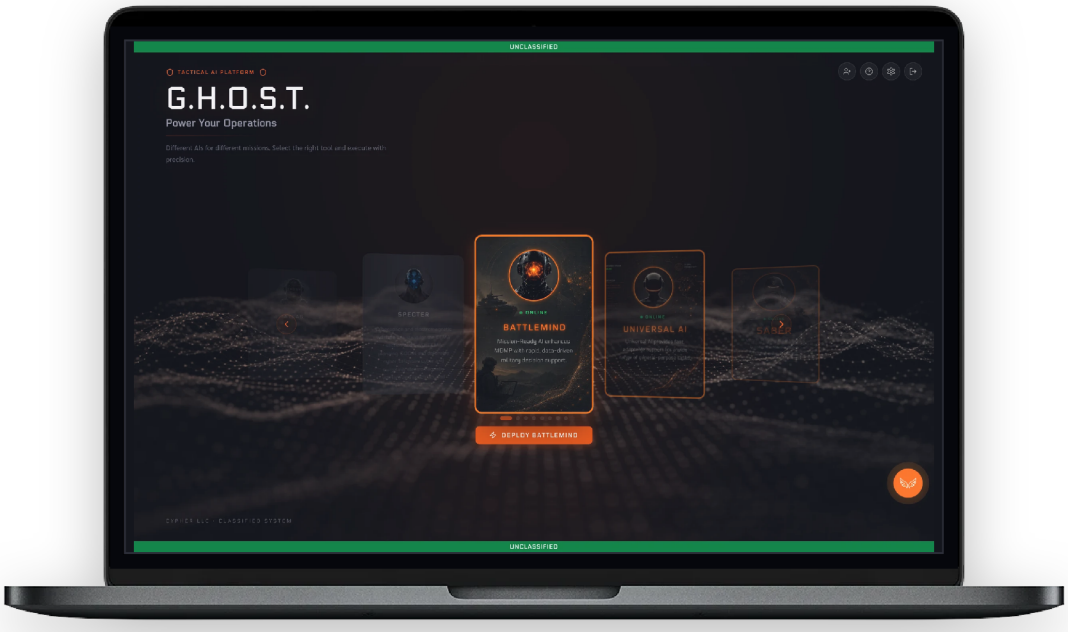
G.H.O.S.T. (Guided Heuristic On-Prem Support & Troubleshooting)

G.H.O.S.T.™ (Guided Heuristic On-Prem Support & Troubleshooting) is a commercial, modular agentic AI platform - hardened and operationally proven for the demands of national security and defense - delivering agentic, multi-step automated workflows across both cloud and on-device deployment models, giving operators the flexibility to leverage enterprise cloud infrastructure when available while retaining full mission capability with zero cloud dependency, zero data exfiltration risk, and no connectivity requirement.

G.H.O.S.T. hosts, orchestrates, and extends a growing library of mission-specific AI agents purpose-built for virtually any mission objective - including **BATTLEMIND**, a doctrine-trained agent for Army mission planning and warfighting functions that automates the Military Decision-Making Process and accelerates course of action development; **NODE** for network diagnostics and automation; **AMBASSADOR** for real-time translation and cultural advisory support; and **ARIA** for advanced intelligence research and analytical fusion - each capable of executing complex, multi-step tasks with minimal human intervention regardless of network environment.

OPERATIONALLY VALIDATED

The platform has been operationally validated in austere and denied environments across multiple theater-level exercises, including Pacific Pathways in the Philippines, JPMRC Yama Sakura in Japan, and STAFFEX, demonstrating its ability to automate staff workflows and accelerate decision cycles whether connected to enterprise systems or operating entirely at the tactical edge.



> G.H.O.S.T. is not a prototype or a proof of concept - it is a fielded, operational agentic AI system actively supporting live mission requirements within elements of the government, proven across both secure cloud environments and fully disconnected forward deployments.

PLATFORM ARCHITECTURE

Platform architecture.

USERS & ENTRY POINTS How your teams interact with the platform	**Business Analysts** **Finance Teams**	**Legal & Compliance**	**Military / Contracts**	**IT Operations**
G.H.O.S.T. ORCHESTRATION ENGINE Automated workflows | Agentic routing | Context management	**Workflow Router** One-click agentic automated workflows. Domain-aware routing	**Agentic Reasoning** Multi-step planning. Chain-of-thought logic	**Ghost Guardrails** Safety & compliance layer. Output validation	
RAG ENGINE Retrieval-Augmented Generation	**Document Ingestion** PDF, DOCX, XLSX, PPTX, images	**Semantic Indexing** Vector embeddings & smart chunking	**Cited Retrieval** Inline source links. Audit-trail ready	
SPECIALIZED AI AGENTS Pre-built | Configurable | Domain-expert	**Military Missions** Online	**Network Operations** Online	**Policy & Regulatory** Online	**IT Infrastructure** Online
SECURITY & GOVERNANCE Zero-trust | Audit-ready | Compliant	**RBAC & SSO** Role-based access. Identity federation	**Audit Logging** Full session trace. Immutable log chain	**Kill-Switch** Instant shutdown. Incident isolation	
CONFIGURABLE INFERENCE BACKEND Where the reasoning runs	**On-Device LLM** Fine-tuned local models. Zero external latency	**Private Cloud LLM** Your AWS / Azure tenant. No data boundary breach	**Air-Gap Fallback** Offline mode. Disconnected environments	**Workflow Templates** Pre-built & custom. One-click automation
ENTERPRISE DATA SOURCES What the agents read	**Documents** **ERP / CRM**	**Databases**	**REST APIs**	**SharePoint**

HOW G.H.O.S.T. WORKS

Select your agent. Upload your data. Get structured outputs.

01 SELECT YOUR AGENT

Choose from pre-built industry agents, or configure one to your workflow using the platform's standard tools.

02 UPLOAD YOUR DATA

Drag-and-drop documents, PDFs, reports, or connect data sources. No data leaves your environment.

03 GET STRUCTURED OUTPUTS

Receive citation-backed, audit-ready results - summaries, analyses, drafts, decisions in seconds.

CORE CAPABILITIES

DATA SOVEREIGNTY

Runs fully on-premise or private cloud. Your data never touches external servers.

ONE-CLICK WORKFLOWS

Automated multi-step processes reduce hours of work to a single button click.

MULTI-FORMAT INGESTION

Natively parses PDF, DOCX, PPTX, XLSX, plaintext, and images - no prep required.

RAG-ENABLED CITATIONS

Every output references the source document - full traceability for audits and reviews.

MODULAR & EXTENSIBLE

Open framework lets your IT team add custom agents, integrate APIs, or swap LLMs.

ROLE-BASED ACCESS

Granular user permissions, system-level logging, and SSO support for enterprise IT.

AT A GLANCE

< 48 hrs
Fully operational

90%
Reduction in time to resolution

100%
Data stays in your environment

5+
Industry verticals supported

HOW IS G.H.O.S.T. DIFFERENT THAN OTHER GEN A.I. PLATFORMS?

Mission-grade Gen A.I., off the shelf.

KEY FEATURES

01 On-device fine-tuned Large Language Models (LLMs)

02 Modular AI Agents for domain-specific, "One Click" Automated Workflows - no more experimental prompting

03 Optional Zero-connectivity, no-call-home architecture

04 Containerization and encryption meeting DoD & IC requirements

05 Role-based access control and system-level logging

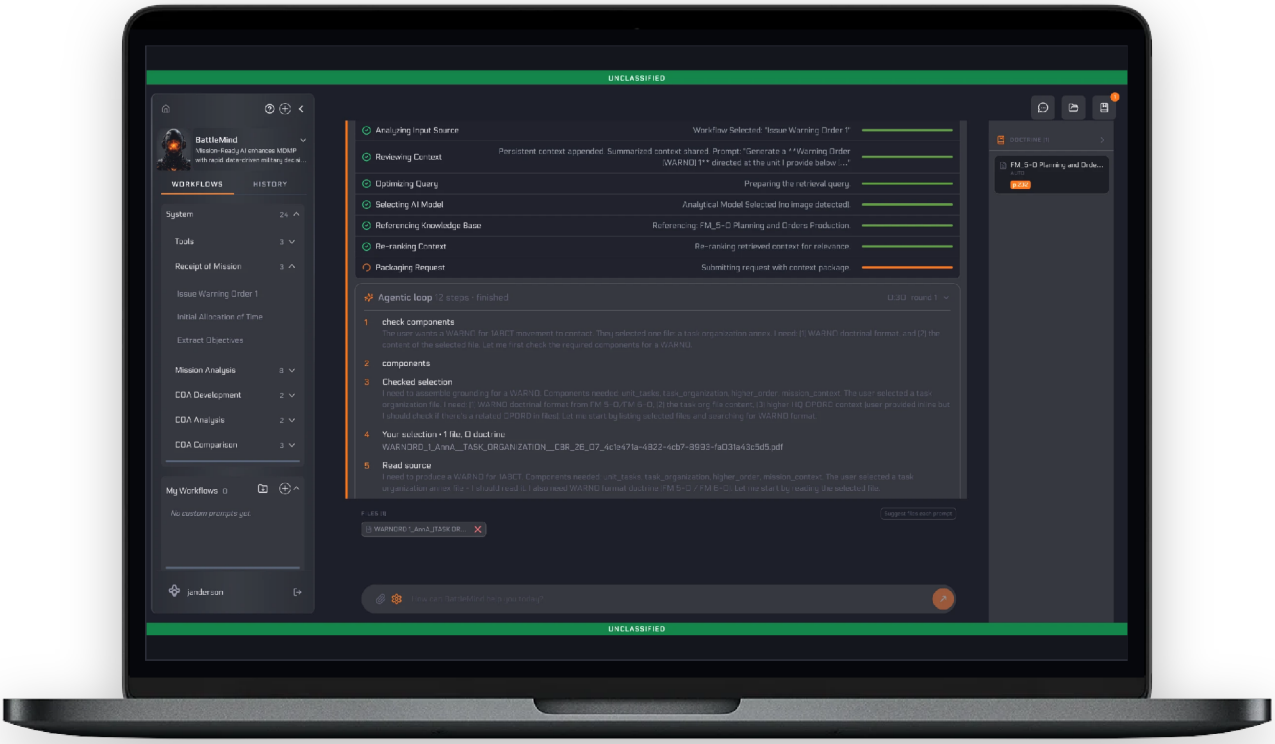
06 Multi-file UPLOAD processing and RAG-enabled citations

This enables a plug-and-play AI ecosystem deployable in forward-deployed field environments, SCIFs, Closed Enclaves, TOCs or Enterprise environments - without sacrificing capability or security.

WHY CHOOSE US?

Most enterprise GenAI platforms excel at generalized writing and analysis but remain fundamentally chat-centric tools. G.H.O.S.T.'s AI agents are built for operational time-to-resolution, driven by doctrine-aligned structures, unit-grounded references with full traceability, and mission-specific, one-click automated workflows that mirror how staffs actually plan, decide, and execute - rather than delivering a generic model wrapped in a user interface.

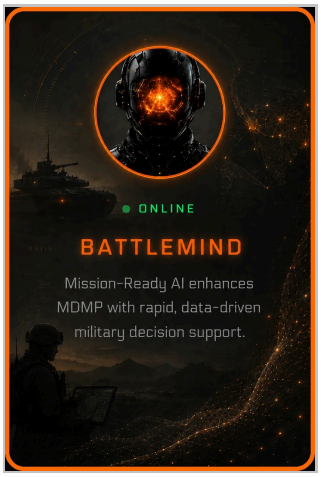
G.H.O.S.T. does not rely on the model as the differentiator. Value is created through the orchestration of workflows, retrieval-augmented generation, and context window that encode institutional knowledge, producing outputs with the fidelity, judgment, and structure typically achievable only by seasoned operators.



An agent working a tasking: the plan is shown first, every step is on screen, and the answer arrives cited.

FEATURED AI AGENTS

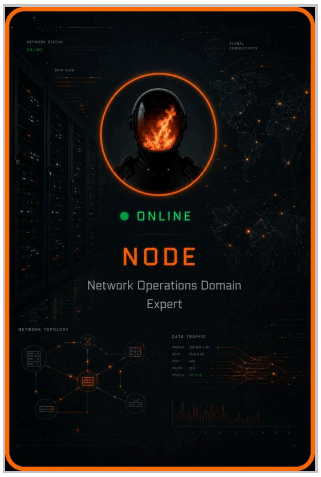
Three of the agents in the G.H.O.S.T. roster.



BATTLEMIND

Mission-Ready Military Decision Support

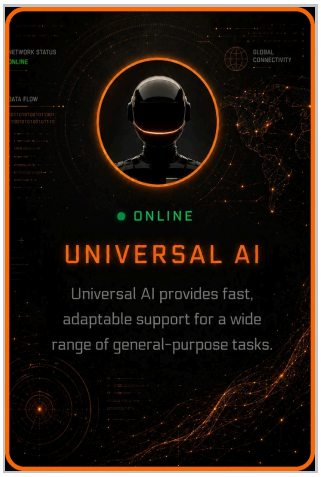
An MDMP-focused agent trained on Joint and Army doctrine, designed to accelerate Course of Action (COA) development, WFF product generation, image terrain analysis, mission analysis, and staff planning. Used in live Division HQ operations.



NODE

Network Operations & Infrastructure Analysis

A network diagnostic and automation agent capable of identifying misconfigurations, generating CLI commands, analyzing logs, and simplifying network troubleshooting - ideal for disconnected tactical networks, secured enclaves, and data centers.



UNIVERSAL AI

General-Purpose Assistant & Concierge

A general-purpose cognitive assistant pre-trained for high-assurance environments. Accepts structured and unstructured inputs (PPTX, DOCX, PDFs, imagery), and supports everything from intelligence production to acquisition planning.

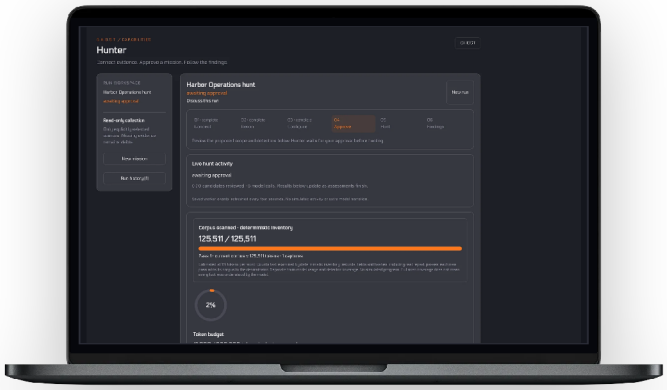
USING G.H.O.S.T. ON SCREEN

Choose the agent for the job. Approve the mission. Watch it work. Review the findings.



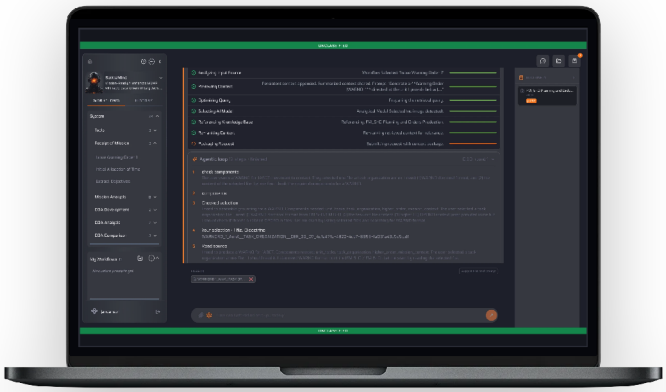
Step 1

Choose the agent for the job. Each agent is built for one mission area. Pick one from the roster, or let G.H.O.S.T. Mode choose the team for you.



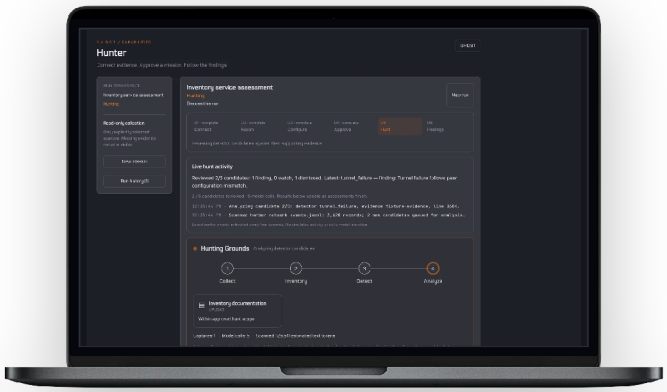
Step 2

Approve the mission. State the task in plain language and attach your files. The plan is shown first, and nothing runs until a person approves it.



Step 3

Watch it work. Every step shows on screen: what it searched, which documents it read, and which model it used.

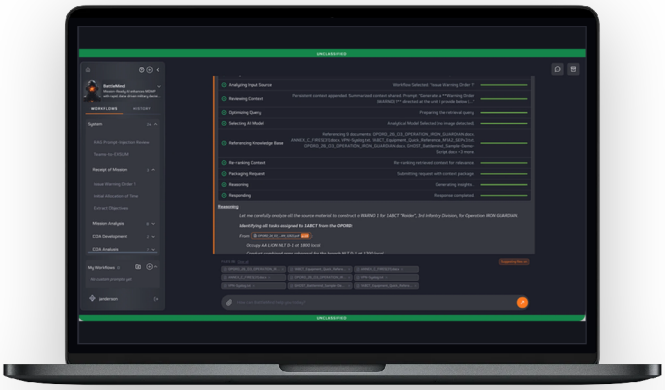


Step 4

Review the findings. Results arrive ranked, with evidence, assumptions and gaps, ready for a person to review and approve.

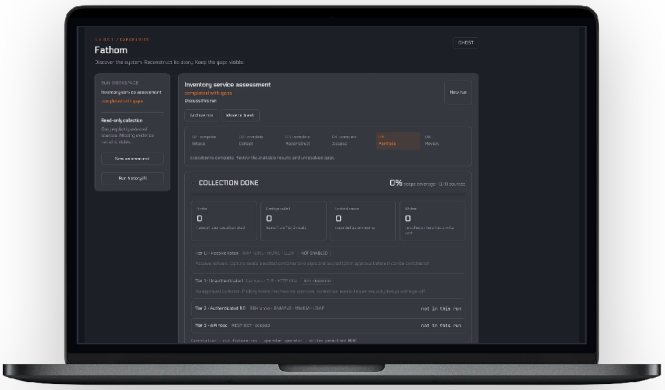
INSIDE THE PLATFORM

Every step is shown. A person stays in charge.



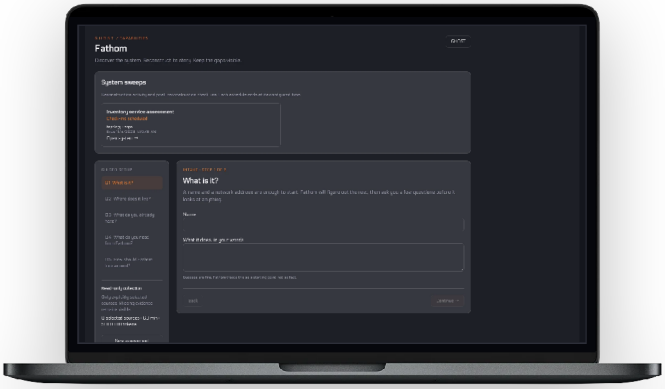
Agent workspace

Doctrine-grounded planning. A planning agent works an operations order: each step, source and model choice is on screen, and the answer cites the document it came from.



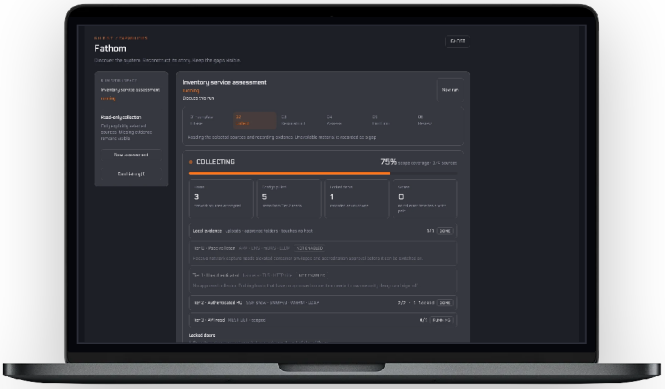
Assessment

Gaps are shown, not hidden. Fathom reconstructs a system from evidence and says plainly what the evidence does not prove.



Guided setup

A name and an address are enough. Fathom asks a few plain questions before it looks at anything, and scheduled sweeps keep the system watched afterward.



Collection

Read-only by design. Collection goes tier by tier. Locked doors are recorded as unknowns instead of being forced.

GOVERNMENT ALIGNMENT & COMPLIANCE

Aligned with the DoD AI Ethical Principles and CDAO Responsible AI guidance.

01 **Secure-by-design:** Operates entirely offline or on closed networks.

02 **Mission-focused:** Optimized for tactical and operational workflows.

03 **Auditable:** RAG-enabled outputs with inline citation tracing.

04 **Responsible AI:** No model training on operational data; no data leakage; auditable chain-of-thought reasoning.

COMPLIANT WITH OMB M-25-21

Supports modernization of federal digital infrastructure by:

01 Enabling secure, on-premise AI adoption.

02 Operating within Zero Trust architectures.

03 Providing modular, standards-based components for federal interoperability.

SUPPORTS EO 13960 AND EO 14179

01 Enhances AI transparency, accountability, and traceability.

02 Delivers safe, testable AI for sensitive mission workflows.

03 Empowers workforce augmentation and augmented intelligence, not full autonomy.

FEDRAMP PATHWAY SUPPORT

While G.H.O.S.T. is primarily offline, its architecture supports easy modular transition to FedRAMP-authorized cloud enclaves for hybrid deployment scenarios. G.H.O.S.T. is ATO ready for DoD and IC networks.

MISSION USE CASES

G.H.O.S.T. is not hypothetical - it's operational.

Our platform has been deployed and validated in high-stakes, real-world environments across multiple combatant commands and partner nations. Below are examples of successful field implementations.

1 U.S. ARMY - OPERATION FREEDOM SHIELD	Battlemind was deployed in support of live MDMP execution at the Division HQ in Schofield Barracks, HI. The AI agent enabled planners to accelerate mission analysis and generate WFF products that were briefed directly to the Commanding General.
2 U.S. ARMY - PACIFIC PATHWAYS (PHILIPPINES)	Battlemind supported combined planning operations in austere forward environments during bilateral exercises in the Indo-Pacific, validating G.H.O.S.T.'s disconnected deployment model.
3 3RD INFANTRY DIVISION - STAFFEX (ZAGAN, POLAND)	Battlemind was integrated into staff planning activities at the Division level in Europe, providing mission analysis products and COA development support during joint training in support of U.S. Army Europe and NATO partners.
4 JPMRC YAMA SAKURA, JAPAN	Used during coalition joint exercises, Battlemind demonstrated its ability to support multi-national staff planning by generating doctrinally aligned MDMP outputs without internet access or cloud dependency.

These use cases demonstrate G.H.O.S.T.'s effectiveness in enabling mission command and accelerating staff workflows under real-world operational constraints - no connectivity, high operational tempo, and doctrinal precision.

SECURITY AND GOVERNANCE

Built so a person stays the authority.

Controls that come with the platform, from the way it is packaged to the way every answer is recorded.

ROLE-BASED ACCESS AND SINGLE SIGN-ON

Granular permissions per user and per agent, with identity federation and optional multi-factor sign-in.

SMART CARD SIGN-IN

Optional CAC and PIV sign-in for deployments that require it.

A FULL AUDIT TRAIL

Every session is traced and logged in a tamper-evident chain, with an audit log viewer for administrators.

A KILL SWITCH

Instant shutdown and incident isolation when something looks wrong.

CONTAINERIZED AND ENCRYPTED

Packaged and encrypted to meet DoD and Intelligence Community requirements, built around NIST 800-53 controls.

UPLOADS ARE SCANNED

Files are checked for malware before they enter the document store.

ADVISORY ONLY, BY DESIGN

G.H.O.S.T. drafts, analyzes and recommends. A person stays the authority, and every product is staged for review.

NO OUTSIDE SERVICES AT RUNTIME

Fonts, models and search can all run on your side of the boundary. Web search is off unless you turn it on.

CLASSIFICATION BANNERS

Banners show on every screen, so people always know the level they are working at.

SEPARATE WORKSPACES

Agents, workflows and appearance can be set up separately for each organization or unit.

MODEL PROVENANCE

Every answer records which model actually produced it, so a reviewer can see what was behind the words.

A BUILT-IN GUIDE

G.H.O.S.T. Wingman explains features and what the signed-in user is allowed to do. It never acts for them and stores no conversation on the server.

UNIQUE ADVANTAGES

What G.H.O.S.T. does that others do not.

DEPLOYMENT	Operates fully offline or on classified networks; no internet required
SECURITY	No data exfiltration risk; operates within DoD containerized environments
SPEED TO USE	Pre-trained, pre-configured agents available out of the box
DOCTRINAL RELEVANCE	Agents trained using FM, JP, and MIL STD references for accuracy and trust
PORTABILITY	Runs on hardened laptops or TAC servers for field and austere deployment
EXTENSIBILITY	Open agent framework enables custom workflows, third-party LLM integration

ACQUISITION VALUE & OPERATIONAL ROI

G.H.O.S.T. delivers rapid, low-friction value at the edge - no integration delays, no network dependencies, and no guesswork. Unlike legacy enterprise AI solutions that require cloud migration, lengthy ATO processes, and heavy resourcing, G.H.O.S.T. can be fielded within days and demonstrate value within hours.

SPEED TO MISSION

01 Fully operational in under 48 hours in most field or SCIF environments.

02 Pre-trained models and embedded doctrine eliminate lengthy configuration or model tuning.

03 Runs on standard, hardened laptops or TAC servers with no additional infrastructure required.

PLUG-AND-PLAY ADOPTION, ROI AND LICENSING

Fielded within days. Value within hours.

PLUG-AND-PLAY ADOPTION

01 No cloud dependency. No internet required.

02 Drag-and-drop interface with built-in file parsers (PDF, DOCX, PowerPoint, plaintext, xlsx).

03 Role-based access controls and system logs allow seamless integration with existing OPSEC protocols.

04 Works on SIPR, JWICS, online, or disconnected networks.

QUANTIFIABLE ROI

01 Reduces planning timelines by up to 80% for complex COA and WFF generation.

02 Reduces reliance on contractor support and analyst overhead by automating high-friction workflows.

03 Enables warfighter and staff productivity gains without replacing personnel - augments human decision-making.

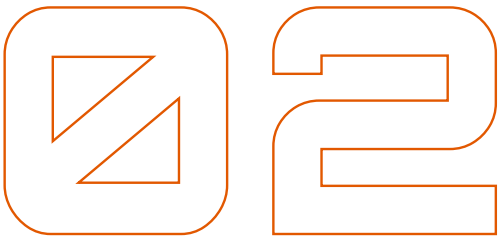
04 One G.H.O.S.T. license often offsets the cost of hours for planning and analysis.

FLEXIBLE LICENSING & PROCUREMENT

01 Available under FAR Part 12 - Commercial Products and Commercial Services - as a firm-fixed-price (FFP) annual license.

02 Already listed with a network of resellers, with access to GWACs and IDIQs.

03 Multi-year license options available for SCIF or on-premise customers. G.H.O.S.T. is commercial-off-the-shelf software: licensed to the general public under standard commercial terms and delivered to every customer - commercial and government.



Capabilities

Seven capabilities sit on top of every G.H.O.S.T. agent. Each one is the same platform, deployed your way.

AI Council	Convene more than one expert
M.A.C.S.	Wargame the plan before the world does
Agentic Loop	One prompt, many rounds
Smart Router	One box. The router decides.
Memory	Every session starts where the last one ended
Hunter	It doesn't wait for an alert
G.H.O.S.T. Mode	One request. The right agents.

CAPABILITY 01

AI Council

STANDING TEAM
MULTI-MODEL SEATS
ONE ACCOUNTABLE VERDICT

Convene more than one expert | Multi-agent deliberation

Real problems cross domains - and single-agent AI forces a staff to break the problem into pieces, ask separate questions, and stitch the answers together by hand. The stitching is where the mistakes live. The AI Council convenes several specialist agents on the same problem: state your tasking, and the platform recommends a chair agent to organize the effort and specialist seats that each bring their own persona, workflows, and document corpus to the same thread. The council deliberates in front of the operator and returns one verdict - then stands by for the next tasking.

HOW IT WORKS

1 TASKING
State the main goal

2 RECOMMEND
GHOST proposes chair + seats

3 THE COUNCIL
Chair and seats. Each seat: its own persona, workflows, corpus - and its own model, chosen per role

4 DELIBERATE
Challenge and reconcile - dissent stays visible

5 VERDICT
One synthesis, all reasoning kept

A STANDING TEAM - TASK IT AGAIN

ONE SEAT, ONE MODEL

Council seats are not clones of a single model. Each seat can run on a different underlying model - frontier or local - so the deliberation carries genuinely different perspectives, not one model agreeing with itself five times.

A STANDING TEAM

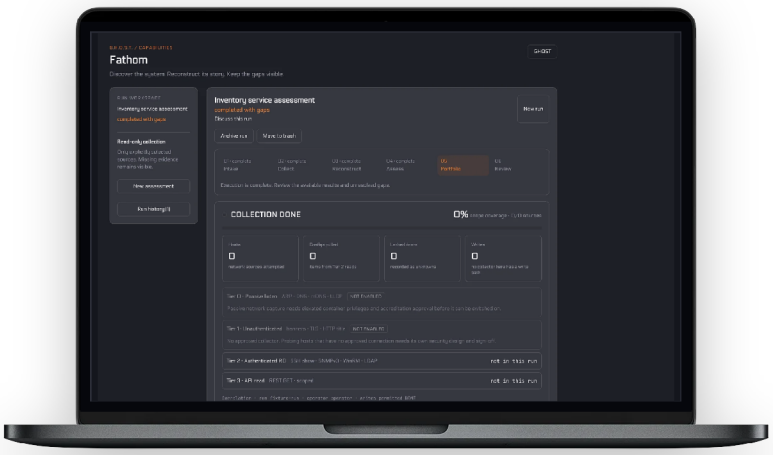
Convene once, then keep tasking it. The council is a long-running capability that carries context forward as the mission evolves - the team you assembled yesterday takes today's follow-on tasking without re- briefing.

ADVISORY-ONLY BY DESIGN

The operator keeps the decision. The council explains, assesses, debates, and recommends; the verdict arrives with each seat's reasoning traceable, and dissent recorded rather than hidden.

DEPLOYS YOUR WAY

Air-gapped on-prem, cloud-connected, or hybrid, the same platform - and the same council - deployed to match your security posture and mission need, built from the base agents you already field.



The AI Council convenes the G.H.O.S.T. agent roster - one chair, many perspectives, one accountable verdict.

CAPABILITY 02

M.A.C.S.

AN OPPONENT WITH A DOCTRINE
ROUNDS YOU CONTROL
A WRITTEN RECORD

Wargame the plan before the world does | Multi-agent competitive simulation

Every plan survives until it meets the other side. M.A.C.S. lets you meet the other side early - on your own terms. Stand up a friendly agent that knows your unit: your doctrine, your OPORD, your constraints. Stand up the adversary and arm it with everything you hold on them - a foreign tactics manual, intelligence, observed behavior. Then let them fight. Every hole the opposition finds is one you close before the real world finds it. Same engine in the boardroom: your company on one side, your three biggest competitors on the other - and the product launch or price change wargamed before you make the move.

HOW IT WORKS

1 YOUR SIDE

The friendly agent, built from your material. Doctrine, OPORD, unit details - or strategy, pricing, the plan you intend to execute

2 THE OPPOSITION

The adversary or your three biggest competitors. Foreign tactics manual, intel, observed behavior, market posture - everything you hold

3 THE ENGAGEMENT

Probe, counter, exploit - round after round, in front of the operator

Set the rounds: quick test or deep campaign

4 THE RECORD

Exposed assumptions, holes found, counters that worked - in writing

5 REVISE AND RE-RUN

Fix the plan, keep the sides standing, and fight it again until it holds

AN OPPONENT WITH A DOCTRINE

Not a devil's advocate - an adversary. The opposition argues from the tactics manuals, intelligence, and competitor context you load, so its moves are theirs, not a generic critique of your plan.

ROUNDS YOU CONTROL

One control sets how many rounds the sides trade. Dial it down for a fast pressure test before a decision brief; dial it up for a sustained campaign against your hardest problem.

ONE ENGINE, TWO ARENAS

Pit your unit against a foreign adversary - or your company against the competitors most likely to answer your next move. New product, price change, market entry: wargame it first.

ADVISORY-ONLY BY DESIGN

The commander - or the executive - keeps the decision. A simulation is an assessment, not an order; every finding in the record traces to the round and the material that produced it.

> M.A.C.S. puts a red team on retainer - wargame the plan before the world does.

CAPABILITY 03

Agentic Loop

MULTI-ROUND RETRIEVAL
SOURCE-CITED OUTPUT
GAPS STATED, NOT FILLED

One prompt, many rounds | Retrieve, verify, assemble

Ask most AI for a work product and it answers from one pass at your files - then writes confidently over whatever it did not find. The gap is where the mistakes live. The Agentic Loop is how a G.H.O.S.T. agent works a tasking the way a trained analyst would: it plans its own lines of inquiry, searches your document corpus and your trusted sources of truth round after round - regulation, doctrine, vendor documentation, standard operating procedures, whatever governs your work - opens and reads the passages that matter, checks what it still lacks, and only then assembles the product, delivered with its sources cited and its assumptions declared.

HOW IT WORKS

1 TASKING
One plain-language request

2 STAGE
Workflow, files, model and context selected

3 THE AGENTIC LOOP
Plan, search, read, verify, re-query. Round n - it keeps going until the evidence holds. Your corpus and your trusted sources of truth, searched and re-searched - every hit scored and traceable

4 ASSEMBLE
Built section by section from what it read

5 DELIVER
Product, sources, and stated gaps

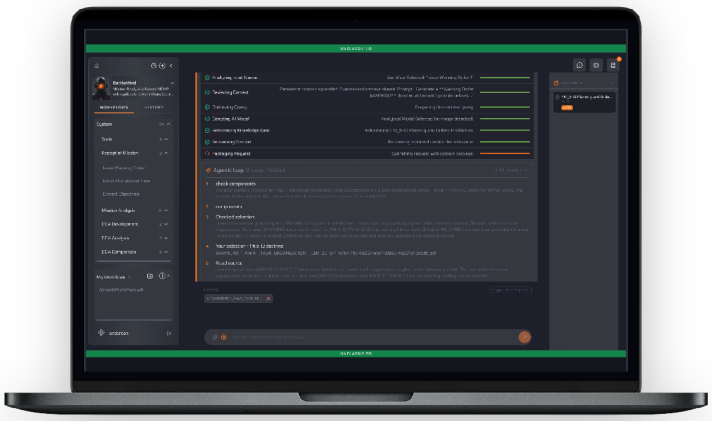
REFINE IT IN PLACE - THE NEXT PROMPT PICKS UP THE WHOLE THREAD

IT SHOWS ITS WORK
Every round is on screen as it happens - the queries it chose, the files and reference pages it hit, the relevance score of each, and the passages it opened. Nothing about how the answer was reached is hidden from the operator.

GROUNDED, NOT GUESSED
Products arrive with a grounded-on list naming the source documents and authoritative references behind them, so a reviewer can walk any line of the output back to the page it came from.

GAPS STATED, NOT FILLED
What the corpus does not contain is declared as an assumption or gap for your people to resolve - intent, timings, decisions still owed - rather than quietly invented to complete the page.

DEPLOYS YOUR WAY
Air-gapped on-prem, cloud-connected, or hybrid, the same platform - and the same loop - deployed to match your security posture and mission need, built from the agents you already field.



> The Agentic Loop is how a G.H.O.S.T. agent earns its answer - retrieve, verify, assemble, and cite your trusted sources of truth.

CAPABILITY 04

Smart Router

ALWAYS-ON CLASSIFIER
POLICY-SET THRESHOLD
REVERSIBLE REDACTION

One box. The router decides. | Complexity-aware model routing

Local-only AI gives up the hardest reasoning; frontier-only AI gives up control of the data - and of the bill. The work is not uniform, so the answer should not be either. The Smart Router scores every request for reasoning depth before anything moves, applies your policy and spend window, and sends the task to the AI model that actually fits - a small model on your own server for high-volume pattern work, a large model on your own server for full-document analysis, and an off-site frontier model only when the reasoning requires it. When frontier is selected, sensitive values are tokenized at the boundary and rehydrated on return: the reasoning leaves, the data does not.

HOW IT WORKS

1 REQUEST

Operator states the task

2 CLASSIFY

Custom-trained classifier scores reasoning depth - always on

Policy, threshold, spend window: sovereignty-weighted, balanced, quality-first, or frontier off - you control where the tokens go

3 ROUTING PROXY

Three AI models, one interface - the operator never picks

4 ONE MODEL SELECTED

Small AI model on your server: high volume, strict schema, pattern work. Large AI model on your server: full-document context, multi-step comparison. Frontier AI model off-site: long reasoning chains, maximum context

5 SANITIZE, REHYDRATE

PII becomes tokens before the call; the mapping stays in an on-prem vault and is never transmitted

ONE ANSWER RETURNS COMPLETE - NOTHING SENSITIVE LEFT

THE RIGHT MODEL, AUTOMATICALLY

Nobody picks a model. The classifier runs on every request and the proxy makes the call, so routine work is absorbed locally and frontier reasoning is reserved for the problems that actually need it.

REVERSIBLE REDACTION

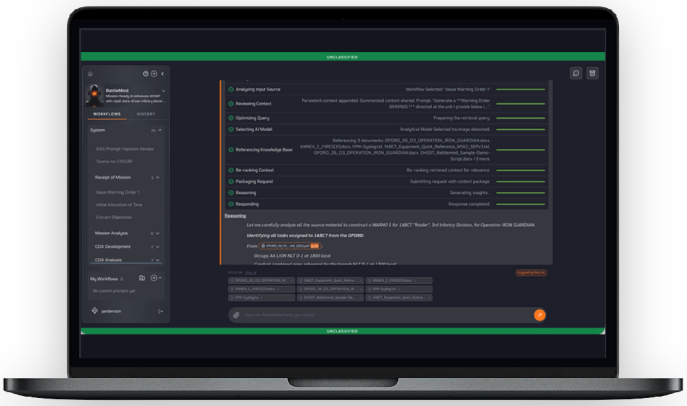
Sanitization is not deletion. Names, dates of birth, record numbers, and party names become tokens at the boundary and are restored from a local vault on return - a full-fidelity answer with no exposure.

COST STAYS UNDER YOUR CONTROL

Token spend is where most organizations quietly bleed money on AI. Frontier use is bounded by an envelope you set; when consumption runs past it, the router falls back to local tiers until the window resets - the work continues, the invoice does not climb.

ADVISORY-ONLY BY DESIGN

The router decides where reasoning happens, never what the answer should be. Routing decisions, policy trips, and sanitization steps are visible to the operator, and the human keeps authority over the work.



The Smart Router is the routing layer beneath the G.H.O.S.T. agent roster - every agent inherits it.

CAPABILITY 05

Memory

LEARNED, NOT UPLOADED
NOMINATED, THEN PROMOTED
PINNED, NEVER FORGOTTEN

Every session starts where the last one ended | Personal + team recall

Your files answer what you know; the Authority Library answers what the rules are. Neither answers the question every returning team asks: what have we already learned? MEMORY does. As GHOST works, it captures corrections, decisions and their rationale, conventions, and constraints - and users can pin a fact themselves with "remember this." Personal memory files itself and stays yours to edit. Nothing enters Team Memory silently: a curator promotes every entry, so the shared store stays small, current, and trusted. The pinned core rides in every session at a fixed, capped cost; what retires stays searchable.

HOW IT WORKS

1 CAPTURED

The session yields candidates. GHOST proposes what's worth keeping - or the user says "remember this"

2 CURATED

Personal memory files itself and stays user-editable; team entries wait in the queue until the curator promotes them

3 THE PINNED CORE

Rides in every session. Hard-capped, so the cost is fixed and known - retired entries stay searchable

4 IT STAYS CURRENT

Provenance and a last-confirmed date on every entry - stale ones retire to the archive

The memory budget: lean core, past this it's RAG

AUTHORITY OUTRANKS MEMORY: DOCTRINE, POLICY, AND SYSTEM FILES ALWAYS WIN; A CONFLICTING MEMORY IS FLAGGED, NOT OBEYED

LEARNED, NOT UPLOADED

RAG holds what you gave it; MEMORY holds what working together produced - the correction that ended a bad habit, the decision and why it was made, the convention nobody wrote down. Corrections are captured first because re-teaching is the most expensive mistake.

NOTHING ENTERS TEAM MEMORY SILENTLY

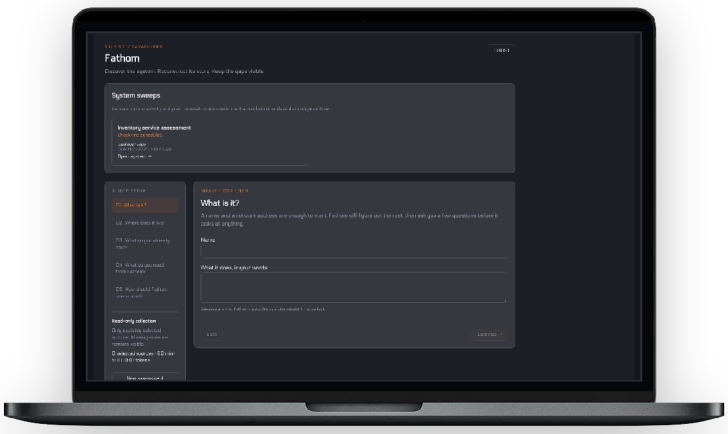
GHOST and users can only nominate a shared entry; a named curator promotes it from the queue. The whole store is visible to every member, and every entry carries its provenance - who, when, and the session it came from.

YOURS TO EDIT, YOURS TO SILENCE

Personal memory is fully user-editable - view, change, or delete any entry, any time - and any session can run with memory off. Team memory is off by default and enabled per team, deliberately.

MEMORY NEVER OUTRANKS THE RULES

Task instructions, the Authority Library, and System Files all take precedence. When a memory drifts out of step with doctrine or policy, GHOST raises the conflict for a human to resolve - memory doubles as a drift detector.



MEMORY is the G.H.O.S.T. platform's institutional recall - it learns as your team works, so no session starts from zero.

CAPABILITY 06

Hunter

SWEEP EVERYTHING, CHEAPLY
TRIAGE ONLY WHAT SURVIVES
CHASE WHAT IT CONFIRMS

It doesn't wait for an alert | Autonomous 24/7 hunt

Detection waits for a rule to fire, which means it only ever finds what somebody already knew to look for. HUNTER goes looking instead. Point it at a log repository, a SIEM export, a database, or a published API; it maps the terrain deterministically, takes its mission in plain language, then sweeps every record with detectors that cost nothing to run - and shows a model only the handful of candidates that survive. Full coverage is what makes it cheap, and cheap is what lets it run around the clock. Confirmations pivot into their own sub-hunts, and the loop keeps going inside the GPU duty cycle you set.

HOW IT WORKS

1 THE TERRAIN

Logs, SIEM, database, API. Enumerated and baselined first: inventory, entities, what normal looks like - 0 tokens

2 THE HUNT

Detectors sweep the whole corpus for free; the model triages only survivors - and confirmations pivot

The mission, in plain words: you state the intent and arm the behaviors; the plan comes back for approval, and that intent anchors every hunt that follows it

3 THE FINDINGS

Severity, confidence, evidence, next step. Incident case, timeline, coverage and assurance, rolling shift summary

Set the GPU duty cycle: share the card, or run flat out

4 IT DOESN'T STOP

Each finding becomes a hypothesis and runs down in its own sub-hunt; what that turns up sets the next - always against your intent

COVERAGE YOU CAN PROVE

Deterministic detectors read every record, so the assurance report can state exactly which detector examined what - including the hypotheses eliminated and the work deferred to the next period.

CANDIDATES ARE NOT FINDINGS

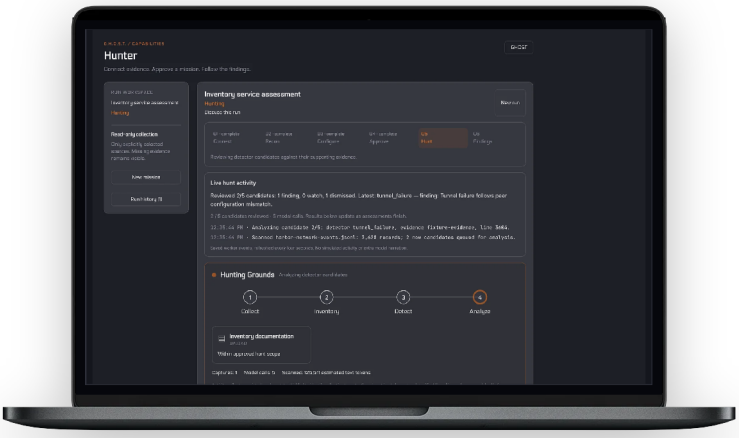
Triage weighs every candidate against the terrain baseline and reports the rejections alongside the escalations, with the reasoning that decided each.

EVERY FIND OPENS THE NEXT HUNT

A confirmation becomes a hypothesis and runs down in its own sub-hunt - foothold to traversal to staging - and whatever that turns up seeds the hunt after it, always against the intent you set.

ADVISORY-ONLY BY DESIGN

HUNTER detects, correlates, and recommends. Analysts and incident commanders keep the decision; every finding carries its detector, its triage reasoning, and its evidence on the face of the report.



HUNTER maps the terrain, works the mission, and never stops looking - the hunt runs while the queue sleeps.

CAPABILITY 07

G.H.O.S.T. Mode

DOMAIN DETECTED
AGENTS AUTO-SELECTED
ONE SYNTHESIZED PRODUCT

One request. The right agents. | Full-agentic orchestration

Real taskings cross specialties - a warning order carries foreign-language threat reporting and another service's fires; a protest response turns on acquisition procedure, regulatory reference, and cloud architecture at once. Nobody should have to know which agent to open first. In G.H.O.S.T. Mode the user states the task in plain language and attaches what they have. G.H.O.S.T. detects the mission domain, assembles the specialist team, binds each agent to the Automated Workflow that fits its part of the job, and runs them in sequence - returning one deliverable with grounding citations and an explicit assumptions-and-gaps list.

HOW IT WORKS

1 REQUEST
Plain language plus your documents

2 DOMAIN LOCK
Mission lexicons matched - ambiguity is reported, not guessed

3 ORCHESTRATION
Lead + specialists + bound workflows. Each agent is selected with a stated rationale and bound to the Automated Workflow for its part of the job

4 EXECUTION
Sequenced, logged - pauses to ask rather than assume

5 DELIVERABLE
One synthesis, cited - staged for human review

IT ASKS INSTEAD OF ASSUMING

When a step needs an input the request never supplied, execution pauses and asks one targeted question. Completed context is held in place - answering resumes the run rather than restarting it, and the answer is carried through the finished product.

A RUN BECOMES A WORKFLOW

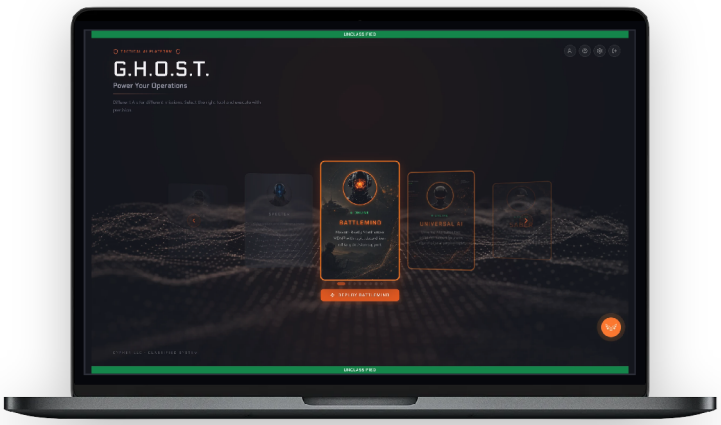
Any successful run can be saved as a named Automated Workflow with its steps, triggers, and input parameters captured. Questions asked mid-run become parameters asked up front, and the workflow can be shared to the unit or office catalog.

ADVISORY-ONLY BY DESIGN

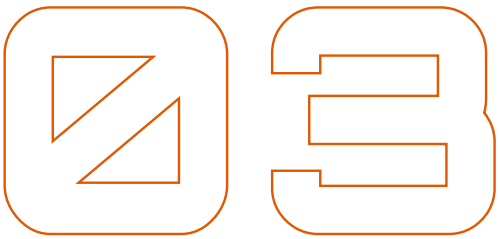
G.H.O.S.T. Mode assembles, analyzes, drafts, and recommends. Commanders, Contracting Officers, and counsel remain the sole authorities; every product is staged for human review with grounding citations, assumptions, gaps, and flagged conflicts on its face.

DEPLOYS YOUR WAY

Air-gapped on-prem, cloud-connected, or hybrid, the same platform deployed to match your security posture and mission need - orchestrating the base agents and workflows you already field.



> G.H.O.S.T. Mode is the orchestration layer over your G.H.O.S.T. agent roster - one request in, one reviewed product out.



The agent roster

Thirty-four agents, each built for one mission area. Pick one from the roster, or let G.H.O.S.T. Mode choose the team for you.

Mission planning	6 agents
Intelligence and cyber	8 agents
Acquisition and legal	5 agents
Networks and infrastructure	9 agents
Knowledge and people	5 agents
AI assurance	1 agent

THE ROSTER BY MISSION AREA

Thirty-four agents in six families.

MISSION PLANNING

- **BATTLEMIND** Army operations planning and the military decision-making process
- **J-ARC** Joint planning for combatant command and task force staffs
- **SABER** Marine Corps planning, step by step through MCWP 5-10
- **SUB-ZERO** Navy planning with an information warfare focus
- **AEROCENT** Air Force joint planning, from guidance to the air tasking order
- **SYNQ** Staff coordination and decision support for front offices

INTELLIGENCE AND CYBER

- **TALON** Cyber threat intelligence that turns reporting into defensive action
- **PERISCOPE** Open source collection planning and source vetting
- **VANTAGE** All-source analysis and finished intelligence
- **InWard** Insider threat and counterintelligence analytics
- **SIFT** Digital forensics and incident response analysis
- **BASTION** Security assessment and authorization through the risk management framework
- **QUARRY** Scoring mission data to decide what earns engineering time
- **PULSE** Finding the blind spots in security telemetry

ACQUISITION AND LEGAL

- **PRYME** Pre-award acquisition, from market research to solicitation
- **ARIA** Post-award contract administration
- **CADENCE** Requirements, deliverables and program rhythm
- **JURIS** Legal evaluation and regulatory analysis
- **REDACT** FOIA processing and declassification review

NETWORKS AND INFRASTRUCTURE

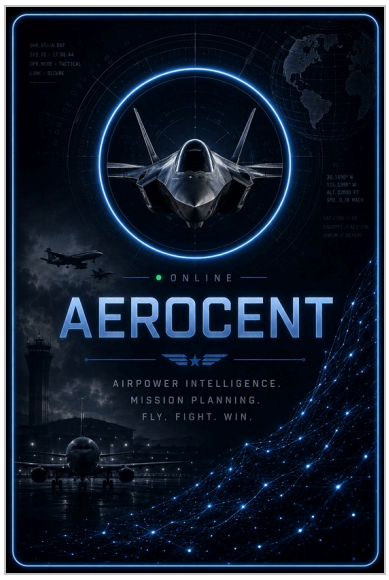
- **NODE** Network operations and troubleshooting across vendors
- **LYNX** Linux and cloud router engineering
- **SysTower** Server operations and infrastructure analysis
- **TRIAD** Multi-cloud engineering and cost
- **DataMind** Datacenter facilities and critical infrastructure
- **GridLock** Microgrid engineering, operations and testing
- **CASCADE** What a proposed change will actually touch
- **FATHOM** Rebuilding legacy systems from the evidence left behind
- **LATTICE** Database, data pipeline and data-quality analysis

KNOWLEDGE AND PEOPLE

- **UNIVERSAL AI** The general assistant every user starts with
- **AMBASSADOR** Translation and linguistic intelligence
- **LORE** Capturing what experts know before they leave
- **VELLUM** Governing what a retrieval system may answer from
- **HAVEN** Federal HR and workplace guidance

AI ASSURANCE

- **GAUNTLET** Offline testing to prove an AI system works before it ships



G.H.O.S.T. PLATFORM AGENT | 01 / 34

AEROCENT

Air Component Planning Intelligence

WHAT AEROCENT IS

AEROCENT is the U.S. Air Force operational-level planning specialist in G.H.O.S.T., built for the planners who run the Joint Planning Process for Air - AOC Strategy Division teams, Combat Plans and ISR Division planners, AFFOR A3/A5 staffs, and the APG and OPT leads who carry a plan to a JFC-approved JAOP.

AEROCENT ingests higher guidance, JFC orders, JIPOE extracts, staff estimates, force-status exports, wargame records, and prior plans - then structures the analysis, drafts planning products, audits strategy-to-task traceability, and recommends next actions with doctrine citations behind every finding.

WHAT SUCCESS LOOKS LIKE

- Mission-analysis and COA packages that survive planning-lead and commander decision-brief review with minimal editing
- Strategy-to-task traceability preserved from JAOP objective to AOD task to ATO mission - breaks flagged, never papered over
- Every doctrinal claim cited to the current publication; every quantity carries a source and data-as-of time
- A faster planning cycle overall: compressed timelines spent deciding, not assembling and reconciling
- A clean handoff from the approved JAOP into the air tasking cycle: AOD shells, apportionment recs, and conflict scans

68 Base Automated Workflows Across the Joint Planning Process for Air

1 - Plan Initiation 8
Stands the planning effort up right: Guidance Matrix, Planning Timeline, APG Roster Builder, Authority Matrix, and Compression Plan.

2 - Mission Analysis 10
Builds shared understanding: Task Derivation, COG Analysis, Strategy-to-Task Map, Mission Analysis Brief.

3 - COA Development 10
Generates distinct, valid options: COA Options Generator, Target-System Scheme, Apportionment Logic, COA Validity Test.

4 - COA Wargaming 9
Tests every COA against the red cell: Wargame Design, Red Cell Pack, Resource Ledgers, Decision Point Harvest.

5 - COA Comparison 8
Keeps the comparison honest: Criteria Dictionary, Weighting Audit, Comparison Matrix, Tradeoff Narrative, Dissent Ledger.

6 - COA Approval 7
Protects the decision record: Decision Brief Pack, Freshness Audit, Decision Menu, Approval Status Board, Guidance Tasker.

7 - JAOP Development 10
Turns the approved COA into a coordinated plan: JAOP Base Plan, Nesting Crosswalk, Branch & Sequel Cards.

8 - Tasking Bridge 6
Hands strategy to the tasking cycle: Strategy Handoff Pack, AOD Shell, Apportionment Rec, Cycle Conflict Scan.

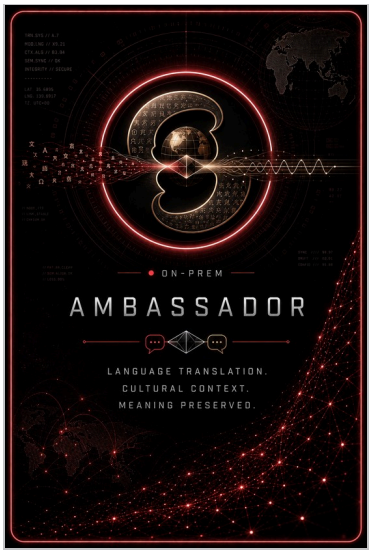
BUILT FOR THE OPERATIONAL ENVIRONMENT

DEPLOYS YOUR WAY
Air-gapped on-prem, cloud-connected, or hybrid, the same platform, deployed to match your security posture and mission need.

ADVISORY-ONLY BY DESIGN
The JFACC and JFC keep every decision - COA selection, apportionment, risk acceptance, plan approval, release. AEROCENT structures, drafts, and audits; unconfirmed authorities are marked UNCONFIRMED and every product ships under a DRAFT stamp.

GROUNDED IN TRUSTED SOURCES OF TRUTH
AEROCENT's analysis draws on your uploaded planning record and current public doctrine - JP 5-0, JP 3-30, AFDP 5-0, AFMAN 13-1AOCV3 - and flags anything requiring live verification.

A BASE AGENT, TAILORED ON-SITE
AEROCENT ships as a base agent, refined on-site with your teams - your SOPs, battle rhythm, and the nuance only your planners know.



G.H.O.S.T. PLATFORM AGENT | 02 / 34

AMBASSADOR

Translation & Linguistic Intelligence

WHAT AMBASSADOR IS

Ambassador is the translation and linguistic-intelligence agent in G.H.O.S.T., built for government and contractor linguists, all-source intelligence analysts, translation program managers, and QA reviewers. It works the whole lifecycle, from source intake and scoping through translation, dialect and origin analysis, terminology audit, and analyst handoff.

Ambassador ingests source documents, scans, transcripts, machine-translation drafts, glossaries, do-not-translate lists, and prior products - then translates, reads dialect and register from named markers, audits terminology, reviews drafts, and extracts entities and relationships, with a glossary entry, prior product, or source segment behind every non-obvious call.

WHAT SUCCESS LOOKS LIKE

- Translations a reviewer can validate in one pass - every non-obvious rendering shows its source segment and the standard applied
- Dialect, register, and origin calls that state their markers and a confidence level, and return INSUFFICIENT rather than a guess
- Ambiguity surfaced rather than resolved: competing readings held side by side for the analyst
- Terminology, transliteration, and do-not-translate consistency held across a product line
- A faster pipeline overall: intake triage, effort scoping, and queue priority tracked continuously

40 Base Automated Workflows Across the Full Translation Lifecycle

1 - Intake & Scoping 8

Starts the work on solid ground: source intake triage, scope and effort planning, OCR triage and routing, language and script ID, marking checks, tasking package build, corpus inventory, and queue priority planning.

2 - Translation Core 8

Carries the meaning across: cultural-context passes, gist and summary translation, full-document rendering, transliteration planning, register adaptation, idiom and metaphor notes, and consistent entity-name rendering.

3 - Dialect & Origin 8

Reads what the language itself reveals: regional dialect reads, register and period assessment, speaker-origin hypotheses, code-switch analysis, authorship consistency, sociolect markers, ambiguity surfacing, and an attribution confidence brief.

4 - Terminology & QA 8

Holds the product line honest: terminology consistency audits, second-set-of-eyes review, back-translation checks, DNT list building, glossary entry drafting, transliteration standard audits, MT post-edit planning, and onboarding packs.

5 - Exploit & Report 8

Turns translated material into product: entity and relationship mapping, multi-source cross-reference, read-only queries against approved translation databases, watchlist hit review, analyst handoff packages, QC metrics, and citation packs.

BUILT FOR THE MISSION ENVIRONMENT

DEPLOYS YOUR WAY

Air-gapped on-prem, cloud-connected, or hybrid, the same platform, deployed to match your security posture and mission need. Ambassador ships today as a Local-tier agent, so it runs where the material cannot leave.

ADVISORY-ONLY BY DESIGN

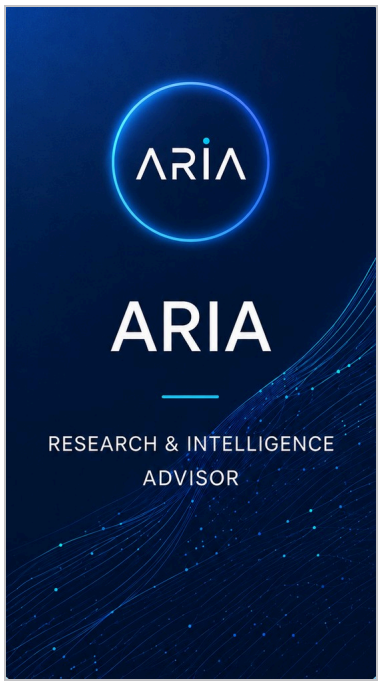
A qualified human linguist remains the sole authority over every rendering; the analyst owns operational meaning and the security officer owns markings. Every line is separated into Source-supported, User-provided, and Inference, with a confidence level and a citation.

GROUNDED IN TRUSTED SOURCES OF TRUTH

Ambassador's analysis draws on your uploaded sources, approved glossaries and terminology databases, do-not-translate lists, transliteration standards, and prior products - and flags anything requiring verification.

A BASE AGENT, TAILORED ON-SITE

Ambassador ships as a base agent and is refined on-site with your teams - incorporating your language pairs, glossaries and DNT lists, and the nuance only your linguists know.



G.H.O.S.T. PLATFORM AGENT | 03 / 34

ARIA

Post-Award Acquisition Intelligence

WHAT ARIA IS

ARIA (Acquisition Resource & Intelligence Assistant) is the post-award contract administration copilot in G.H.O.S.T. It picks up where PRYME hands off at award, supporting the teams that administer task orders on behalf of client agencies - Contracting Officers, CORs, project managers, and performance managers carrying a contract through performance to closeout.

ARIA ingests task orders, modifications, PWS / SOW, CDRL matrices, MSRs, invoices, burn reports, AFDPs, and CPARS evaluations - then identifies risks, explains obligations in plain language, drafts compliant artifacts, and recommends next actions with FAR, agency supplement, and vehicle citations behind every finding.

WHAT SUCCESS LOOKS LIKE

- Deadline risks flagged before they become incidents - OASIS+ 30-day task order reporting, quarterly CAF remittance, deliverable due dates, and CPARS response windows
- Burn rate and funding runway visibility, with over-burn and under-burn trajectories caught early
- Award fee cycles supported end to end: contractor self-assessments reviewed against AFDP criteria, and CPARS narratives prepared on schedule
- Modifications and scope changes kept clean: SF-30 discipline, constructive-change exposure flagged on sight
- Faster administration overall - every recommendation traceable to a FAR part, agency supplement section, contract clause, or vehicle provision

13 Base Automated Workflows Across the Post-Award Lifecycle

1 - Reporting & Status 5

Keeps the record current and honest: MSR generation, review, and analysis, deliverable matrix tracking against CDRL deadlines, and burn rate analysis with funding runway and EAC / ETC projections.

2 - Compliance & Audits 3

Protects the compliance posture: invoice package review before payment approval, OASIS+ vehicle compliance (SRP reporting and CAF math), and QCP gap analysis against contractual requirements.

3 - Award Fee & Ratings 3

Keeps the performance record defensible: contractor self-assessment (CSA) review mapped to AFDP criteria, and CPARS narrative preparation inside the evaluation window.

4 - Transitions & Mods 2

Manages change without exposure: contract modification review with SF-30 and scope discipline, and transition readiness checks for transition-in and transition-out / closeout.

BUILT FOR THE CONTRACT ENVIRONMENT

DEPLOYS YOUR WAY

Air-gapped on-prem, cloud-connected, or hybrid, the same platform, deployed to match your security posture and mission need.

GROUNDED IN TRUSTED SOURCES OF TRUTH

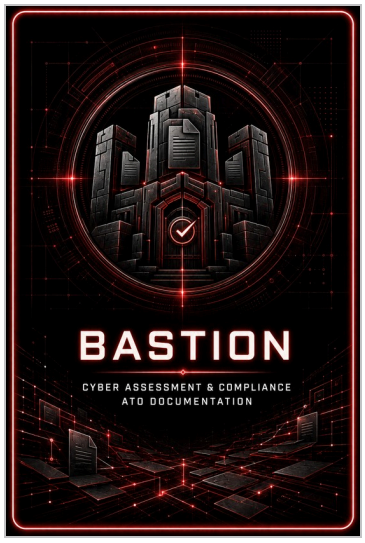
ARIA's analysis draws on your uploaded contract documents, the FAR and your agency supplement, and vehicle master contracts - the uploaded contract is always primary, and anything that requires verification against a live authoritative source gets flagged.

ADVISORY-ONLY BY DESIGN

The Contracting Officer remains the sole authority to bind the government or change contract terms. ARIA explains, assesses, drafts, and recommends; every finding is separated into OBSERVED, ASSESSED, and RECOMMENDED, with a confidence level and a traceable regulatory basis.

A BASE AGENT, TAILORED ON-SITE

ARIA ships as a base agent and is refined on-site with your teams - incorporating agency templates, thresholds, and the nuance only your workforce knows.



G.H.O.S.T. PLATFORM AGENT | 04 / 34

BASTION

Security Assessment & Authorization

WHAT BASTION IS

Bastion is the cyber security assessment and authorization (A&A) specialist in G.H.O.S.T., built for the people who carry a system through the RMF - ISSOs and ISSMs maintaining SSPs and POA&Ms, security control assessors turning test results into Security Assessment Reports, ISSEs writing control-implementation statements, and system owners and package managers preparing ATO submissions.

Bastion ingests scan exports, SCAP and STIG results, pen-test reports, system documentation and prior A&A packages - then maps findings to NIST SP 800-53 controls and CCIs and drafts POA&M entries, control statements, gap analyses and SAR sections, each assertion carrying a [Source / Section / Finding-ID] citation.

WHAT SUCCESS LOOKS LIKE

- Cited draft artifacts a human assessor can finalize with minimal editing - risk write-ups, control statements, POA&M entries, gap analyses and security impact analyses
- Authoritative scores stay with the source. Scanner CVSS, STIG CAT levels and assessor Satisfied / Other Than Satisfied calls are reported as stated; Bastion's prioritization is labeled ADVISORY
- POA&M entries entry-ready for eMASS or Xacta, every required field present and dated
- One consistent package story: SSP, SAR, POA&M and categorization checked for boundary and control-set drift
- Faster assessment overall: scan deltas, cross-tool deduplication, SLA aging and staleness flags on old exports

16 Base Automated Workflows Across Assessment, Compliance, Authorization and Change Review

1 - Vuln & Testing ┐
Turns raw exports into defensible findings: vulnerability assessment with POA&M drafting, security testing documentation, scan delta analysis, STIG rollup by CAT with CCI mapping, pen-test mapping to MITRE ATT&CK, and SLA aging.

2 - Compliance ┐
Holds the control story straight: control mapping and gap analysis, implementation statements (what / how / who / evidence), CNSSI 1253 baseline tailoring, overlay applicability, inheritance matrices, and CCI traceability.

3 - Authorization 1
Gets the package submission-ready: ATO package assembly with completeness, marking and cross-reference checks, gap lists, and draft risk-summary sections.

4 - Change Review 1
Keeps the authorized baseline honest: RFC security analysis of affected controls, boundary changes, new data flows and residual risk, drafted as the security impact analysis.

BUILT FOR THE ACCREDITATION ENVIRONMENT

DEPLOYS YOUR WAY

Air-gapped on-prem, cloud-connected, or hybrid, the same platform, deployed to match your security posture - with a frontier persona and a small-model Local persona.

GROUNDED IN TRUSTED SOURCES OF TRUTH

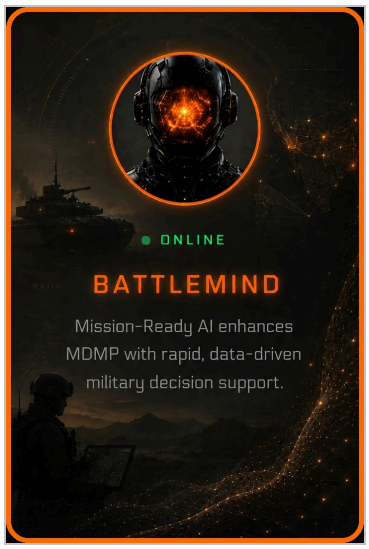
Bastion reads your uploaded exports and packages first - scanner output, STIG and SCAP checklists, pen-test reports, SSPs, SARs and POA&Ms - against NIST SP 800-53, 800-53A, 800-30, ICD 503, CNSSI 1253 and DISA STIGs. Stale dates and out-of-boundary assets are flagged.

ADVISORY-ONLY BY DESIGN

The SCA owns compliance determinations and the risk recommendation; the AO owns risk acceptance and authorization. Bastion never implies an ATO outcome, declares a control compliant, or marks risk accepted. Products separate fact, context and inference, stamped DRAFT - FOR ASSESSOR REVIEW.

A BASE AGENT, TAILORED ON-SITE

Bastion ships as a base agent and is refined on-site with your teams - incorporating your A&A templates, baselines, overlays and SLA thresholds. The 16 workflows are the built catalog today.



G.H.O.S.T. PLATFORM AGENT | 05 / 34

BATTLEMIND

Mission-Ready Military Decision Support

WHAT BATTLEMIND IS

Battlemind is the Army operations-planning and MDMP specialist in G.H.O.S.T., built for the brigade and battalion staffs that run the Military Decision-Making Process - S-3 / G-3 operations officers and plans chiefs, S-2 / G-2 intelligence officers, the XO or COS integrating the process, and the staff sections building schemes of support and annexes.

Battlemind ingests the higher headquarters order, IPOE products, running estimates, and terrain data - then produces step-correct staff products: restated missions, task and constraint lists, fact-assumption ledgers, COA statements, decision matrices, and OPORD / WARNORD / FRAGORD text - every product step-labeled, cited to its source.

WHAT SUCCESS LOOKS LIKE

- Step-labeled DRAFT products a planner carries into the mission analysis, decision, or orders brief with minimal editing
- Every specified, implied, and essential task traced to its source paragraph in the higher order - subordinate tasks written as task plus purpose
- Facts separated from assumptions - each assumption tested for validity and carrying the action that converts it to fact
- Doctrine you can check - every workflow anchored to the FM 5-0 paragraph it implements, backed by the CALL MDMP handbook
- A faster planning cycle overall: the one-third / two-thirds rule protected, and products that survive XO reconciliation and higher-headquarters crosswalk

77 Base Automated Workflows Across All Seven MDMP Steps

1 - Receipt of Mission 9
Gets planning started right: staff alert roster, planning-tools check, objective extraction, initial time allocation, WARNORD 1 drafting, and RFI log initiation.

2 - Mission Analysis 18
Builds the defendable foundation: OAKOC terrain analysis, task determination, fact-assumption ledger, CCIR and EEFI development, and problem and mission statements.

3 - COA Development 10
Generates options worth analyzing: relative combat power, COA options generation, force array, COA statements, and screening against the five validity criteria.

4 - COA Analysis 14
Stress-tests every option: wargame method selection, wargame turn capture, decision points, branches and sequels, DST / DSM builders, and synch matrix first drafts.

5 - COA Comparison 6
Keeps the comparison honest: advantage and disadvantage analysis, the weighted decision matrix with the math shown, sensitivity review, and the staff recommendation.

6 - COA Approval 6
Turns decision into direction: the COA decision brief, refined commander's intent, CCIR and EEFI refinement, WARNORD 3, and commander-approval items isolated in one list.

7 - Orders Production 10
Closes plan to order: the five-paragraph OPORD draft, annex builder, reconciliation and crosswalk, confirmation briefs, and FRAGORD drafting.

Tools 4
Staff utilities riding alongside the process: text summarization, Teams-to-EXSUM conversion, RAG prompt-injection review, and an acronym and term decoder.

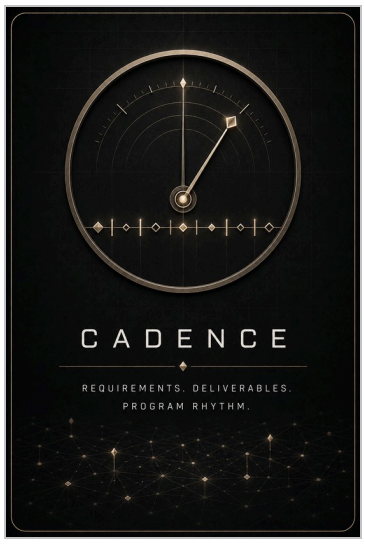
BUILT FOR THE ARMY ENVIRONMENT

DEPLOYS YOUR WAY
Air-gapped on-prem, cloud-connected, or hybrid - the same platform, deployed to match your classification level and mission environment, from garrison to the tactical edge.

ADVISORY-ONLY BY DESIGN
The commander decides - COA approval, CCIRs, EEFIs, risk acceptance, and the signature stay human. Battlemind analyzes, drafts, recommends - and stamps everything DRAFT.

GROUNDED IN TRUSTED SOURCES OF TRUTH
Your uploaded orders, running estimates, and IPOE products, plus FM 5-0, FM 3-90, ATP 2-01.3, and ATP 5-19 - sources cited behind every claim, gaps named rather than filled.

A BASE AGENT, TAILORED ON-SITE
Battlemind ships as a base agent and is refined on-site with your formation - incorporating unit SOPs, reporting formats, and the nuance only your staff knows.



G.H.O.S.T. PLATFORM AGENT | 06 / 34

CADENCE

Requirements, Deliverables & Program Rhythm

WHAT CADENCE IS

Cadence is the program management, requirements, and deliverables specialist in G.H.O.S.T., built for the people who keep a program's battle rhythm running - program and project managers, PMO staff, requirements and CDRL leads, control-account managers, schedulers, and risk managers.

Cadence ingests the SoW/PWS, CDRL lists, requirement sets, IMS exports, EVM data, notes, status inputs, and risk and action logs - then structures requirements, traces deliverables, computes schedule and earned-value figures, and drafts artifacts with a source behind every entry.

WHAT SUCCESS LOOKS LIKE

- A citable DRAFT a PM can finalize with minimal editing - every entry tied to a named source, every gap flagged instead of filled
- Numbers that reproduce. SV, CV, SPI, CPI, EAC, VAC and TCPI computed only from supplied BCWS/BCWP/ACWP, each with its formula and inputs shown
- Requirements that hold at a review board: singular, verifiable, one of four methods (Inspection, Analysis, Demonstration, Test), traced both directions
- Deliverables tracked across exactly six states - Not Started, In Work, Internal Review, Submitted, Accepted, Rejected/Rework - with CDRL, IMS and status reconciled
- Faster reporting overall, and honest status: RAG defined by rule, risks scored on the program's 5×5 scale, conflicts never averaged away

48 Base Automated Workflows Across the Full Program Lifecycle

1 - Program Setup 6
Stands the program up right: charter and PMP shells, WBS builder, stakeholder register, RACI with single-A enforcement, battle-rhythm design, and the baseline register.

2 - Requirements 7
Makes the requirement set defensible: requirements extract with preserved IDs, RTM builder, baseline compare (Added/Modified/Removed), orphan and conflict scan, and I/A/D/T verification.

3 - Deliverables 7
Holds CDRL discipline end to end: tracker across the six states, completeness checks, revision reconcile, submission packages, and acceptance logs.

4 - Schedule & EVM 6
Reads the schedule honestly: IMS health checks for missing logic, constraints and negative float, critical-path narration, milestone tracking, EVM metrics packs, EAC scenarios.

5 - Status, Comms & Risk 12
Runs the reporting and risk cycle: status drafts, KPI tables with formula and data-as-of, trend deltas, minutes and action aging, risk harvest, and register updates.

6 - Change Control & Closeout 10
Protects the baseline and ends clean: change requests with impact and authority named, CCB packages, the baseline change log, closeout checklists, and the final CDRL audit.

BUILT FOR THE PROGRAM ENVIRONMENT

DEPLOYS YOUR WAY

Air-gapped on-prem, cloud-connected, or hybrid, the same platform, deployed to match your security posture and mission need.

ADVISORY-ONLY BY DESIGN

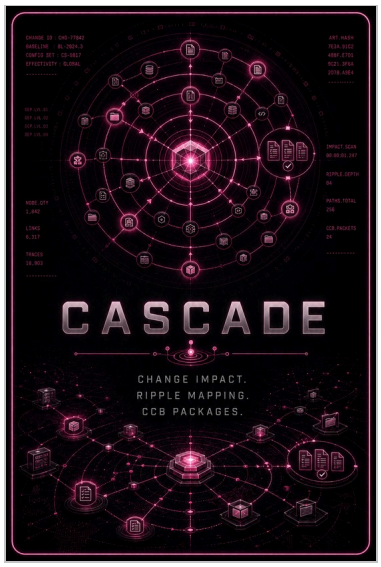
Cadence reads provided exports and never connects to a scheduling or EVM system: no live IMS, no live cost tool, no writes back. The PM, control-account managers, the contracting officer and the customer keep final authority on every commitment, baseline change, acceptance and released report; products are marked DRAFT - FOR PM REVIEW.

GROUNDED IN TRUSTED SOURCES OF TRUTH

Cadence works only from what you provide - SoW/PWS, CDRLs, requirement sets, IMS exports, EVM data, registers and prior snapshots - and names the snapshot behind every trend.

A BASE AGENT, TAILORED ON-SITE

Cadence ships as a base agent and is refined on-site with your teams - your report formats, CDRL templates, variance thresholds, and risk scale.



G.H.O.S.T. PLATFORM AGENT | 07 / 34

CASCADE

Semantic Change-Impact Analysis

WHAT CASCADE IS

CASCADE is the semantic change-impact analysis specialist in G.H.O.S.T., built for the people who must know what a proposed change actually touches before a board votes on it: configuration managers and CCB secretariats, systems and network engineers, ISSMs, requirements managers and test leads, runbook and training owners, contract managers and CORs.

CASCADE ingests a change package - ECR, ECP, policy revision, architecture update, interface change, or control-baseline shift - plus an explicit document manifest, then extracts a numbered delta register, traces dependency edges into a ringed impact map, hunts stale artifacts, and drafts the CCB decision package, with an edge type, a cited path, a ring, a severity, and a confidence grade behind every impact.

WHAT SUCCESS LOOKS LIKE

- Every claimed impact rides a named dependency edge. No edge, no claim: each entry carries a ring (0, 1, or 2+), an edge type, a citation, and a CONFIRMED, PROBABLE, POSSIBLE, or SPECULATIVE confidence grade.
- Escaped impacts caught before approval, not after. Interface changes get both-ends treatment, decommissions get a full name sweep, and boundary-touching impacts route to the ISSM.
- Impact registers where every affected artifact names an owner, a required action, and a priority
- Faster change packages overall: delta extraction, ring assignment, and consistency checks in one pass
- Pre-existing debt kept off the change's bill: stale artifacts in their own register, severities set by cited basis

50 Base Automated Workflows Across the Change-Impact Lifecycle

1 - Manifest & Delta 6
Bounds the analysis before tracing begins: Manifest Census, Change Package Intake, Baseline Delta, and Manifest Gap Report.

2 - Impact Radius 7
Rings the ripple and proves it: Impact Radius, Ring Assignment Review, Edge Trace Audit, Latent Triage Sheet, and Radius Compare.

3 - Requirement & Interface Ripple 13
Keeps traceability honest and treats the highest fan-out delta class at both ends: Requirement Ripple, RTM Break Check, Interface Ripple, Both-Ends Mirror Check, Decommission Sweep.

4 - Control Ripple 6
Protects the authorization posture: Control Ripple, Boundary Touch Check, SSP Consistency Sweep, POA&M Impact Map, and Boundary Alert Pack.

5 - Ops, Training & Contract Ripple 11
Keeps operational and contractual documents current: SOP Runbook Sweep, Training Ripple, Forms & Reports Sweep, AWF & Persona Ripple, Contracts Routing Brief.

6 - Decision & Hygiene 7
Closes the loop for the board: CCB Decision Pack, Register Consistency, Stale Artifact Hunt, Escaped Impact Review, Re-Baseline Brief.

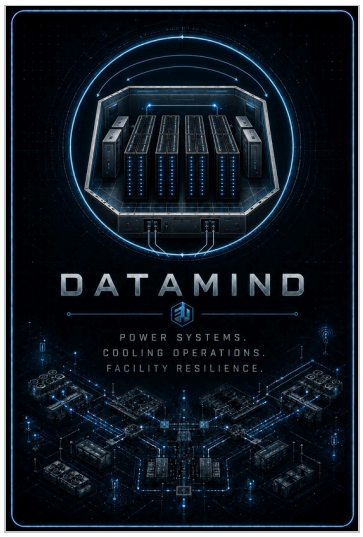
BUILT FOR THE CONFIGURATION ENVIRONMENT

DEPLOYS YOUR WAY
Air-gapped on-prem, cloud-connected, or hybrid, the same platform, deployed to match your security posture and mission need.

GROUNDED IN TRUSTED SOURCES OF TRUTH
CASCADE's analysis draws on your change package and manifest - baselines, RTMs, ICDs, SSP and POA&M exports, SOPs, CDRLs - flagging anything needing live-source check.

ADVISORY-ONLY BY DESIGN, AND MANIFEST-BOUNDED
CASCADE works over a bounded, analyst-supplied document set, not a live repository crawl, and every product closes with a mandatory UNSCANNED SURFACES statement listing absent artifact classes and trace exits. Configuration managers, CCB chairs, chief engineers, document owners, ISSMs, and COs keep final authority.

A BASE AGENT, TAILORED ON-SITE
CASCADE ships as a base agent, refined on-site with your teams - your baseline structure, CCB formats, and workforce nuance.



G.H.O.S.T. PLATFORM AGENT | 08 / 34

DataMind

Datacenter Facility & Critical Infrastructure Analysis

WHAT DATAMIND IS

DataMind is the datacenter facility agent in G.H.O.S.T.: the building and its physical plant, not the servers in the racks. It covers the power path from utility and generator to the rack PDU, the cooling path to the cabinet inlet, and the space itself, for facility managers, MEP and critical-facility engineers, data hall and NOC staff, capacity planners and colocation managers.

DataMind ingests BMS and DCIM exports, branch-circuit metering, one-line and riser diagrams, commissioning and IST reports, PM records and alarm logs - then traces redundancy, recomputes headroom, assesses inlet conditions, reconstructs outages and drafts facility products, every number tied to a measurement point and source file.

WHAT SUCCESS LOOKS LIKE

- A false 2N found before it fails. Both cords of every critical load traced to source and each shared upstream device named, with redundancy reported separately from concurrent maintainability - N+1 behind one path still drops load.
- Headroom you can place load against. Design, installed, allocated and measured demand reconciled at every level and reported in the N-1 case against the 80 percent continuous-load derating.
- Inlet conditions judged at the cabinet inlet against the ASHRAE TC 9.9 recommended envelope and the allowable A1-A4 class
- PUE stated with its measurement category (Cat 0 through Cat 3) and period, so the figure is comparable and the gaming patterns are visible
- Availability read as downtime minutes per year with MTBF and MTTR, every finding carrying an S1 Critical to S5 Advisory severity and a confidence level

40 Base Automated Workflows Across the Critical Facility Stack

1 - Power & Redundancy 5
Proves the power path or finds where it is single-corded: One-Line Review, UPS Health Read, Generator Readiness, ATS & Transfer Review, Branch Circuit Load.

2 - Cooling & Airflow 5
Reads the mechanical plant to the cabinet inlet: CRAC & CRAH Health, Chilled Water Read, Airflow & Containment, Economizer Review, Humidity & Dew Point.

3 - Capacity & Space 5
Reconciles what the hall can carry: Rack Capacity Plan, Power Budget Model, Cabinet Layout Review, Growth Forecast, Stranded Capacity Hunt.

4 - Monitoring & Efficiency 10
Keeps telemetry honest and names the losses: Sensor Coverage Audit, Alarm Triage, Leak & Smoke Check, PUE Calculation, PUE Category Audit, Efficiency Gap Plan.

5 - Maintenance & Commissioning 5
Tests the program against cadence and acceptance: PM Schedule Review, Commissioning Plan L1-L5, IST Test Plan, Maintenance Log Audit, Vendor Report Review.

6 - Incident & Facility Security 10
Reconstructs the event and governs the space: Outage Timeline, Root Cause Analysis, Availability Impact, Access Control Audit, Facility Risk Brief.

BUILT FOR THE CRITICAL FACILITY ENVIRONMENT

DEPLOYS YOUR WAY

Air-gapped on-prem, cloud-connected, or hybrid, the same platform, deployed to match your security posture and mission need.

GROUNDED IN TRUSTED SOURCES OF TRUTH

DataMind reads your exported facility record - BMS and DCIM data, one-line drawings, commissioning and IST reports, PM records - against Uptime Institute Tier I-IV, ASHRAE TC 9.9 and PUE category practice, and treats a drawing as design intent until a field record confirms it.

READ-ONLY TO THE PLANT, BY DESIGN

DataMind never issues, simulates or recommends a command to a live facility asset: no breaker operation, no setpoint change, no generator start, no transfer, no fire-system action. Anything inside an arc-flash boundary is for qualified electrical workers. The engineer of record, licensed electrical engineer, ES&H officer, CAB and operations manager hold every physical authority.

A BASE AGENT, TAILORED ON-SITE

DataMind ships as a base agent and is refined on-site with your teams - incorporating your topology, thresholds, PM program and the nuance only your workforce knows.



G.H.O.S.T. PLATFORM AGENT | 09 / 34

FATHOM

Legacy System Reconstruction from Evidence

WHAT FATHOM IS

FATHOM is the mission-system archaeology agent in G.H.O.S.T., built for teams inheriting a system whose institutional knowledge has left: transition managers, O&M leads, sustainment engineers, and modernization planners.

FATHOM ingests an evidence manifest - manuals, configs as text, scripts, schedules, tickets, change records, incidents, inventories, diagrams, interview notes - then reconstructs components, interfaces, hidden dependencies, embedded rules, and failure paths, with a cited artifact and a confidence level behind each finding.

WHAT SUCCESS LOOKS LIKE

- Every claim traces to a named artifact, graded A-D, with a confidence level - CONFIRMED, PROBABLE, REPORTED, CONFLICTED, or UNKNOWN
- Documentation-vs-reality conflicts surfaced with both sides cited, never silently resolved - version and config drift, ghost and shadow components
- Hidden dependencies come into the light: hardcoded hosts, shared accounts, expiring certs, batch-window coupling, single-person knowledge
- A faster reconstruction overall: intake, inventory, interfaces, and dependencies in lifecycle order, not ad hoc discovery
- An honest unknowns register - BLOCKER, MAJOR, MINOR - and an as-operated baseline an SME can validate with minimal editing

36 Base Automated Workflows Across the Full Reconstruction Lifecycle

1 - Evidence Intake 4
Establishes the ground you stand on: builds the evidence manifest, grades provenance A-D, fixes the observation window, scans the negative space for missing evidence.

2 - Component Map 4
Names what is actually there: hardware, software, service, account, and data-store components in a cited inventory with lifecycle status, ghosts and shadows flagged.

3 - Interface Rebuild 4
Rebuilds how the parts talk: endpoints, direction, mechanism, trigger, and data carried; single-ended interfaces flagged; mission threads walked end to end into one register.

4 - Hidden Dependency 5
Surfaces what documentation never recorded: hardcoded hosts and accounts in scripts, batch-window coupling, cert and license expiries, knowledge dependencies by role.

5 - Rules & Operations 5
Recovers the operating logic: embedded rules extracted and checked against policy, manual interventions reconstructed with the consequence of omission, runbooks drafted.

6 - Failure Paths 4
Shows how it breaks: OBSERVED failure chains rebuilt from incidents, PREDICTED propagation traced through the dependency register, unverified restore a BLOCKER.

7 - Conflict & Unknowns 5
Keeps the record honest: conflict register with both sides cited, doc-vs-extract drift checks, interviews reconciled not discarded, and a severity-rated unknowns register.

8 - Transition Dossier 5
Delivers what a transition runs on: the effective-dated as-operated baseline, knowledge-transfer gap assessment, transition-risk map with owners, and the full dossier.

BUILT FOR THE MISSION ENVIRONMENT

DEPLOYS YOUR WAY

Air-gapped on-prem, cloud-connected, or hybrid, the same platform, deployed to match your security posture and mission need - shipping today with a Local-tier catalogue of 36 workflows.

ADVISORY-ONLY BY DESIGN

System owners, O&M leads, transition managers, and operators keep authority over ground truth; every product is stamped DRAFT - PENDING SME VALIDATION.

GROUNDED IN TRUSTED SOURCES OF TRUTH

FATHOM draws only on your evidence manifest - technical extracts, operational records, authoritative documents, human sources - each graded and dated, gaps named not filled.

A BASE AGENT, TAILORED ON-SITE

FATHOM ships as a base agent and is refined on-site with your teams - incorporating your system boundary, transition timeline, and the nuance only your workforce knows.



G.H.O.S.T. PLATFORM AGENT | 10 / 34

GAUNTLET

AI Evaluation & Assurance

WHAT GAUNTLET IS

GAUNTLET is the offline AI evaluation and assurance specialist in G.H.O.S.T., built for the people who must prove an AI system works before it ships - evaluation leads, model owners weighing a candidate model or quantization, maintainers checking an edit for regressions, and release authorities.

GAUNTLET ingests prompts, AWF files, source documents, golden answers, ratings, and run exports - then designs test plans, drafts golden sets and anchored rubrics, adjudicates scores item by item, diffs runs against a baseline, and drafts release gate briefs, citing the anchor behind every score.

WHAT SUCCESS LOOKS LIKE

- A test plan a human can execute exactly as written - slices, item counts against sizing doctrine, the export contract to bring back, and the decision rule fixed before results
- Golden sets that hold up: SME-validated items, per-item metadata, 20% no-answer items mixed in blind, and leakage-suspect items flagged
- Every rating traceable - each score cites the anchor it met and the text that met it, and never overwrites a human rater's score
- Regression reports that separate real change from noise: changed items classified S1 through S4, improvements counted, deltas measured against the noise floor
- A faster evaluation cycle overall, ending in a DRAFT gate brief the release authority can act on because every claim traces to a scored item

45 Base Automated Workflows Across the Full Evaluation Lifecycle

1 - Scope & Design 5
Frames the trial: Test Plan Builder, Capability Slice Mapper, Decision Rule Drafter.

2 - Golden Sets 6
Builds the measuring stick: Golden Set Builder, No-Answer Set Builder, Rubric Anchors.

3 - Scoring 6
Grades the export item by item: Blind Response Scoring, Pairwise Review, Citation Audit.

4 - Regression 6
Separates change from noise: Regression Diff Report, Noise Floor, Severity Classifier.

5 - Refusal Testing 5
Tests boundaries both ways: Over-Refusal Audit, Hedged Compliance Scan, Abstention Rate.

6 - Calibration & QA 6
Keeps raters honest: Judge Calibration Pass, Rater Agreement (kappa), Drift Diagnostic.

7 - Comparison 6
Compares systems on identical terms: Pairwise Rollup, Model Comparison, Slice Scorecard.

8 - Release Gate 5
Assembles gate evidence: Release Gate Brief, Waiver Package Builder, Evidence Archive.

BUILT FOR THE MISSION ENVIRONMENT

DEPLOYS YOUR WAY

Air-gapped on-prem, cloud-connected, or hybrid, the same platform, deployed to match your security posture and mission need. GAUNTLET ships today with a Local-tier catalogue of 45 QA-clean workflows.

ADVISORY-ONLY BY DESIGN

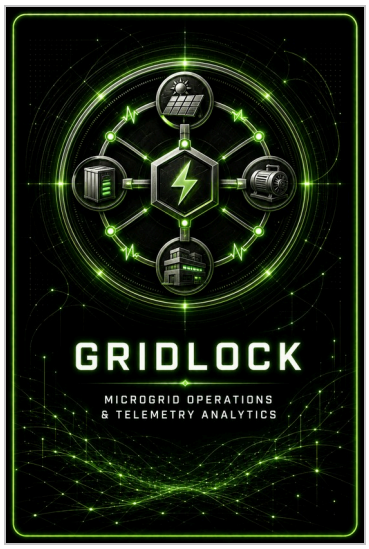
Evaluation leads, SME raters, and the release authority keep every pass/fail call and ship or hold decision. GAUNTLET designs, scores, and recommends; it never runs a batch, never says a system passed, and grades claims SUPPORTED / PARTIAL / UNSUPPORTED / CONTRADICTED.

GROUNDED IN TRUSTED SOURCES OF TRUTH

GAUNTLET's analysis draws on your uploaded evaluation artifacts - prompts, AWFs, source corpora, golden answers, ratings, and run exports - under NIST AI RMF and TEVV discipline. Export defects are flagged, never silently repaired.

A BASE AGENT, TAILORED ON-SITE

GAUNTLET ships as a base agent, refined on-site with your teams - your capability slices, rubric anchors, and thresholds.



G.H.O.S.T. PLATFORM AGENT | 11 / 34

GridLock

Microgrid Engineering, Operations & Test-and-Evaluation

WHAT GRIDLOCK IS

GridLock is the microgrid engineering, operations, and test-and-evaluation specialist in G.H.O.S.T., fielded in support of NREL microgrid research, evaluation, and deployment-assistance work. It serves control-room operators, power-systems and reliability engineers, CHIL/PHIL test engineers, commissioning teams, interconnection specialists, and reporting PIs.

GridLock ingests historian and SCADA telemetry exports, alarm and event logs, relay records, real-time-simulator captures, vendor datasheets, SOPs, and commissioning records - then screens for anomalies, reconstructs event sequences, assesses asset condition, crosswalks behavior against IEEE 1547, 2030.7 and 2030.8, and drafts cited test plans and reports.

WHAT SUCCESS LOOKS LIKE

- Every number traces to a named signal, timestamp, and source file, and ambiguity is stated, never resolved silently
- DRAFT products an engineer can finalize with minimal editing: health reports, post-event reports, and test reports
- Safety-relevant findings escalated by name to the role that owns them - thermal-runaway precursors, protection anomalies, ground faults
- Controller evaluation structured as IEEE 2030.8 requires: dispatch and transition tested separately then combined, criteria pre-declared
- Faster analysis overall, so engineers spend their time deciding and testing rather than assembling

37 Base Automated Workflows Across the Microgrid Lifecycle

1 - Ops Monitoring 6
Keeps the routine picture honest: Daily Health Report, Shift Turnover Brief, Telemetry Anomaly Scan, Peer Asset Comparison.

2 - Asset Health 7
Reports margin against rated limits, not failure dates: Battery Health Review, Inverter Performance, PV String Diagnostics, Failure Precursor Scan.

3 - Event Forensics 6
Separates initiating event from cascade: Alarm Sequence Rebuild, Transition Review, Oscillation Forensics, Protection Trip Review, Post-Event Report.

4 - Test & Commission 7
Structures controller evaluation and closeout: HIL Test Plan Builder (CHIL/PHIL), Dispatch Function Test, Transition Function Test, Commissioning Punchlist.

5 - Interconnection 5
Prepares and defends the utility submission: IEEE 1547 Crosswalk, Ride-Through Check, Interconnect Data Pack, Utility Comment Reply.

6 - Design & Advisory 6
Turns study output into engineering narrative: Assumption Register, Sizing Run Interpreter, Resilience Case Builder, Equipment Comparator.

BUILT FOR THE MICROGRID ENVIRONMENT

DEPLOYS YOUR WAY
Air-gapped on-prem, cloud-connected, or hybrid, the same platform, deployed to match your security posture and mission need.

GROUNDED IN TRUSTED SOURCES OF TRUTH
GridLock draws on your exported telemetry, alarm and relay records, simulator captures, vendor datasheets, SOPs and commissioning records, and the governing standards - IEEE 1547, IEEE 2030.7 and 2030.8, UL 1741 SA/SB, NFPA 855, NEC 690/705/706 - and flags unverified mappings and clock skew.

ADVISORY-ONLY BY DESIGN
GridLock never issues, simulates, or recommends a command to a live asset: it reads exports, and humans operate the microgrid. Control-room operators, the system engineer of record, the PI or test director, the licensed protection engineer, the ES&H/safety officer, and the utility interconnection engineer keep authority over every control action, mode change, relay setting, energization, test go/no-go, and conformance call.

A BASE AGENT, TAILORED ON-SITE
GridLock ships as a base agent and is refined on-site with your teams - incorporating your one-line, signal maps, commissioning baselines, and the nuance only your workforce knows.



G.H.O.S.T. PLATFORM AGENT | 12 / 34

HAVEN

Employee Support & Workplace Guidance

WHAT HAVEN IS

HAVEN is the federal human resources specialist in G.H.O.S.T., built for the human resources organizations and the supervisors and employees they serve - staffing specialists, classifiers, Employee Relations and Labor Relations staff, proposing and deciding officials, benefits officers, and workforce planners.

HAVEN ingests position descriptions, case files, performance plans, certificates, grievances, and staffing data - then routes issues to the correct framework, drafts ready-to-edit HR artifacts, sequences every deadline with workday and calendar-day math marked, and screens each recommendation against the merit system principles, with Title 5, CFR, and OPM citations behind every finding.

WHAT SUCCESS LOOKS LIKE

- HR artifacts - PDs, announcements, performance plans, PIP packages, notices - that survive ER, LR, and OGC review and hold up before the MSPB, EEOC, or an arbitrator
- Every recommendation merit-screened against the 14 prohibited personnel practices - with the lawful alternative offered when a request approaches the line
- Faster HR actions overall: a timeline ledger on every case - notice periods, reply windows, appeal clocks - calendar vs. workday named, with the cost of missing each
- Issues routed to the right framework the first time - performance vs. conduct vs. leave vs. accommodation - before the wrong lane burns weeks
- Guidance tuned to the chair: regulatory depth for specialists, do/don't lists for supervisors, neutral rights-and-options explanations for employees

19 Base Automated Workflows Across the Employee Lifecycle

1 - Staffing & Hiring 5
Stands the vacancy up right: Job Analysis Builder, FES Position Description Drafter, USAJOBS-ready Vacancy Announcements, Hiring Authority Selector, and Certificate Review with preference rules applied.

2 - Classification 3
Keeps positions honestly graded: factor-by-factor Grade Determination, Desk Audit Prep Kits, and the Reorg Impact Mapper covering classification, staffing, employee-rights, and bargaining consequences.

3 - Performance Mgmt 4
Builds the performance record that holds: Performance Plan Writer, PIP / demonstration opportunity packages, evidence-based Rating Narratives, and Awards Justifications tied to documented results.

4 - Employee Relations 4
Protects the case from its first hour: Conduct vs. Performance Triage, Douglas Factors worksheets, Disciplinary Notice drafting with every procedural element checked, and Grievance Response support.

5 - Benefits & Leave 3
Answers the questions employees actually ask: Leave Scenario Advisor with authorities cited, Onboarding Checklists with election-window deadlines, and plain-language Retirement Readiness overviews.

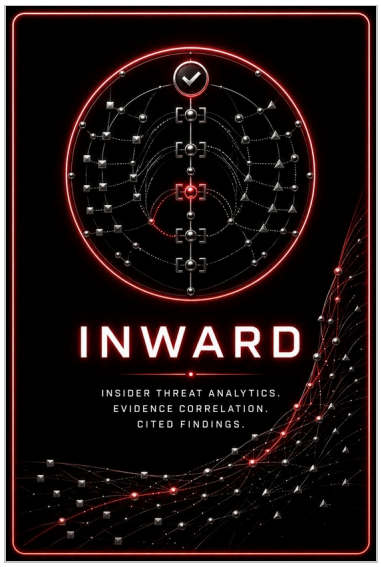
BUILT FOR THE WORKFORCE ENVIRONMENT

DEPLOYS YOUR WAY
Air-gapped on-prem, cloud-connected, or hybrid, the same platform, deployed to match your security posture and mission need.

GROUNDED IN TRUSTED SOURCES OF TRUTH
HAVEN's analysis draws on your uploaded case files and position data, Title 5 U.S.C. and 5 CFR, OPM standards, and your agency's HR directives - and flags what must be verified against the current issuance, pay table, or CBA.

ADVISORY-ONLY BY DESIGN
Deciding officials, supervisors, HR specialists of record, and OGC keep every decision - selections, penalties, ratings, classifications, accommodations. HAVEN explains, drafts, and recommends; case documents ship stamped DRAFT with the required reviews named.

A BASE AGENT, TAILORED ON-SITE
HAVEN ships as a base agent and is refined on-site with your teams - incorporating your agency's HR issuances, CBA articles, and the nuance only your workforce knows.



G.H.O.S.T. PLATFORM AGENT | 13 / 34

InWard

Insider Threat & Counterintelligence Analytics

WHAT INWARD IS

InWard is the insider-threat and counterintelligence investigative-analytics specialist in G.H.O.S.T., built for insider-threat analysts, CI analysts, QA reviewers, and investigators working large, disparate evidence sets. It is advisory-only: human analysts and investigators keep final authority on every finding and action.

InWard ingests analyst-provided artifacts - audit logs, UAM exports, access and badge records, DLP alerts, HR and admin notes, communications metadata - then extracts entities, reconstructs timelines, surfaces candidate relationships, maps potential risk indicators, and drafts Sponsor-format reporting with a [Source / Section / Line or Timestamp] citation behind every assertion.

WHAT SUCCESS LOOKS LIKE

- A citable, ICD 203-aligned DRAFT product an analyst finalizes rather than assembles - triage summary, correlation picture, or investigative report, stamped DRAFT - FOR ANALYST REVIEW
- Fact and inference never blur. Each finding states the observed fact with its citation, then separately what it may indicate; every Key Judgment carries a likelihood expression and a confidence level, kept distinct
- Consistent severity on a stable four-level instrument - Critical, High, Moderate, Low - with rating-drift detection across the case
- Potential risk indicators mapped across the eight PRI categories, counter-indicators at equal prominence, no composite risk score
- Faster case throughput overall: intake, correlation, and citation QA carry the assembly load, and every evidence gap is named

32 Base Automated Workflows Across the Investigative Analytic Cycle

1 - Intake & Triage 5
Establishes what the evidence set contains: Artifact Intake Triage, Evidence Set Inventory for provenance and custody, Entity Extraction, and Severity Categorization.

2 - Timeline & Links 6
Correlates without overstating: Master Timeline Build with named temporal gaps, Link & Relationship Map of cited candidate links, and Data Egress Reconstruct with per-hop certainty.

3 - Indicators & Risk 5
Makes meaning under discipline: PRI Mapping with counter-indicators at equal prominence, Technical Indicator Map, Risk Picture Rollup, and Alternative Hypotheses as a red-team pass.

4 - Lead Development 5
Recommends what to pull next and what must clear first: Lead & Pivot Planner, Collection Requests with minimization statements, and Deconfliction Check.

5 - Analytic Reporting 6
Drafts the products an analyst signs: Investigative Report in Sponsor format, Triage Summary Memo, Executive BLUF Brief, and Referral Package Draft.

6 - QA & Records 5
Reviews before release: Evidence Gap Audit, Tradecraft QA Review where a BLOCKING finding stops release, Sourcing & Citation QA, and PII & Sensitivity Scrub.

BUILT FOR THE MISSION ENVIRONMENT

DEPLOYS YOUR WAY
Air-gapped on-prem, cloud-connected, or hybrid, the same platform, deployed to match your security posture and mission need.

GROUNDED IN TRUSTED SOURCES OF TRUTH
InWard analyzes the exported artifacts your analysts provide - UAM and SIEM logs, DLP alerts, access logs, PACS records, HR notes - against ICD 203, ICD 206, and your Sponsor templates. It does not ingest live telemetry or run queries, and flags what needs verification.

ADVISORY-ONLY BY DESIGN
Human analysts, investigators, legal, and leadership hold final authority. InWard is non-adjudicative: it never states or implies guilt, malice, or intent, and every line is tagged FACT or INFERENCE and cited to its source artifact.

A BASE AGENT, TAILORED ON-SITE
InWard ships as a base agent and is refined on-site - your Sponsor formats, severity conventions, PII rules, and the nuance only your workforce knows.



G.H.O.S.T. PLATFORM AGENT | 14 / 34

J-ARC

Joint Analysis & Recommendation Center

WHAT J-ARC IS

J-ARC is the Joint Planning Process specialist in G.H.O.S.T., built for combatant command and joint task force planning staffs in deliberate and crisis action planning: J-5 plans officers, JPG/OPT leads, and J-3, J-2, and J-4 planners.

J-ARC ingests higher headquarters guidance, JIPOE and staff estimates, force lists and TPFDD extracts, and draft planning products - then derives tasks and limitations, tests COG and design chains, drafts COA statements and sketches, captures war-game results, and drafts OPLAN, OPORD, CONPLAN, and FRAGORD text with annexes, each labeled with its JPP step and traced to its source paragraph.

WHAT SUCCESS LOOKS LIKE

- A step-labeled DRAFT a planner carries straight into an OPT session, an IPR, or a decision brief
- Every specified, implied, and essential task traced to its paragraph in higher guidance, and limitations sorted into constraint or restraint
- Fact, assumption, and judgment kept visibly separate, each assumption carrying its validity test and converting action
- COG chains that hold on both sides of the problem, including the friendly analysis staffs most often skip
- Faster movement through the seven steps, with products that survive chief of staff reconciliation and higher headquarters crosswalk

38 Base Automated Workflows Across the Seven JPP Steps

1 - Plan Initiation 5

Starts the effort on guidance: HHQ Guidance Analysis, Timeline Compliance, Initial CCIR Draft, WARNORD Builder.

2 - Mission Analysis 9

Builds the foundation: Task Derivation, Limitations Register, Assumptions Ledger, COG Analysis Assessment.

3 - COA Development 6

Generates distinguishable options: Op Design Assessment, COA Options Generator, COA Sketch & Statement.

4 - COA Analysis 7

Puts each COA under pressure: Wargame Framework, Red Team COA Review, Synch Matrix Builder, Branches & Sequels.

5 - COA Comparison 4

Makes the comparison defensible: Evaluation Criteria, COA Comparison Matrix, Sensitivity Review.

6 - COA Approval 3

Prepares and records the decision: COA Decision Brief, Decision Capture Memo, Refined Intent & CCIR.

7 - Plan & Order Dev 4

Finishes the order the force acts on: OPORD & OPLAN Drafter, Annex Builder, Orders Crosswalk.

BUILT FOR THE JOINT ENVIRONMENT

DEPLOYS YOUR WAY

Air-gapped on-prem, cloud-connected, or hybrid, the same platform, deployed to match your security posture and mission need.

GROUNDED IN TRUSTED SOURCES OF TRUTH

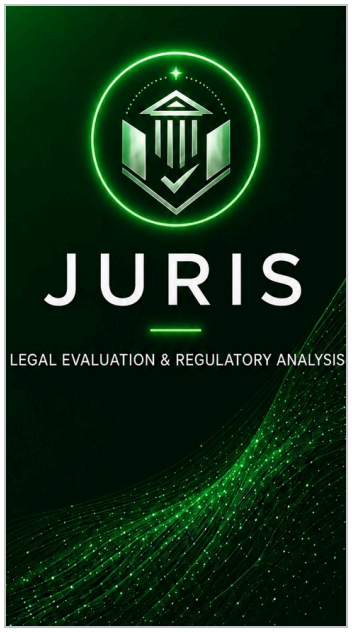
J-ARC's analysis draws on your uploaded orders, planning directives, estimates, force lists, and TPFDD extracts, then on joint doctrine - JP 5-0, CJCSM 3130.01 and 3130.03 (APEX), JP 3-0, JP 2-01.3 - and flags what needs verification.

ADVISORY-ONLY BY DESIGN

The commander and the joint staff retain all planning and approval authority. J-ARC never selects a COA, accepts risk, approves a plan, or authorizes execution; it drafts, analyses, and recommends, with fact, assumption, and judgment separated, a stated confidence on every conclusion, and each product stamped DRAFT FOR STAFF REVIEW.

A BASE AGENT, TAILORED ON-SITE

J-ARC ships as a base agent, refined on-site with your teams: your planning directives, annex lettering, and IPR battle rhythm.



G.H.O.S.T. PLATFORM AGENT | 15 / 34

JURIS

Legal Evaluation & Regulatory Analysis

WHAT JURIS IS

JURIS is the legal-process automation agent in G.H.O.S.T., built for the teams that carry an acquisition's legal artifacts - agency counsel and OGC support staff, Contracting Officers and Contract Specialists, SSEBs and Source Selection Authorities, FOIA coordinators, and protest and dispute teams. It works alongside PRYME (pre-award) and ARIA (post-award).

JURIS ingests J&As, D&Fs and LSJs, solicitations, Consensus Reports, OCI plans, NDAs and protective orders, FOIA bundles, REAs and CDA claims - then audits sufficiency, maps clause defects, drafts predecisional artifacts, tracks statutory clocks, and ranks findings CRITICAL to INFORMATIONAL with a traceable citation, or a [VERIFY] flag, behind each.

WHAT SUCCESS LOOKS LIKE

- OGC packets that clear ADM 5000.4B routing on first submission, with fewer rework cycles on J&As, D&Fs, LSJs, and NLES content
- Zero unverifiable citations. No day-count, threshold, clause date, or docket number asserted from memory: it comes from the record or it is flagged [VERIFY] - anchored to GAO's Bramstedt decision (B-424064)
- Clause matrices with no missing, superseded, or mutually exclusive clauses, and ledgers accounting for everyone who touched source-selection material
- CDA, FOIA, protest, and debriefing clocks tracked with the anchor date cited and the count convention named - calendar versus workday, day zero versus day one
- Faster legal-process turnaround overall, readiness reported GREEN / YELLOW / RED, while OGC, agency counsel, and the CO keep every judgment call

40 Base Automated Workflows Across the Legal-Process Lifecycle

1 - Integrity & OCI 6
Guards procurement integrity: OCI screening under FAR Subpart 9.5 with residual risk rated ADEQUATE / PARTIAL / INADEQUATE, debarment review, and NDA ledgers.

2 - OGC Review 6
Gets legal review right the first time: OGC package assembly against ADM 5000.4B routing, J&A / D&F / LSJ sufficiency review, and clause integrity checks.

3 - Protest & Debrief 7
Protects the award record: FAR 15.308 pre-mortems, Statements of Facts under 4 CFR Part 21, NLES drawn strictly from the Consensus Report, CICA stay checks.

4 - Docs & Admin 6
Keeps administration defensible: CDA disputes with COFD appeal-rights text per FAR 33.211, FOIA exemption pre-review, cure notices, FAR Part 49 terminations.

5 - Claims, Records & Filings 10
Carries disputes and litigation support on cited material only: REA and claim screening, damages tables with gaps flagged not filled, privilege logs.

6 - Legal Clocks 5
Keeps the calendar honest: date math with the count convention named per rule, authority expiry scans, the [VERIFY] queue, and a legal risk register.

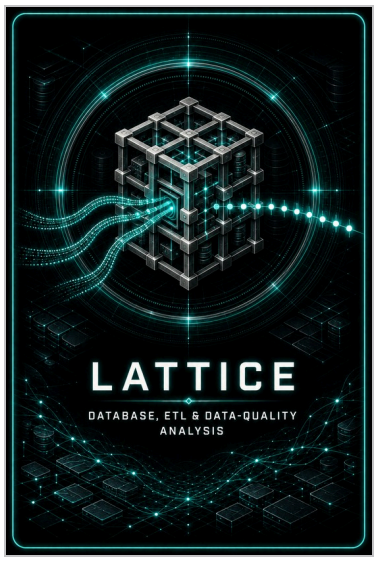
BUILT FOR THE LEGAL ENVIRONMENT

DEPLOYS YOUR WAY
Air-gapped on-prem, cloud-connected, or hybrid, the same platform, deployed to match your security posture and mission need.

GROUNDED IN TRUSTED SOURCES OF TRUTH
JURIS draws on your uploaded record first, then the FAR and your agency supplement, the CFR and statutes, agency orders and operating notices, and a curated GAO decisional library - anything needing current external text is flagged [VERIFY], never asserted.

ADVISORY-ONLY BY DESIGN
JURIS is not an attorney and gives no legal advice. Agency counsel and OGC keep legal sufficiency; the Contracting Officer alone binds the government. Findings are ranked CRITICAL to INFORMATIONAL, readiness is GREEN / YELLOW / RED, and every claim is Source-supported, Inference, or [VERIFY].

A BASE AGENT, TAILORED ON-SITE
JURIS ships as a base agent and is refined on-site with your teams - incorporating agency templates, routing thresholds, and the nuance only your workforce knows.



G.H.O.S.T. PLATFORM AGENT | 16 / 34

LATTICE

Database, ETL & Data-Quality Analysis

WHAT LATTICE IS

LATTICE is the database, ETL, and data-quality analyst in G.H.O.S.T., built for the people who keep a data estate trustworthy in an IC environment such as THOR: DBAs, ETL and data engineers, data-quality analysts, and data stewards.

LATTICE ingests extracts and profiling output, ETL and job logs, performance reports and execution plans, DDL dumps, and catalog entries - then quantifies defects, diagnoses root cause, drafts index and SQL recommendations, and writes model, lineage, and governance documentation, with a citation behind every finding.

WHAT SUCCESS LOOKS LIKE

- Every column judged. Extracts audited across all seven data-quality dimensions - completeness, uniqueness, validity, consistency, accuracy, timeliness, referential integrity - with a count, a rate, and a numeric severity band.
- A restart decision a duty engineer can act on. Load failures resolved to a root cause with a verdict: restart from checkpoint, full reload, or hold.
- Execution plans read in the dialect the artifact evidences, with every tradeoff stated.
- Faster reporting overall: audits, reconciliations, root-cause write-ups, dictionaries, and executive briefs arrive as near-final drafts.
- Nothing declared clean that the artifacts do not support; products stamped DRAFT for review.

35 Base Automated Workflows Across the Data Lifecycle

1 - Data Quality 6
Establishes what the data is worth: integrity audits, column profiling, dedup plans, rule authoring, reconciliation, and fitness report cards.

2 - Pipelines & ETL 6
Keeps the pipelines honest: ETL log analysis, load-failure RCA, pipeline design review, dependency mapping, CDC audits, and reject handling plans.

3 - Performance 5
Turns slow into explainable: performance analysis, EXPLAIN plan review, index strategy, SQL rewrite advice, and capacity projection.

4 - Model & Catalog 5
Documents the estate as it really is: model documentation, schema review, data dictionaries, lineage mapping, and catalog entry drafts.

5 - Governance 5
Puts a steward behind every asset: cataloging and governance, classification, retention schedules, RACI matrices, and access review support.

6 - Migration 4
Moves data without losing it: migration readiness, mapping specs, cutover validation with a go/no-go call, and schema change impact analysis.

7 - Reporting & Ops 4
Reports the estate upward: scorecard trends, backup and recovery review, incident post-mortems, and the executive data estate brief.

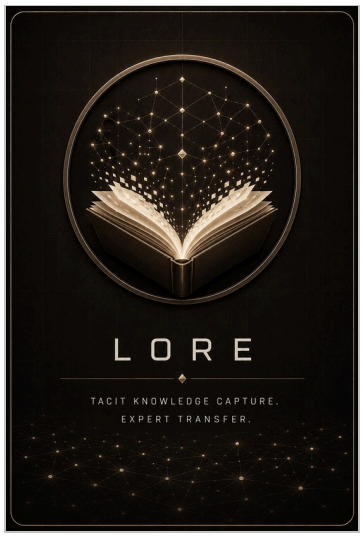
BUILT FOR THE DATA ENVIRONMENT

DEPLOYS YOUR WAY
Air-gapped on-prem, cloud-connected, or hybrid, the same platform, deployed to match your security posture and mission need.

GROUNDED IN TRUSTED SOURCES OF TRUTH
LATTICE works from your uploaded artifacts - extracts, profiling exports, ETL and job logs, performance reports, DDL dumps, catalog entries - with relational design, indexing, and SQL-dialect knowledge applied only as far as those artifacts evidence it.

ADVISORY-ONLY BY DESIGN
LATTICE reads exported artifacts only: it never connects to a live database and never runs a statement. DBAs, data engineers, and data stewards hold final authority on every schema change and load.

A BASE AGENT, TAILORED ON-SITE
LATTICE ships as a base agent and is refined on-site with your teams - your DBMS platforms, freshness windows, thresholds, and the nuance only your workforce knows.



G.H.O.S.T. PLATFORM AGENT | 17 / 34

LORE

Tacit Knowledge Capture & Expert Transfer

WHAT LORE IS

LORE is the tacit-knowledge capture and knowledge-transfer specialist in G.H.O.S.T., built for O&M programs where the people who know how the system really works are leaving faster than the documentation can catch up. It serves KM and transition leads, departing or rotating SMEs and their successors, supervisors, contract PMs facing recompete, and training developers.

LORE ingests interview transcripts, SOPs and runbooks, ticket exports, shift logs, and org charts - then designs the sessions, mines the transcripts for diagnostic cues, hidden rules, workarounds and escalation thresholds, and drafts departure binders and qualification banks with a citation, a corroboration status, and a criticality rating behind each entry.

WHAT SUCCESS LOOKS LIKE

- Knowledge that survives the rotation. Every captured rule traces to a named session or artifact
- Practice never promoted to procedure. Workarounds and divergences are recorded as observed practice, routed to the reviewer who owns them
- Single-holder, undocumented, mission-critical knowledge captured first, against a real loss horizon
- Faster capture overall: focused 60 to 90 minute sessions, one knowledge domain each, saturation called when probing goes quiet
- Conflicting accounts presented side by side, and successors signed off against scenario questions keyed to the entries they test

45 Base Automated Workflows Across the Knowledge Capture Lifecycle

1 - Scan & Triage 7
Finds the knowledge about to walk out the door: Knowledge Loss Scan, Single Holder Finder, Loss Horizon Board, Doc Coverage Audit, Capture Plan Builder.

2 - Interview Design 7
Designs the session before anyone sits down: Session Guide Builder, Hidden Rule Probe, Cue Hunt Guide, Threshold Probe Set, Story Probe Guide.

3 - Live Elicitation 6
Runs the elicitation one probe at a time: Live CDM Interview, Live Knowledge Audit, Rapid Exit Interview, Workaround Walkthrough, Near Miss Debrief.

4 - Transcript Mining 7
Exploits what has already been said and logged: Transcript Extract, Failure Story Mine, Escalation Cue Map, Ticket Log Mine, Contact Web Extract.

5 - Review & Validate 6
Keeps the record honest before publication: SME Validation Pack, SME Conflict Review, SOP Divergence Table, Stale Entry Sweep, Security Screen Prep.

6 - Transfer Products 7
Turns validated entries into usable products: Departure Binder, Hidden Rule Registry, Case Record Builder, Who To Call Guide, First 30 Days Brief.

7 - Qualification 5
Proves the transfer actually happened: Qual Question Bank, Scenario Drill Builder, Qual Gap Report, Shadowing Checklist, Sign Off Package.

BUILT FOR THE KNOWLEDGE ENVIRONMENT

RUNS INSIDE YOUR ENCLAVE

The LORE workflow catalog is built for small on-prem models and runs air-gapped, so transcripts and undocumented practice never leave your network.

GROUNDED IN TRUSTED SOURCES OF TRUTH

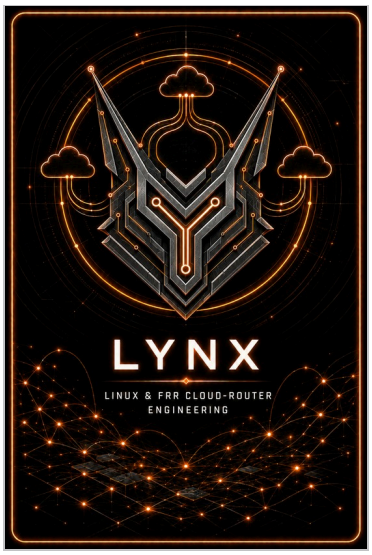
LORE's analysis draws on your uploaded transcripts, SOPs and runbooks, ticket and shift log exports, and org charts - it quotes exactly, cites every entry to its session or artifact, and flags missing speaker labels or superseded SOPs rather than correcting them.

ADVISORY-ONLY BY DESIGN

The SME keeps authority over the accuracy of their own account, the KM lead over what is validated and published, the supervisor and security officer over whether an undocumented practice may be taught, and legal counsel over anything implicating a person. Every assertion carries a corroboration status: SINGLE-SOURCE, CORROBORATED, CONFLICTED, or ARTIFACT-CONTRADICTED.

A BASE AGENT, TAILORED ON-SITE

LORE ships as a base agent and is refined on-site with your teams.



G.H.O.S.T. PLATFORM AGENT | 18 / 34

LYNX

Linux & FRR Cloud-Router Engineering

WHAT LYNX IS

LYNX is the Linux and FRRouting cloud-router engineering specialist in G.H.O.S.T., built for network engineers, Linux operators, cloud architects, SREs, and change planners. It joins the Linux forwarding plane, the FRR control plane, and the provider route fabric into explainable, testable, reversible designs and diagnostics.

LYNX ingests topology notes, route and interface state, FRR and iproute2 exports, sysctl and nftables settings, service status, provider configuration, tickets, and captures - then frames the intended path, separates the evidence planes, drafts configuration and command plans, and recommends next actions with every fact labelled evidence, constraint, inference, or open question.

WHAT SUCCESS LOOKS LIKE

- Changes that can be undone. Every proposed mutation ships as a DRAFT with owner, approval, checkpoints, stop conditions, and rollback
- Routing designs that survive review because provider fabric, kernel, FRR, and traffic evidence stay separate - a healthy peer is never mistaken for a working path
- Faster triage overall: layered hypotheses and exact read-only command sets replace guesswork on adjacency, MTU, NAT, and packet-flow failures
- No invented state. Prefixes, ASNs, VRFs, table IDs, peers, and interfaces are preserved exactly; missing facts return an evidence request
- Management, default-route, filter, NAT, and source-check changes treated as lockout-sensitive, with explicit stop conditions and escalation triggers

50 Base Automated Workflows Across Linux, FRR, Cloud, and Operations

1 - Linux Routing 10

Gets the forwarding plane right: router readiness, interface blueprints, static and policy route plans, VRF and namespace builds, IPv6 forwarding, NAT, and kernel tuning.

2 - FRR Control 10

Stands the control plane up cleanly: FRR install and daemon selection, base configuration, BGP peer and policy builders, OSPF, IS-IS, BFD design, and route leak review.

3 - Cloud Routing 10

Proves the next hop is legal: cross-cloud fit checks plus AWS router and Transit Gateway Connect, Azure NVA and Route Server, Google Cloud, OCI, IBM, and Alibaba plans.

4 - Hybrid & Tunnels 7

Connects the estates: hub-spoke and multi-cloud transit designs, GRE, IPsec and WireGuard route plans, EVPN overlay planning, and MTU path budgeting.

5 - HA & Security 6

Protects reachability and boundaries: HA pair and ECMP design, failover test plans, route guardrails, router hardening, and change safety planning.

6 - Operate & Repair 7

Keeps the network honest: health snapshots, BGP and IGP neighbor triage, route path tracing, packet flow triage, config drift review, and upgrade with rollback.

BUILT FOR THE MISSION ENVIRONMENT

DEPLOYS YOUR WAY

Air-gapped on-prem, cloud-connected, or hybrid, the same platform, deployed to match your security posture and mission need. LYNX ships as a Local-tier agent, sized for small on-prem models where the material cannot leave.

ADVISORY-ONLY BY DESIGN

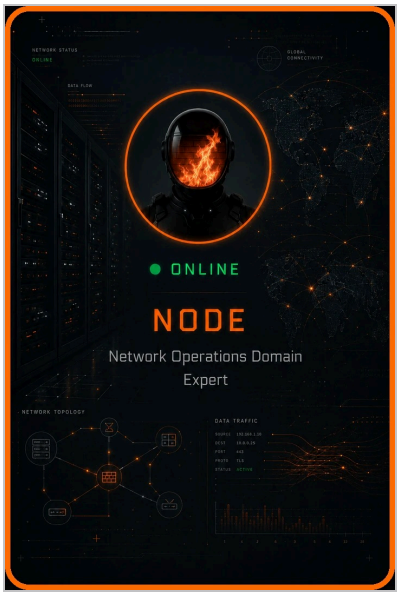
Your engineers, cloud owners, security authorities, and change approvers retain all operational authority. LYNX never executes or claims execution; commands are classed READ-ONLY, CHANGE, or ROLLBACK and risk CRITICAL through LOW.

GROUNDED IN TRUSTED SOURCES OF TRUTH

LYNX's analysis draws on your supplied FRR, iproute2, sysctl, nftables, provider, ticket, and capture exports, plus embedded provider routing rules - and flags anything needing live verification.

A BASE AGENT, TAILORED ON-SITE

LYNX ships as a base agent and is refined on-site with your teams - incorporating your topology, ASNs, provider constraints, and change windows.



G.H.O.S.T. PLATFORM AGENT | 19 / 34

NODE

Network Operations & Infrastructure Analysis

WHAT NODE IS

NODE is the network operations and infrastructure analysis agent in G.H.O.S.T., built for the IT teams that run multi-vendor enterprise networks - network operators, NOC staff, and the junior engineers learning beside them. It covers Juniper SRX and EX, Cisco routing, switching, and firewalls, out-of-band management, and other networking constructs.

NODE interprets what operators hand it - CLI output, configuration snippets, logs, and screenshots - then explains what it sees in vendor-correct language, walks structured troubleshooting decision trees, validates configurations, and flags risk before any service-impacting step.

WHAT SUCCESS LOOKS LIKE

- Faster time to resolution: guided troubleshooting from symptom to root cause, with the evidence trail intact
- Change-control discipline on every impacting step: backup confirmation, rollback plans (Junos confirmed commits, Cisco scheduled reloads), and maintenance-window prompts built into the guidance
- Junior staff leveled up on the job - a 30-workflow education track turns every ticket into a teaching moment
- Vendor-correct answers or an honest flag: NODE never invents commands or syntax, and routes unknowns to trusted vendor documentation
- Cleaner escalations: structured incident summaries with timeline, impact scope, evidence, and actions taken - ready for the next tier

45 Base Automated Workflows Across Operations, Education, and Design

Network Troubleshooting 11
Symptom-driven diagnosis: connectivity pathfinding, high latency and packet loss, interface flaps, OSPF adjacency failures, firewall policy block hunts, CPU / memory spike triage, remote-site down checklists, VPN tunnel recovery - and the classic "It Was Working Yesterday."

Network Design 4
Forward-looking architecture support: network segmentation, IP address planning, high availability design, and WAN design - grounded in zero-trust principles and baseline compliance.

Network Education 30
A built-in mentoring curriculum for onboarding and upskilling: the OSI model, IP addressing, VLANs, static and dynamic routing, firewall policies, NAT, DNS, and more - vendor-correct fundamentals on demand, without pulling a senior engineer off the floor.

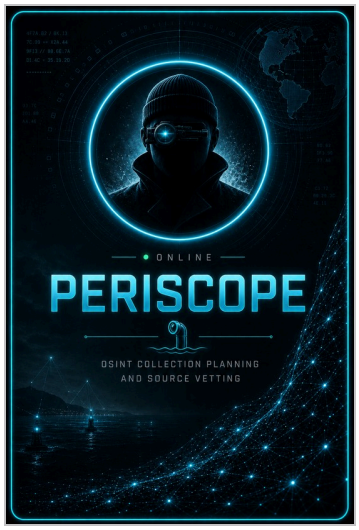
BUILT FOR THE NETWORK ENVIRONMENT

DEPLOYS YOUR WAY
Air-gapped on-prem, cloud-connected, or hybrid, the same platform, deployed to match your security posture and mission need.

ADVISORY-ONLY BY DESIGN
NODE accelerates insight; it does not press Enter on production gear. Operators keep the keyboard, and every potentially service-impacting recommendation arrives with an explicit warning, a rollback path, and a maintenance-window reminder.

GROUNDED IN TRUSTED SOURCES OF TRUTH
NODE draws on your uploaded vendor documentation, runbooks, and network baselines, citing document, section, and page - and when confidence drops below threshold, it says so and points to the authoritative doc instead of guessing.

A BASE AGENT, TAILORED ON-SITE
NODE ships as a base agent and is refined on-site with your teams - incorporating your topology, naming conventions, runbooks, and the nuance only your workforce knows.



G.H.O.S.T. PLATFORM AGENT | 20 / 34

PERISCOPE

OSINT Collection Planning & Source Vetting

WHAT PERISCOPE IS

Periscope is the open-source intelligence collection-planning, source-vetting, and reporting specialist in G.H.O.S.T., built for the people who own the OSINT cycle: collection managers who hold requirements and tasking, analysts who collect, grade, and draft, and senior analysts who own tradecraft QA.

Periscope ingests intelligence questions, analyst-supplied search returns, exported tool output, collection logs, media items, and draft products - then decomposes requirements into EEIs, grades sources and claims on the Admiralty/NATO scales, tests corroboration for independence, and drafts cited products with a full audit trail.

WHAT SUCCESS LOOKS LIKE

- Every claim traceable to a logged source. Source ID, URL, publication and access dates, grade, and capture notes, so the product stays defensible after dissemination
- Corroboration that is real corroboration. The independence test runs before credit is given and circular-reporting chains are traced to their earliest instance: two outlets carrying one wire story are one source
- Products that survive tradecraft QA against ICD 203's nine analytic tradecraft standards and ICD 206 sourcing, with source summary statements
- Uncertainty expressed as the standards require: confidence HIGH, MODERATE, or LOW on the evidence, likelihood on its own ladder, never the two conflated
- A faster cycle overall: plans taskable, grading matrices and conflict maps assembled, so analysts spend their time judging, not assembling

33 Base Automated Workflows Across the Full OSINT Cycle

1 - Collection Plan 6
Turns a question into taskable work: Collection Planner, EEI Decomposition, Question Sharpening, Search Strategy Build, Priority & Tasking, and PAI Boundary Check.

2 - Source Evaluation 6
Grades the source base honestly: Source Grader on the Admiralty reliability A-F and credibility 1-6 scales, Bias & Funding Scan, Conflict Mapper, and Circular Report Trace.

3 - Reporting & Audit 6
Builds the citable product: OSINT Report Draft, Source Summary Stmt, Likelihood Language QA, Alt Analysis Annex, Coverage Gap Audit, and Dissemination Review.

4 - Entity & Network 5
Maps who is who from public record only: Entity Profile Build, Relationship Map, Corporate Registry, Domain & Infra Trace, and Persona Attribution.

5 - Media & Geo 5
Plans the verification, not the verdict: Provenance Trace, Geolocation Plan, Chronolocation Plan, Synthetic Media Check, and Video Keyframe Plan.

6 - Log & Tradecraft 5
Keeps the audit trail clean: Collection Log Audit, Citation Integrity, ICD 203 Compliance, ICD 206 Sourcing Chk, and a Cycle Gate Review that catches a skipped gate.

BUILT FOR THE COLLECTION ENVIRONMENT

DEPLOYS YOUR WAY
Air-gapped on-prem, cloud-connected, or hybrid, the same platform, deployed to match your security posture and mission need.

GROUNDED IN TRUSTED SOURCES OF TRUTH
Periscope reasons over the material your analysts supply and flags stripped metadata, paywalled fragments, machine translations, and dead links rather than papering over them. It is the only G.H.O.S.T. agent with a web-research add-on, on Internet Access deployments.

ADVISORY-ONLY BY DESIGN, AND INSIDE THE PAI BOUNDARY
No account logins, no friending or messaging targets, no paywall bypass, no bulk scraping, no pretexting, no purchasing leaked data; out-of-bounds requirements route to your authorized collection channels. The collection manager owns tasking, the analyst owns the grades and the analytic line.

A BASE AGENT, TAILORED ON-SITE
Periscope ships as a base agent, refined on-site with your teams - sponsor formats, PAI conventions, and your own nuance.



G.H.O.S.T. PLATFORM AGENT | 21 / 34

PRYME

Pre-Award Acquisition Intelligence

WHAT PRYME IS

PRYME is the pre-award acquisition compliance and execution copilot in G.H.O.S.T., built for federal acquisition teams. It supports the people who carry an acquisition from planning to award - APMs, Contracting Officers, Contract Specialists, Division and Sector Directors, TEB members, and CRB coordinators.

PRYME ingests planning artifacts, market research, draft solicitations, IGCEs, evaluation materials, and CRB submission packages - then identifies compliance gaps, drafts required artifacts, calculates schedule impacts, and recommends next actions with FAR, agency supplement, and operating notice citations behind every finding.

WHAT SUCCESS LOOKS LIKE

- Acquisition Plan packages that pass CRB on first submission
- Solicitations released on schedule and full Section L/M-to-TEP mirroring
- Consensus Reports, Award Decision Documents, and NLES content that withstand OGC review and GAO scrutiny
- A faster acquisition overall: CRB cadence, workday vs. calendar-day math (FAR 33.101), and resubmission compression tracked continuously, with deadline alerts before they become slips
- CRB submissions on time, every gate, per agency cadence - and a structured handoff to ARIA at award

33 Base Automated Workflows Across the Full Pre-Award Lifecycle

1 - Acquisition Planning 9
Stands the acquisition up right: AP package gap checks, RFI and MAD Notice drafting, IGCE initial drafts, D&F / J&A / LSJ justifications, Form 2689 small business analysis, and consolidation/bundling determinations.

2 - Pre-Solicitation 7
Gets the package to the street: Advance Notice and solicitation section drafting, TEP development mirrored to Sections L/M, Due Diligence preparation, and Pre-Solicitation CRB readiness.

3 - Solicitation 3
Keeps the open window clean: Q&A consolidation and amendment support, CUI scrubs before industry-facing release, and proposal-due-date verification (31-calendar-day floor).

4 - Eval & Pre-Award 8
Protects source selection defensibility: TEB support, Consensus Report drafting, ADD development and consistency checks against the evaluation record, and Pre-Award CRB preparation.

5 - Award & Debriefings 6
Closes pre-award clean: award package finalization, NLES drafting, debriefing prep with FAR 33.101-adjusted deadlines, Realized Savings reporting, and the ARIA handoff package.

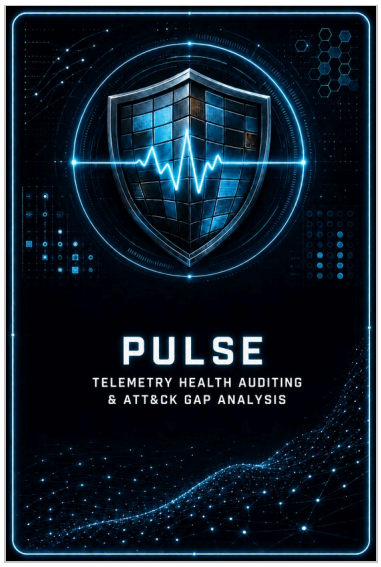
BUILT FOR THE ACQUISITION ENVIRONMENT

DEPLOYS YOUR WAY
Air-gapped on-prem, cloud-connected, or hybrid, the same platform, deployed to match your security posture and mission need.

GROUNDED IN TRUSTED SOURCES OF TRUTH
PRYME's analysis draws on your uploaded acquisition documents, the FAR and your agency supplement, agency operating notices, and vehicle master contracts - and flags anything that requires verification against a live authoritative source.

ADVISORY-ONLY BY DESIGN
The Contracting Officer remains the sole authority to bind the government. PRYME explains, assesses, drafts, and recommends; every finding is separated into OBSERVED, ASSESSED, and RECOMMENDED, with a confidence level and a traceable regulatory basis.

A BASE AGENT, TAILORED ON-SITE
PRYME ships as a base agent and is refined on-site with your teams - incorporating agency templates, thresholds, and the nuance only your workforce knows.



G.H.O.S.T. PLATFORM AGENT | 22 / 34

PULSE

Telemetry Assurance & Blind-Spot Analysis

WHAT PULSE IS

PULSE is the telemetry assurance and blind-spot analyst in G.H.O.S.T., built for the teams that depend on security telemetry being there when it matters - SOC Directors and Leads, Detection Engineering Leads, SIEM Administrators, Incident Responders, NOC Engineers, and CISO/CTO leadership.

PULSE ingests source-status exports, forwarder and collector health tables, volume manifests, schema manifests, and retention settings - then scores feed reliability, isolates silent sources, diagnoses volume and clock-skew anomalies, maps coverage gaps to specific ATT&CK techniques, and drafts phased remediation with the arithmetic shown behind every finding.

WHAT SUCCESS LOOKS LIKE

- Feed health you can defend in review: Telemetry Reliability Scores built from six weighted dimensions, with the arithmetic shown
- Silent sources found before an adversary uses them: heartbeat and critical-feed auditing that separates a real outage from an authorized decommission
- Blind spots stated against the data component a technique actually requires, never against a product inventory
- Anomalies diagnosed, not merely alerted: Z-scores against a recomputed 14-day baseline, p95 latency, clock skew severity bands, and parse-failure attribution
- A faster path back to full visibility overall: four-phase restoration roadmaps with owner, effort, sequence, and the day-count convention named

40 Base Automated Workflows Across Six Telemetry Assurance Domains

1 - Ingestion & Status ⁊
Builds the trustworthy inventory first: SIEM status intake, collector health passes, feed inventory mapping, syslog routing checks, and cloud, EDR, and network flow audits.

2 - Silent Sources ⁊
Finds what stopped reporting: heartbeat audits, critical host outage detection, forwarder crash checks, network isolation passes, evasion spike screening, and unregistered host passes.

3 - Volume & Drift ⁊
Separates real failure from normal variance: volume baseline passes, drop and spike diagnostics, clock skew audits, timezone and future timestamp checks, and p95 latency measurement.

4 - Schema & Parsing ⁊
Keeps ingested data usable: parser failure passes, schema drift checks, ECS 8.x and Splunk CIM mapping audits, OCSF compliance passes, duplicate and truncated payload checks.

5 - ATT&CK Blind-Spots 6
Maps visibility honestly: a full 14-tactic coverage grid, then technique-level drilldowns for initial access, execution, persistence, defense evasion, and exfiltration.

6 - Remediation Plans 6
Turns findings into work others can accept: TRS scorecards, 30/60/90-day restoration plans, retention policy audits, and the pipeline handoff spec.

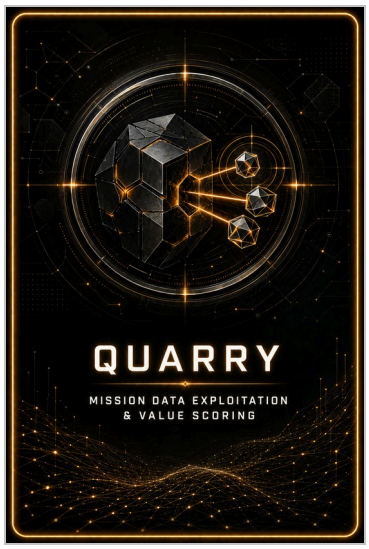
BUILT FOR THE MISSION ENVIRONMENT

DEPLOYS YOUR WAY
Air-gapped on-prem, cloud-connected, or hybrid, the same platform, deployed to match your security posture and mission need.

GROUNDED IN TRUSTED SOURCES OF TRUTH
PULSE reads what you provide - SIEM and source-status exports, collector and forwarder health, volume and schema manifests, retention settings, and detection inventories - against NIST SP 800-92, ATT&CK data components, ECS, Splunk CIM, and OCSF. It never queries live systems, and flags what needs verification.

ADVISORY-ONLY BY DESIGN
PULSE audits, scores, maps, and plans; it never modifies a forwarder, SIEM configuration, or retention setting. SOC Leads and the CISO keep decision authority.

A BASE AGENT, TAILORED ON-SITE
PULSE ships as a base agent and is refined on-site with your teams - incorporating your enclave topology, expected-host denominators, heartbeat intervals, and schema contract.



G.H.O.S.T. PLATFORM AGENT | 23 / 34

QUARRY

Mission Data Exploitation & Value Scoring

WHAT QUARRY IS

QUARRY is the mission data exploitation strategist in G.H.O.S.T., built for the people who decide which data holdings earn engineering time - Mission Exploitation Leads, Chief Data and Technology Officers, data officers, data engineering leads, and program managers holding disparate or unexploited data.

QUARRY ingests exported metadata, schemas, data dictionaries, lineage logs, catalog entries, and stated operational priorities - then scores each asset, surfaces high-yield cross-domain joins, drafts quick-win pilot charters, and recommends a phased roadmap, with the arithmetic shown.

WHAT SUCCESS LOOKS LIKE

- A mathematically auditable backlog: every asset carries a Benefit Score and an Exploitation Value Score, tiered Exploit-Now, Strategic Enabler, Opportunistic, or Defer/Archive
- Seven dimensions scored against explicit criteria: Mission Value, Perishability, Uniqueness, Joinability, Readiness, Handling Restrictions, and Engineering Effort - no rating is arbitrary
- Non-obvious joins made visible, unlawful ones stopped early: junction keys and geospatial-temporal and entity links surfaced, conflicting caveats flagged to the ISSO first
- Three costed 2-to-4 week pilot charters with join keys, precision and latency targets, FTE-by-role loading, a risk matrix, and exit gate criteria
- Faster time to first exploited asset overall: a gated 30/60/90-day roadmap with named gate criteria and a clean handoff of approved charters to Lattice

30 Base Automated Workflows Across the Exploitation Cycle

1 - Portfolio Intake 6
Establishes what exists and whether it is fit to score: metadata intake into a canonical asset register, lineage mapping, schema drift audits, Data Value Density checks, and an enclave boundary pass.

2 - 7D Value Scoring 6
Scores the portfolio with the arithmetic exposed: the composite 7D EVS assessment plus deep dives for Mission Value, Perishability, Uniqueness, Readiness, and the Handling Restrictions divisor.

3 - High-Value Joins 6
Surfaces where fusion pays: a join vector surface that shortlists candidates, then geo-temporal, entity resolution, cyber telemetry, semantic graph, and cross-enclave fusion passes.

4 - Pilot Charters 6
Turns priority into a signable plan: quick-win screening, the nine-section charter generator, ROI and value density, FTE allocation, and exit gate criteria.

5 - 30-60-90 Roadmaps 6
Sequences the work and closes the loop: dependency-aware backlog sequencing, the 30-day quick-win phase, the 60-day ETL handoff, the 90-day scaling plan, and a portfolio audit.

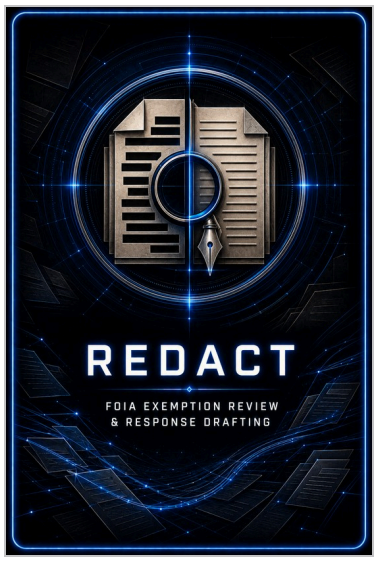
BUILT FOR THE MISSION ENVIRONMENT

DEPLOYS YOUR WAY
Air-gapped on-prem, cloud-connected, or hybrid, the same platform, deployed to match your security posture and mission need.

GROUNDED IN TRUSTED SOURCES OF TRUTH
QUARRY works from your exported metadata, schemas, lineage logs, and catalog entries from tools such as Collibra, Alation, and Unity Catalog, read alongside VAULTIS, ICD 501 and 208, the CUI Registry, and the CAPCO Register - flagging anything needing live verification.

ADVISORY-ONLY BY DESIGN
QUARRY reads exported inventories and samples; it never queries live systems, executes ETL, or modifies a schema. The data officer, Mission Exploitation Lead, and CDO/CTO keep every decision on capital, prioritization, and access; handling stays with the ISSO/CISO.

A BASE AGENT, TAILORED ON-SITE
QUARRY ships as a base agent and is refined on-site with your teams - incorporating your catalog conventions, mission value criteria, scoring weights, and tier thresholds.



G.H.O.S.T. PLATFORM AGENT | 24 / 34

REDACT

FOIA Processing & Declassification Review

WHAT REDACT IS

Redact is the FOIA processing, declassification, and records-review agent in G.H.O.S.T., built for Government Information Specialists doing line-by-line review, FOIA Officers and Public Liaisons, declassification reviewers applying E.O. 13526, agency counsel preparing Vaughn indexes, and records managers tracking deadlines.

Redact ingests requester correspondence, responsive-record sets, classification guides, and queue exports - then flags exemptions passage by passage under (b)(1) through (b)(9), maps redactions to the page and paragraph, verifies segregability, drafts determinations and Vaughn-style indexes, and triages the queue by statutory age, with a citation behind every proposal.

WHAT SUCCESS LOOKS LIKE

- Withholdings that survive de novo review. Every deletion carries its exemption citation, a record-specific foreseeable-harm articulation under 5 U.S.C. 552(a)(8)(A), and a segregability finding.
- Deadlines that name their convention. 20 workdays to a determination from perfection, 10 more on unusual circumstances, expedition in 10 calendar days, appeals in 20 workdays.
- Vaughn indexes and declarations screened for the fatal defects before counsel sees them: conclusory entries, missing segregability findings
- Over-redaction caught on the second pass, with letter, redaction map and index citing the same exemptions for the same passages
- A faster queue overall, with hard stops for b(1), exclusions and RD/FRD raised before a response goes out

34 Base Automated Workflows Across the Full FOIA and Declassification Lifecycle

1 - Exemption Review 8
Proposes every deletion defensibly: Exemption Screener across (b)(1) through (b)(9), Redaction Map, Segregability Check, Foreseeable Harm Test, Privacy Balancing.

2 - Correspondence 6
Drafts what goes out the door: Response Letter with appeal rights and OGIS notice, Acknowledgment Letter, Delay and Interim Notice, Vaughn Index Draft, Glomar.

3 - Queue & Tracking 7
Keeps the program honest about its clocks: Request Intake, Backlog Triage, Clock and Tolling Audit across workday and calendar-day conventions, Fee Determination.

4 - Search Adequacy 4
Proves the agency looked where it should have: Custodian Scoping, Search Term Design, Search Declaration under 28 U.S.C. 1746, No-Records Adequacy audit.

5 - Appeals & Defense 6
Finds the weak withholding first: Appeal Analysis de novo, Appeal Response Draft, Litigation Risk Review, Vaughn Defect Audit, Consistency Cross-Check.

6 - Declass & MDR 3
Flags and refers, never decides: E.O. 13526 Review with the 25-year and 25X sweep, MDR Processing kept separate from FOIA, Kyl-Lott Screening on RD or FRD.

BUILT FOR THE RECORDS ENVIRONMENT

DEPLOYS YOUR WAY

Air-gapped on-prem, cloud-connected, or hybrid, the same platform, deployed to match your security posture and mission need.

GROUNDED IN TRUSTED SOURCES OF TRUTH

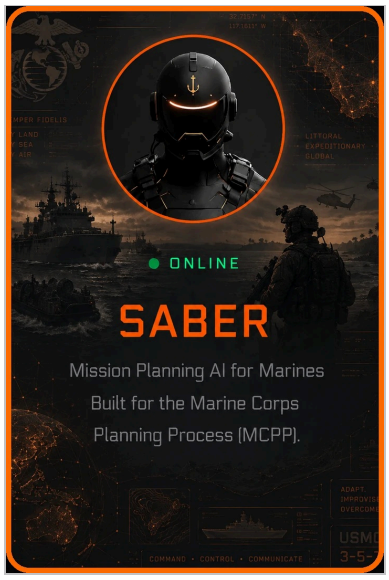
Redact draws on your uploaded records and requester correspondence first, then 5 U.S.C. 552, E.O. 13526, E.O. 12600, DOJ OIP guidance and your FOIA regulations - and flags anything needing verification against a live source.

ADVISORY-ONLY BY DESIGN

FOIA Officers, Government Information Specialists, declassification authorities, certified Kyl-Lott reviewers, agency counsel and Chief FOIA Officers hold final authority on every withholding, release and classification action. Redact screens, plans and drafts, splitting findings into OBSERVED, ASSESSED and RECOMMENDED.

A BASE AGENT, TAILORED ON-SITE

Redact ships as a base agent and is refined on-site with your teams - incorporating your letter templates, component structure, track and fee rules, and your workforce's nuance.



G.H.O.S.T. PLATFORM AGENT | 25 / 34

SABER

Marine Corps Mission Planning Intelligence

WHAT SABER IS

SABER is the Marine Corps operations-planning specialist in G.H.O.S.T., built around the six steps of the Marine Corps Planning Process (MCWP 5-10). It supports the staffs that carry a MAGTF from receipt of a higher order to a signed OPORD - G-3/S-3 and OPT leaders, G-2/S-2, the COS/XO, and the fires, aviation, logistics, and communications planners.

SABER ingests the higher order, IPB products, staff estimates, task organization, and terrain data - then builds task and limitation ledgers, COG analyses, COA graphics and narratives, synchronization matrices, war-game records, comparison matrices, and OPORD/WARNORD/FRAGO text, every product step-labeled and stamped DRAFT for staff review.

WHAT SUCCESS LOOKS LIKE

- Step-correct products a planner briefs with minimal editing - every task traceable to the higher order, fact separated from assumption
- COAs that survive all five screening criteria - a COA that failed in the war game is never quietly adjusted
- Decision support the watch floor can execute: triggers, feeding information requirements, and latest-time-of-value discipline
- A faster planning cycle overall: one-third/two-thirds time discipline enforced and subordinate planning time protected
- Orders that hold up: reconciliation, crosswalk against higher and adjacent, and a signature package where nothing open hides

72 Base Automated Workflows Across All Six MCPP Steps

1 - Problem Framing 10

Frames the fight right: task analysis with constraint/restraint ledgers, COG analysis, CCIR building, collection planning, the planning timeline, and the problem-framing brief.

2 - COA Development 17

Builds distinguishable options: IPB updates, battlespace framework, force arraying, options generation, purpose-and-task matrices, task organization, and COA screening.

3 - COA War Game 24

Fights the plan before the enemy does: red and green cell packs, turn capture, adjudication, loss ledgers, decision points, DST/DSM, branches and sequels.

4 - COA Comparison 5

Earns the decision honestly: per-criterion evaluation, the comparison matrix with visible arithmetic, the decision brief, and verbatim decision capture.

5 - Orders Development 7

Turns the decision into a signed order: CONOPS, the five-paragraph base order, annexes, reconciliation, crosswalk, signature package, and dissemination.

6 - Transition 9

Moves the plan from authors to executors: transition briefs, rehearsals, confirmation briefs, question tracking, responsibility transfer, and FRAGOs.

BUILT FOR THE MARINE CORPS ENVIRONMENT

DEPLOYS YOUR WAY

Air-gapped on-prem, cloud-connected, or hybrid, the same platform, deployed to match your security posture and mission need - afloat, forward, or in garrison.

GROUNDED IN TRUSTED SOURCES OF TRUTH

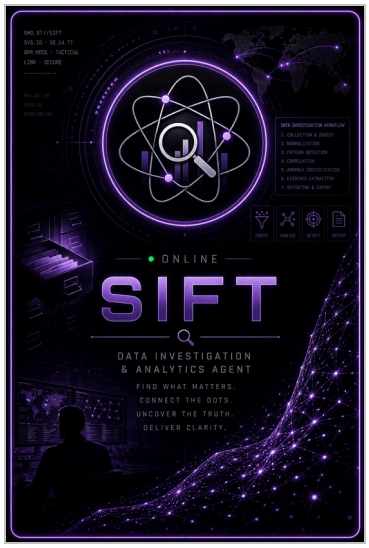
SABER draws on your uploaded orders, IPB, estimates, and task organization, grounded in MCWP 5-10 and the MCDP/MCRP doctrinal set - and when a figure is missing, it names the gap and uses a labeled placeholder.

ADVISORY-ONLY BY DESIGN

The commander keeps the COA decision, risk acceptance, and signature. SABER analyzes, drafts, and recommends - facts, assumptions, and judgments visibly separate, every product stamped DRAFT - FOR STAFF REVIEW.

A BASE AGENT, TAILORED ON-SITE

SABER ships as a base agent and is refined on-site with your teams - incorporating unit SOPs, orders formats, T/O and T/E data, and the nuance only your Marines know.



G.H.O.S.T. PLATFORM AGENT | 25 / 34

SIFT

Digital Forensics & Incident Response Analysis

WHAT SIFT IS

Sift is the digital forensics and incident response analysis specialist in G.H.O.S.T., built for Intelligence Community and defense cyber personnel working incidents and hunts from exported evidence. It supports SOC analysts at every tier, incident responders, threat hunters, detection engineers, and CIRT leads.

Sift ingests alert exports, log excerpts, endpoint and network telemetry, malware-analysis reports, email headers, and case notes - then triages evidence, reconstructs timelines, tests competing hypotheses, maps behavior to MITRE ATT&CK, drafts detections, and writes the report, with a named source artifact and a stated confidence behind every finding.

WHAT SUCCESS LOOKS LIKE

- Case packages an analyst can act on with minimal editing, every observation cited
- Observed fact separated from inference, confidence stated per judgment
- A faster incident overall: triage, timelines, and handoff packages assembled fast
- Detection drafts with tuning parameters, stamped DRAFT - NOT TESTED
- Products that survive peer review: benign alternatives stated, clock skew surfaced

50 Base Automated Workflows Across the Full Incident Lifecycle

1 - Case Foundations 6
Starts the case on solid ground: Evidence Intake Triage, Case Triage Brief, Incident Timeline, IOC Extraction Pack, Event Pivot Finder, Entity Scope Review.

2 - Analytic Reasoning 4
Keeps the analysis honest: Attack Path Sketch labeled observed vs inferred, Alt Hypothesis Test, Evidence Gap Plan, MITRE Technique Map.

3 - Identity & Access 10
Works the credential lane: Auth Failure Triage, Password Spray, Brute Force, Account Takeover and Kerberos Abuse Hunts, Privilege Escalation, New Admin, Service Account, MFA and VPN Review.

4 - Endpoint & Malware 10
Reads host behavior in context: Endpoint Alert Triage, Process Tree and Remote Tool Review, Script, Persistence, Lateral Move, LOLBin and Ransomware Hunts.

5 - Network & Email 10
Follows the traffic and the lure: DNS, Proxy, Firewall, Web Attack and TLS Review, Beaconing and Exfiltration Hunts, C2 Candidate Review, Email Header and Phish Cluster Review.

6 - Hunt & Detection 5
Turns findings into coverage: Hunt Plan Builder with stop conditions, SIEM Query Draft in your dialect, Sigma and conservative YARA Rule Drafts, Detection Test Plan.

7 - Respond & Report 5
Closes the case cleanly: Alert Playbook Draft, IR Runbook Draft, Analyst Handoff Pack, blameless Incident Postmortem, Executive Cyber Brief.

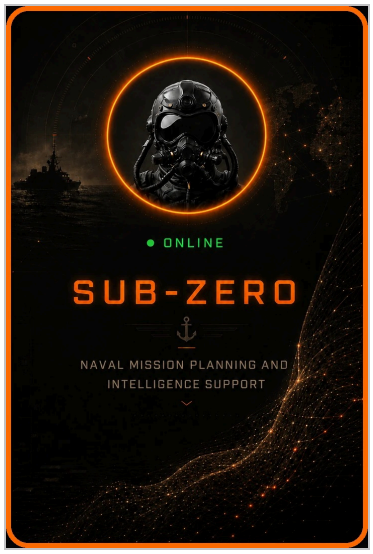
BUILT FOR THE CYBER ENVIRONMENT

DEPLOYS YOUR WAY
Air-gapped on-prem, cloud-connected, or hybrid, deployed to match your security posture and mission need.

GROUNDED IN TRUSTED SOURCES OF TRUTH
Sift draws on your supplied case package first - log and alert exports, telemetry extracts, malware reports, email headers, case notes - plus the NIST-aligned IR lifecycle and your loaded MITRE ATT&CK reference, and flags anything needing live verification.

ADVISORY-ONLY BY DESIGN, ARTIFACT-ONLY BY BOUNDARY
Sift reasons over exported evidence and never touches a live host, queries a live system, executes code, or detonates a sample. Analysts, responders, hunters, detection engineers, and the incident commander hold authority over every disposition, containment action, and detection; observed is separated from inferred, with confidence stated.

A BASE AGENT, TAILORED ON-SITE
Refined on-site with your report templates, SIEM dialects, and thresholds.



G.H.O.S.T. PLATFORM AGENT | 27 / 34

SUB-ZERO

Naval Mission Planning & Information Warfare

WHAT SUB-ZERO IS

Sub-Zero is the naval operations-planning specialist in G.H.O.S.T., built around the six steps of the Navy Planning Process (NWP 5-01) with an information warfare emphasis. It supports the staffs that carry a maritime operations center or strike group from a higher order to a released OPORD - N5 future plans and OPT leaders, N3 operations, N2, the chief of staff, and the N39/IWC, N6, N4, and METOC planners.

Sub-Zero ingests the higher order, IPOE products, staff estimates, force lists, and environmental data - then builds task and limitation ledgers, COG analyses, COA statements and sketches, synchronization matrix row data, war-game records, comparison matrices, and OPORD/WARNORD/FRAGO text with annexes and OPTASKs, each product NPP-step-labeled and stamped DRAFT.

WHAT SUCCESS LOOKS LIKE

- Step-correct products a planner briefs with minimal editing - every task traceable to the higher order, fact separated from assumption
- COAs that survive all five screening criteria - a COA that failed in the war game is never quietly adjusted
- Information warfare planned in, not bolted on: every COA carries an EMCON posture and its command-and-control cost, with spectrum and EEFI deconfliction on the record
- A faster planning cycle overall: one-third/two-thirds time discipline enforced and every product mapped to the board or brief that consumes it
- Orders that hold up: reconciliation and crosswalk run across every annex and every OPTASK - a base order without its OPTASKs is not executable

43 Base Automated Workflows Across All Six NPP Steps

1 - Mission Analysis 9
Frames the problem right: initial IPOE, task and limitation ledgers, facts and assumptions, COG analysis, operational design, CCIR and EEFI, and the mission statement.

2 - COA Development 9
Builds distinguishable options: relative combat power, COA options, force and sustainment arraying, CWC task organization, EMCON and spectrum, sync matrix row data, and screening.

3 - COA Wargaming 9
Fights the plan before the adversary does: wargame design and prep, red and neutral cells, turn capture, decision points, and the risk and culmination log.

4 - COA Comparison 5
Earns the decision honestly: per-criterion evaluation first, then the comparison matrix with its arithmetic shown, the decision brief, and verbatim decision capture.

5 - Orders Development 7
Turns the decision into a releasable order: CONOPS, base order, annexes and OPTASKs, execution timeline, branches and sequels, reconciliation and crosswalk.

6 - Transition 4
Moves the plan from planners to the watch floor: transition handoff, rehearsal planning and capture, and the assessment plan written before execution begins.

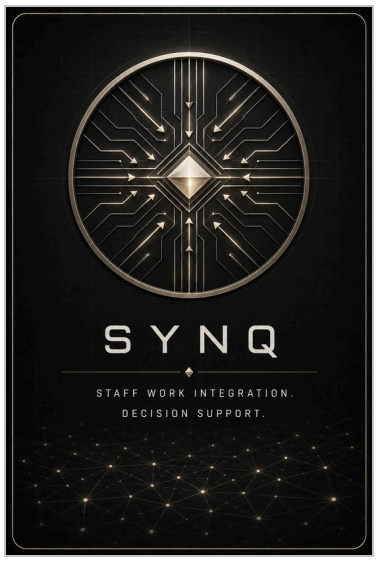
BUILT FOR THE NAVY INFORMATION WARFARE ENVIRONMENT

DEPLOYS YOUR WAY
Air-gapped on-prem, cloud-connected, or hybrid, the same platform, deployed to match your security posture and mission need - afloat, forward, or ashore.

GROUNDED IN TRUSTED SOURCES OF TRUTH
Sub-Zero draws on your uploaded orders, IPOE, staff estimates, and force lists, grounded in NWP 5-01 and the joint publication set - a missing figure is named as a gap, not guessed, and a synchronization matrix is read only as labeled text, never as a graphic.

ADVISORY-ONLY BY DESIGN
The commander keeps the COA decision, EMCON posture, risk acceptance, and release of the order. Sub-Zero analyzes, drafts, and recommends - facts, assumptions, and judgments visibly separate, every product stamped DRAFT.

A BASE AGENT, TAILORED ON-SITE
Sub-Zero ships as a base agent and is refined on-site with your teams - incorporating MOC procedures, order and OPTASK formats, and battle rhythm.



G.H.O.S.T. PLATFORM AGENT | 28 / 34

SYNQ

Staff Work Integration & Decision Support

WHAT SYNQ IS

SYNQ is the executive systems-integration, coordination, and decision-support agent in G.H.O.S.T., built for the people who hold a front office together: government-contractor Systems Integrators, project and program integrators, executive officers, PMO analysts, Chiefs of Staff, and business-operations leads who work through facilitation and context rather than command authority.

SYNQ ingests leadership intent, taskers and suspenses, meeting material, status reports, schedules, RAID registers, and stakeholder positions - then frames the decision, drafts the product, maps the dependency, and assigns the action, keeping fact, assumption, inference, and recommendation apart.

WHAT SUCCESS LOOKS LIKE

- Decision-ready packages, not raw material. Every memo, BLUF, and options analysis names the decision owner, the criteria, the major downside, and the exact ask, at a HIGH, MODERATE, or LOW confidence level.
- Decisions captured without invented consensus. Records carry a state of PROPOSED, IN REVIEW, APPROVED, REJECTED, DEFERRED, or SUPERSEDED, and only the authorized owner reaches APPROVED.
- Meetings built around an outcome rather than a topic, with the smallest purposeful attendee set and no parked item left unclosed.
- Dependencies, blockers, unclear ownership, and decision latency surfaced before a committed milestone slips - GREEN, AMBER, and RED rollups that never average independent reds into green.
- Faster staff work overall: taskers split into numbered requirements with internal suspenses ahead of the authoritative date, and actions that carry one owner, a deliverable, a due date, and completion evidence.

50 Base Automated Workflows Across the Executive Integration Cycle

1 - Intake & Priority 5
Starts the tasking clean: Request Triage, Tasker Breakdown, Priority Crosswalk, Deadline Risk Scan, and Calendar Deconflict.

2 - Meeting Ops 10
Runs the meeting lifecycle: Meeting Purpose, Agenda Architect, Read-Ahead Builder, Exec Prep Brief, Decision Capture, and Action Register.

3 - Program Sync 10
Holds the plan together: Project Charter, Workplan Builder, Dependency Mapper, RAID Register, Change Impact, and Schedule Recovery.

4 - Executive Support 10
Puts the leader in front of the decision: One-Page BLUF, Talking Points, Decision Memo, Options Analysis, and Escalation Brief.

5 - Stakeholder Align 5
Finds the disagreement early: Stakeholder Map, Comms Plan, Alignment Scan, Conflict Prep, and Cross-Team Sync.

6 - Governance 10
Keeps the operating rhythm honest: KPI Designer, Status Rollup, Trend Narrative, Budget Snapshot, Scope Guard, and Handoff Package.

BUILT FOR THE EXECUTIVE ENVIRONMENT

DEPLOYS YOUR WAY

Air-gapped on-prem, cloud-connected, or hybrid, the same platform, deployed to match your security posture and mission need.

GROUNDED IN TRUSTED SOURCES OF TRUTH

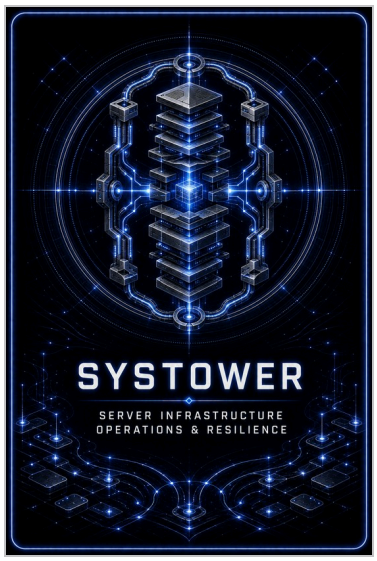
SYNQ works from your exported record - taskers, meeting material, project and task tool exports, ITSM and financial extracts, and status reports - treating each as a snapshot and flagging what needs verification when sources disagree.

ADVISORY-ONLY BY DESIGN

Sponsors, government leads, program managers, technical authorities, contracting officers, finance, and legal retain final authority over priorities, approvals, resources, contracts, personnel, and acceptance. SYNQ drafts and recommends; contractor products are marked CONTRACTOR-PRODUCED DRAFT - FOR GOVERNMENT REVIEW, and FAR 7.503 questions are routed, not decided.

A BASE AGENT, TAILORED ON-SITE

SYNQ ships as a base agent and is refined on-site with your teams - your templates, risk and health scales, suspense conventions, and the nuance only your workforce knows.



G.H.O.S.T. PLATFORM AGENT | 29 / 34

SysTower

Server Operations & Infrastructure Analysis

WHAT SYSTOWER IS

SysTower is the server operations and infrastructure specialist in G.H.O.S.T., built for the teams who keep enterprise compute running - Linux and Windows Server administrators, virtualization engineers, storage and backup administrators, infrastructure architects, and the managers accountable for the estate.

SysTower ingests monitoring exports, hypervisor data, patch and scanner reports, backup job histories, configuration dumps, and incident notes - then reconstructs what happened, rates the risk, models the capacity, and drafts the change requests, runbooks, and reports, with a cited source artifact behind every finding.

WHAT SUCCESS LOOKS LIKE

- Incidents that close with a defensible timeline - probable causes ranked against supporting and contradicting evidence, never a root cause declared before the evidence supports one
- Change requests that reach the board complete: blast radius enumerated, risk rated with its basis, rollback written before the change is scheduled
- Patch waves that reconcile honestly - patched-but-not-rebooted reported as pending, never counted as compliant
- Recovery you can prove: stated RPO and RTO checked against real job history and the last verified restore - an untested backup is a hypothesis, not a capability
- A faster operations cycle overall: capacity projected to a date with the arithmetic printed, so expansion decisions land before the procurement lead time closes

90 Base Automated Workflows Across the Server Operations Lifecycle

1 - Incident Response 9

Restores service without destroying the evidence: triage, severity classification, timeline reconstruction, probable cause analysis, and blameless post-incident review.

2 - Patch & Change Management 17

Closes exposure on a defensible clock and moves changes safely: advisory triage, ring-based rollout, reboot reconciliation, blast radius analysis, and rollback planning.

3 - Capacity & Virtualization 16

Reads the platform layer the guests cannot see: time-to-threshold models, host sizing, cluster headroom, contention and overcommit analysis, and snapshot hygiene.

4 - Storage, Backup & Recovery 15

Proves recoverability instead of assuming it: storage latency, exhaustion risk, RPO and RTO verification, restore tests, DR runbooks, and ransomware readiness.

5 - Hardening & Lifecycle 12

Holds the build standard from provisioning to disposal: benchmark and STIG gap review, drift detection, access review, end-of-support sweeps, and decommission packages.

6 - Runbooks & OS Triage 21

Puts operating knowledge in writing with the depth behind it: runbook drafting, command safety review, on-call handover, plus dedicated Linux and Windows triage.

BUILT FOR THE ENTERPRISE ENVIRONMENT

DEPLOYS YOUR WAY

Air-gapped on-prem, cloud-connected, or hybrid, the same platform, deployed to match your security posture and mission need.

GROUNDED IN TRUSTED SOURCES OF TRUTH

SysTower's analysis draws on your uploaded exports - monitoring and hypervisor data, patch reports, backup histories, configuration and event logs - alongside recognized hardening baselines, and flags anything needing verification against current vendor documentation.

ADVISORY-ONLY BY DESIGN

SysTower drafts; it never executes, approves, or schedules. The change advisory board and infrastructure manager keep approval authority, and every production-affecting step carries its precondition check and its rollback path in the same block.

A BASE AGENT, TAILORED ON-SITE

SysTower ships as a base agent and is refined on-site with your teams - incorporating your estate topology, baselines, cadences, and the thresholds only your operators know.



G.H.O.S.T. PLATFORM AGENT | 30 / 34

TALON

Cyber Threat Intelligence

WHAT TALON IS

TALON is the cyber threat intelligence specialist in G.H.O.S.T., built for the people who turn reporting into defensive action - CTI analysts, threat hunters, detection engineers, SOC leads, IR leads consuming intel support, and CISOs consuming strategic products.

TALON ingests vendor and government reporting, incident reports, IOC sets, and sandbox output - then extracts TTPs and maps them to ATT&CK technique IDs, correlates actors and campaigns, grades every source on the Admiralty scale, and drafts cited products carrying a stated confidence and a named defensive implication.

WHAT SUCCESS LOOKS LIKE

- Attribution that holds up. Every claim states its tier and never exceeds the evidence presented - no flat attribution at any confidence level
- Sources graded, not just cited. Admiralty reliability and credibility on every source, with circular reporting traced so two outlets carrying one upstream report count as one
- Behavior mapped to ATT&CK technique IDs with an evidence anchor per mapping, at the most specific level the reporting supports and no deeper
- Indicators that do not burn SOC trust: decay tiers assigned by age and type, collision risk flagged before anything reaches a block list
- A faster intelligence cycle overall: every product written to one intelligence level, with its defensive implication and owning role named

42 Base Automated Workflows Across the Intelligence Cycle

1 - Intake & Grading 5
Establishes what the reporting is worth: intake and normalization, Admiralty source grading, circular-reporting traces, victimology, and activity-versus-reporting timelines.

2 - ATT&CK Mapping 5
Turns narrative into behavior: TTP extraction, ATT&CK mapping with evidence anchors, kill chain narratives, and sub-technique refinement.

3 - Actors & Campaigns 6
Builds the adversary picture: actor profiles, campaign correlation with competing groupings, alias cross-referencing, Diamond Model builds, and ACH matrices.

4 - IOC Management 5
Keeps indicators honest: structured IOC packages, decay-tier reviews, collision checks before a block list, pivot maps, and retro-hunt sets.

5 - Hunt Packages 4
Turns intelligence into hunts: package building with query logic and false-positive sources, hypothesis framing, result adjudication, and crown-jewel scoping.

6 - Detection Support 5
Closes the loop with the engineers: detection recommendation and gap briefs, evasion considerations, Pyramid of Pain ranking, and false-positive forecasting.

7 - PIR & RFI 4
Points collection at decisions: PIR drafting, coverage maps, RFI responses, gap reports.

8 - Products & QA 8
Ships defensible product: flash reports, landscape briefs, leadership briefings, and sector profiles, plus estimative-language, sourcing, and attribution reviews.

BUILT FOR THE THREAT ENVIRONMENT

DEPLOYS YOUR WAY
Air-gapped on-prem, cloud-connected, or hybrid, the same platform, deployed to match your security posture and mission need.

GROUNDED IN TRUSTED SOURCES OF TRUTH
TALON reasons over the reporting you provide - vendor and government advisories, ISAC sharing, incident data, and platform and sandbox exports - graded on the Admiralty scale, with anything undated, unmarked, or uncorroborated flagged rather than normalized away.

ADVISORY-ONLY BY DESIGN
TALON never runs a query, scan, or platform connection; it analyzes uploaded and exported content only. Attribution released outside the team stays with the senior CTI analyst, detection logic with the detection engineer, and product release with the CISO.

A BASE AGENT, TAILORED ON-SITE
TALON ships as a base agent and is refined on-site with your teams - incorporating your sector, your intelligence requirements, and the nuance only your workforce knows.



G.H.O.S.T. PLATFORM AGENT | 31 / 34

TRIAD

Multi-Cloud Infrastructure & Engineering

WHAT TRIAD IS

TRIAD is the multi-cloud infrastructure and engineering specialist in G.H.O.S.T., built for the teams that design, secure, and pay for cloud estates - Cloud Architects, DevOps Engineers, SREs, Security Engineers, FinOps Managers, and IT leadership - across AWS, Azure, and Google Cloud.

TRIAD ingests architecture diagrams, Terraform / Bicep / CloudFormation manifests, IAM and RBAC exports, telemetry, and billing data - then scores designs against each hyperscaler's Well-Architected framework, drafts hardened infrastructure-as-code, scopes least-privilege access, and builds cost roadmaps with a citation behind every finding.

WHAT SUCCESS LOOKS LIKE

- Architecture reviews scored pillar by pillar against the AWS, Azure, and GCP Well-Architected frameworks and CIS Benchmarks - findings that survive Architect, SRE, and CISO review
- Validated, modular IaC drafts (Terraform, Bicep, CloudFormation) with security and cost rationale written into every module
- IAM, RBAC, and Org Policy recommendations that eliminate wildcard permissions and privilege-escalation paths across all three clouds
- Backup and DR designs built to explicit targets - RTO under 1 hour, RPO under 15 minutes - with cross-region and cross-cloud options costed
- A faster path from billing telemetry to a defensible savings plan: rightsizing candidates, orphan registers, and commitment strategy with break-even math shown

72 Base Automated Workflows Across AWS, Azure, and Google Cloud

1 - Compute & Network 12
Lays the foundations per cloud: EC2 & VPC Design, Transit Gateway Pass, Azure VM & VNet Plan, ExpressRoute Setup, Cloud Interconnect, and network audits on each platform.

2 - Data & Storage 12
Designs and protects the data layer: S3 & EBS Security, RDS & Aurora Design, Cosmos DB Architecture, Azure Backup Vaults, BigQuery Data Ops, and Cloud Spanner Plan.

3 - IAM & Security 12
Guards identity and encryption everywhere: IAM Least Privilege, AWS Org SCP Design, Entra ID Governance, GCP IAM Audit, and KMS key governance passes.

4 - Containers 12
Architects the container estate: EKS, AKS, and GKE Cluster Design, ECR and ACR registry security, Cloud Run Operations, and Anthos fleet operations.

5 - Serverless 12
Designs event-driven systems per platform: Lambda Architecture, Step Functions Plan, Azure Functions Plan, Service Bus Messaging, Pub Sub Architecture, and Eventarc Routing.

6 - Ops & FinOps 12
Keeps the estate observable and affordable: CloudWatch Dashboard, AWS Config Audit, Log Analytics Audit, and cost-rightsizing passes on all three clouds.

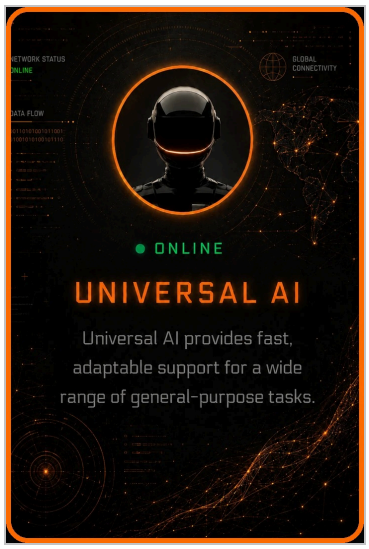
BUILT FOR THE MISSION ENVIRONMENT

DEPLOYS YOUR WAY
Air-gapped on-prem, cloud-connected, or hybrid, the same platform, deployed to match your security posture and mission need.

GROUNDED IN TRUSTED SOURCES OF TRUTH
TRIAD's analysis draws on your uploaded architecture, IaC, IAM, and billing exports, read against the AWS, Azure, and GCP Well-Architected frameworks and CIS Benchmarks - and flags every quota or rate figure for verification against your own exports.

ADVISORY-ONLY BY DESIGN
TRIAD never runs a CLI command or touches a live resource. Cloud Architects, SRE leads, the CISO, and the FinOps Manager keep every deployment, access, and spend decision; findings are labeled Source-supported or Inference, cited, and stamped DRAFT.

A BASE AGENT, TAILORED ON-SITE
TRIAD ships as a base agent and is refined on-site with your teams - incorporating your landing zones, tagging taxonomy, and the nuance only your engineers know.



G.H.O.S.T. PLATFORM AGENT | 32 / 34

UNIVERSAL AI

General-Purpose Assistant & Concierge

WHAT UNIVERSAL AI IS

Universal AI is the general-purpose assistant and front door in G.H.O.S.T. - the agent every user lands on when no specialist owns the request. It serves the whole force: staff, analysts, engineers, and reviewers who need dependable general help between specialist tasks.

Universal AI ingests open-ended requests plus any files you provide - documents, threads, transcripts, and images - then answers, summarizes, extracts, compares, drafts, and explains, separating established fact from inference with a stated confidence. When a request belongs to a specialist, it gives a brief orienting answer and names the agent that owns it.

WHAT SUCCESS LOOKS LIKE

- Everyday work, done right. Answers, summaries, extractions, comparisons, and drafts that separate established fact from inference and state their assumptions - concise products a reviewer can trust.
- An AI-fluent workforce. Plain-language education on responsible AI, generative AI, RAG, reasoning models, and prompt engineering, so people get better results from every G.H.O.S.T. agent.
- Nothing invented to fill a gap: names, numbers, citations, and quotes are marked unknown rather than fabricated, with confidence stated where it matters.
- Quick general help between specialist tasks - drafting, planning, light analysis, and image and document interpretation.
- Every request in the right lane: general work handled directly, specialist work routed by name with context carried over, and high-stakes calls deferred to the responsible human.

37 Base Automated Workflows Across General Work, Education, and Operations

Support

Tools 7
Puts core assistant power one step away: Summarize Content and Summarize in One for fast BLUFs, SOP Creator, Problem Solver, Prompt Creator, Expert-Mode, and a Prompt-Injection Scan that inspects untrusted text before you act on it.

AI Education 15
Builds AI fluency across the workforce: plain-language explainers covering What is Responsible AI, Generative AI, RAG, CAG, a Reasoning Model, and Prompt Engineering, plus How to Get the Best Results from the platform.

Document Management 6
Turns raw material into clean records: Document Summary Generator, Meeting Minutes Formatter, Standard Operating Procedure and Report Template Creators, Federal Register Notice Drafter, and Correspondence Drafter.

Administrative Support 5
Keeps operations moving: Action Item Tracker, Briefing Preparation, Conference/Event Planning Checklist, Onboarding Checklist Generator, and Travel Authorization Preparation.

Communications 4
Keeps stakeholders informed: Stakeholder Update Email and Internal Newsletter Content drafting, for consistent, on-message internal communications.

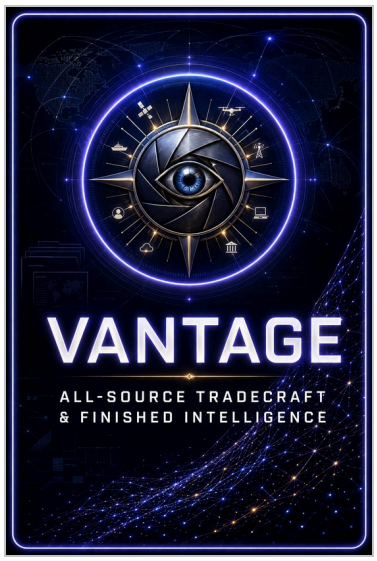
BUILT FOR ANY ENVIRONMENT

DEPLOYS YOUR WAY
Air-gapped on-prem, cloud-connected, or hybrid, the same platform, deployed to match your security posture and mission need - including tactical-edge and disconnected operation on local models.

ADVISORY-ONLY BY DESIGN
The human user keeps final authority; on regulated or high-stakes subjects the owning specialist and responsible role own the product. Universal AI separates fact from inference and never poses as an authority in a domain a specialist owns.

GROUNDED IN WHAT YOU PROVIDE
Universal AI draws on the documents, images, and threads you give it - and marks anything it cannot establish as unknown, flagging what needs verification rather than guessing.

A BASE AGENT, TAILORED ON-SITE
Universal AI ships as a base agent and is refined on-site with your teams - incorporating your templates, terminology, routing map, and the nuance only your workforce knows.



G.H.O.S.T. PLATFORM AGENT | 33 / 34

VANTAGE

Analytic Tradecraft & Finished Intelligence

WHAT VANTAGE IS

VANTAGE is the all-source analytic tradecraft and finished-intelligence specialist in G.H.O.S.T., built for the people who own the analytic line - all-source and fusion-cell analysts, senior analysts and branch chiefs, tradecraft reviewers, and analytic editors.

VANTAGE ingests raw reporting, analyst notes, structured-technique worksheets, prior products, and coordination comments - then builds the matrices, drafts the assessments, grades them against the nine analytic tradecraft standards, and preserves alternatives and dissent, with a citation and a stated confidence basis behind every line.

WHAT SUCCESS LOOKS LIKE

- Draft products that survive review-board scrutiny with minimal editing - key judgments up front, every assertion traceable to supplied reporting
- Uncertainty in the standard lexicon only: one likelihood term per judgment, with confidence stated separately and never merged into it
- Alternatives analyzed rather than acknowledged, and coordinating-office dissent recorded with its attribution and reasoning intact
- A reporting base described honestly - where it is single-threaded, dated, or traces back to one origin behind apparent corroboration
- A faster production cycle overall: sourcing traces, estimative-language checks, and prior-line reconciliation run before review, not discovered during it

30 Base Automated Workflows, Each Built in Frontier and Local Variants

1 - Structured Techniques ⊟

Externalizes the reasoning so it can be tested: ACH matrices with a denial-and-deception hypothesis always in the set, key assumptions and quality-of-information checks, devil's advocacy, what-if pathways, indicator sets, deception testing, and high-impact/low-probability analysis.

2 - Product Drafting ⊟

Turns analysis into products: finished assessments, calibrated key-judgment blocks, executive briefs, source summary statements, warning notices, trend assessments, graphic plans, and an estimative-language calibration pass before review.

3 - Review & QC ⌝

Protects the product before it ships: tradecraft scorecards against all nine standards, comment adjudication, claim-by-claim sourcing verification, alternatives audits, prior-line reconciliation, dissent recording, and peer-review packaging.

4 - Evidence Prep ⌝

Fixes the ground the analysis stands on: evidence baselines with a stated scope boundary, cited timelines, source reliability grids, circular-reporting traces, gap matrices ranked by consequence, entity registers, and raw-pull intake triage.

BUILT FOR THE ANALYTIC ENVIRONMENT

DEPLOYS YOUR WAY

Air-gapped on-prem, cloud-connected, or hybrid, the same platform, deployed to match your security posture - and all 30 workflows ship in both a Frontier and a Local variant, so the same menu runs on either.

GROUNDED IN TRUSTED SOURCES OF TRUTH

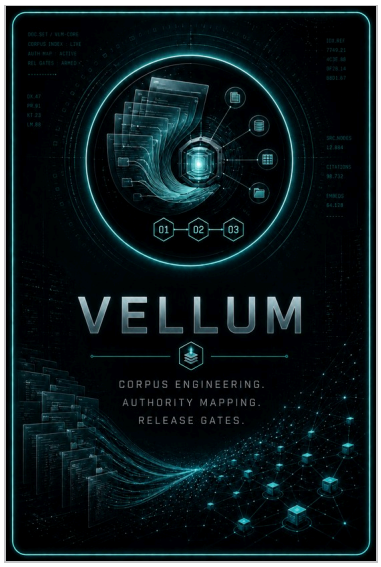
VANTAGE reasons only over the reporting and prior products you supply, applying ICD 203 analytic standards, ICD 206 sourcing, and the standard likelihood lexicon - and prints "Information not currently available" rather than filling a gap with plausible language.

ADVISORY-ONLY BY DESIGN

The drafting analyst owns the line and its confidence; the review chain and releasing authority own coordination and release. VANTAGE labels every line reporting, assumption, or judgment, proposes confidence with its basis, and never raises a judgment to match a preferred conclusion.

A BASE AGENT, TAILORED ON-SITE

VANTAGE ships as a base agent and is refined on-site with your teams - incorporating your production templates, coordination cadences, customer accounts, and the nuance only your workforce knows.



G.H.O.S.T. PLATFORM AGENT | 34 / 34

VELLUM

Corpus Engineering & Knowledge Governance

WHAT VELLUM IS

VELLUM is the secure corpus engineering specialist in G.H.O.S.T., built for the people who decide what a retrieval system may answer from - GHOST administrators, knowledge owners, document SMEs, security officers, and records managers - and for the RAG and knowledge-base engineers who stand the corpus up.

VELLUM ingests document inventories, source documents, supersession records, metadata and partition definitions, and exported retrieval-test results - then builds the corpus Bill of Materials, drafts the authority hierarchy, logs conflict and staleness registers, writes the golden-question baseline, and assembles the draft release package, citing the evidence behind every rating.

WHAT SUCCESS LOOKS LIKE

- Every document owned, tiered, and dispositioned, with cited evidence
- One T1 authority per topic per partition; unverified material never released
- Every conflict cited to both passages, a severity, and a named adjudicator
- A golden set across all six types - 100% on conflict probes and negatives
- A faster corpus release overall, run as repeatable workflows end to end

30 Base Automated Workflows Across the Full Corpus Lifecycle

1 - Intake & Scoping 4

Starts from a known baseline: Corpus BOM Builder, Scope & Partition Brief, BOM Recon Check, and Source Intake Screen.

2 - Authority & Tiers 4

Establishes what governs: Authority Tiering Draft across T1-T4/U, T1 Uniqueness Sweep, Tier Sign-Off Package, and Role Authority Card.

3 - Staleness Review 4

Rates currency on evidence, not age: Stale Source Scan, Supersession Chain Walk, Currency Evidence Check, and Refresh Cadence Planner.

4 - Conflict & Dedup 3

Surfaces what contradicts: Conflict Log Builder with severity and SME, Duplication Sweep, and Adjudication Tracker.

5 - Coverage & Risk 4

Finds what will fail in retrieval: Gap Analysis Report, Ingestion Risk Scan, Cross-Reference Check, and Fix-First Queue Builder.

6 - Metadata & Access 3

Makes the corpus governable: Metadata Schema Design, Partition Plan Builder on need-to-know rules, and Partition Consistency audits.

7 - Golden-Q Testing 4

Proves the corpus answers: Golden-Q Set Builder, Golden Coverage Matrix, Retrieval Miss Review, and Vocabulary Gap Finder.

8 - Release & Retire 4

Closes the gate honestly: Release Gate Review with evidence, Release Package Build, Retirement Package, and Rollback Readiness.

BUILT FOR THE MISSION ENVIRONMENT

DEPLOYS YOUR WAY

Air-gapped on-prem, cloud-connected, or hybrid, the same platform, deployed to match your security posture and mission need. VELLUM ships today with a Local-tier catalogue of 30 QA-clean workflows.

ADVISORY-ONLY BY DESIGN

Knowledge owners decide what is authoritative, administrators decide releases and partitions, security officers decide access, records managers decide retention. VELLUM inventories, rates, tests, and packages - every product stamped DRAFT - FOR ADMINISTRATOR REVIEW, every rating cited.

GROUNDED IN TRUSTED SOURCES OF TRUTH

VELLUM's analysis draws on your own inventories, source documents, supersession records, metadata definitions, and exported retrieval results - and flags anything of unknown provenance or unconfirmed currency for verification.

A BASE AGENT, TAILORED ON-SITE

VELLUM ships as a base agent and is refined on-site with your teams - incorporating your partition boundaries, metadata vocabularies, review cadences, and the nuance only your workforce knows.



In the field and in business

Real results from operational exercises and enterprise teams, and how G.H.O.S.T. supports acquisition and mission planning.

Why G.H.O.S.T.

Specialized agents built for business

Real-world business impact

Contracts and acquisition intelligence

Mission planning and operations

ENTERPRISE AI SHOULD WORK SMARTER

Most platforms promise transformation. G.H.O.S.T. delivers it.

THE CHALLENGES FACING MODERN ENTERPRISES

COMPLEX DEPLOYMENTS

Standing up AI requires costly specialists and months of configuration.

RIGID WORKFLOWS

Generic chatbots add complexity instead of streamlining processes.

RESOURCE DRAIN

Overreliance on outside consultants eats budgets and slows decisions.

SCALABILITY GAPS

Inconsistent outputs undermine confidence and enterprise-wide adoption.

ENTER G.H.O.S.T. - INTELLIGENCE YOU CAN DEPLOY TODAY

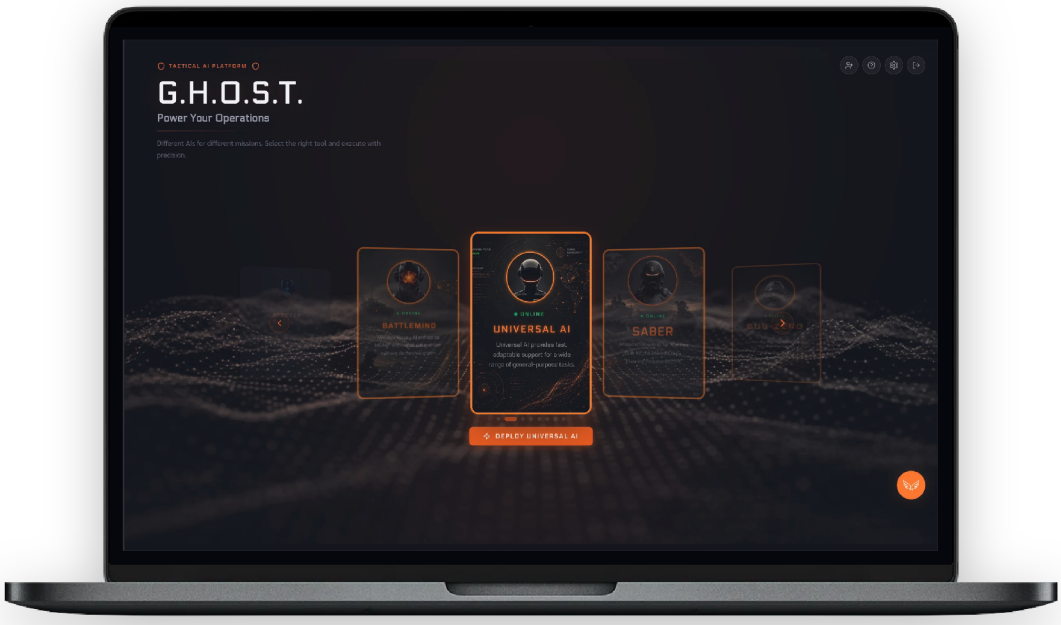
G.H.O.S.T. (Guided-Heuristic On-Prem Support & Troubleshooting) is a commercially available, modular agentic AI platform purpose-built for enterprise environments. It delivers specialized AI agents - pre-configured for your industry - that integrate directly into your existing workflows without demanding a new skillset from your team.

Think of it as hiring a team of domain experts that never sleep, never miss a compliance update, and deliver structured, auditable outputs at the speed of business.

Unlike generic LLM interfaces, G.H.O.S.T. encodes institutional knowledge into every agent. Whether it's regulatory analysis in life sciences, contract review in legal, or network diagnostics in IT, each agent is trained and tuned for your domain.

The platform runs on-premise, in private cloud, or hybrid environments - ensuring your sensitive data never leaves your control.

> G.H.O.S.T. removes the steep learning curves and resource strain by delivering an out-of-the-box AI solution that can be configured to any enterprise challenge.
>
> CYPHER LLC, PRODUCT DOCUMENTATION



SPECIALIZED AI AGENTS BUILT FOR BUSINESS

Each agent is a domain expert - pre-trained, pre-configured, ready on day one.

Data Center Operations Agent

Rack-side expertise, fully disconnected

- Runs on the G.H.O.S.T. Mobile laptop - fully disconnected
- Diagnoses switch, router & firewall faults from CLI output
- Isolates root cause from device logs & error counters
- Instant answers from ingested vendor manuals & runbooks
- Generates MOPs, incident reports & as-built documentation

Reduce unplanned downtime by up to 40%

IT Infrastructure & Network Agent

Intelligent network operations

- Automated network diagnostics & misconfiguration detection
- CLI command generation & runbook automation
- Log analysis, anomaly detection & root cause isolation
- Infrastructure documentation & topology mapping
- Incident response playbook generation

Reduce MTTR by up to 60% for network incidents

Contractor Acquisitions Agent

From requirement to award, faster

- Market research & industry sources analysis
- IGCE, D&F & J&A first-draft generation
- PWS / SOW drafting & evaluation criteria development
- Source selection documentation & debriefing preparation
- CDRL tracking, invoice review & CPARS narratives

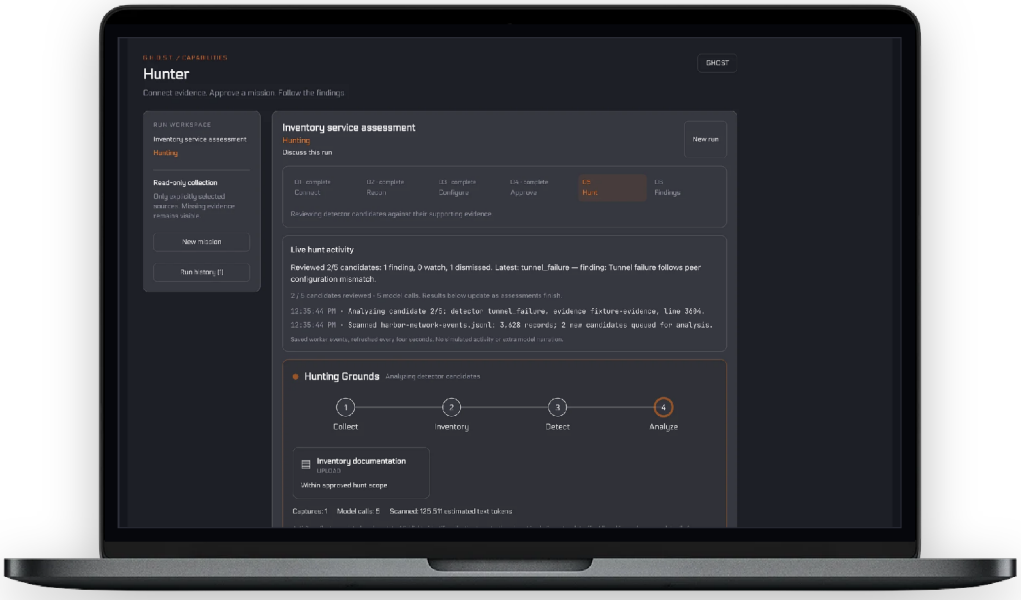
Cut acquisition package prep time by up to 65%

Legal & Regulatory Analysis Agent

Regulatory analysis you can defend

- Legal evaluation and regulatory analysis (JURIS)
- Every finding cited to the source it came from
- Gaps and assumptions stated, not papered over
- Structured drafts ready for attorney review
- Runs inside your environment, so privileged material stays put

Every claim traceable to its source



REAL-WORLD BUSINESS IMPACT

G.H.O.S.T. isn't hypothetical - it's delivering results now.

INDUSTRY USE CASES

Network Operations

CHALLENGE

A distributed network operations center faced recurring outages, inconsistent troubleshooting across shifts, and long escalation cycles caused by fragmented logs, runbooks, and topology records.

G.H.O.S.T. SOLUTION

G.H.O.S.T.'s Network Operations Agent correlated logs and topology data, isolated likely root causes, generated validated CLI remediation steps, and updated incident records with source-linked findings.

-60%

Network incident MTTR reduced

Contracts & Acquisitions

CHALLENGE

A contracting team had to produce market research, requirements and award documentation with fewer people, while keeping every file defensible on review.

G.H.O.S.T. SOLUTION

G.H.O.S.T.'s Acquisition Agent drafted structured first-pass packages from program inputs, flagged gaps before review, and kept every citation tied to its source.

-65%

Acquisition package prep time reduced

Enterprise IT & Operations

CHALLENGE

An enterprise technology team struggled with slow incident response times, inconsistent runbooks across regional teams, and growing technical debt from undocumented infrastructure changes.

G.H.O.S.T. SOLUTION

G.H.O.S.T.'s IT Agent provided real-time diagnostics, auto-generated CLI remediation commands from log analysis, and maintained living infrastructure documentation - all within a secure on-premise deployment.

-58%

Mean time-to-resolution improved

QUANTIFIABLE RETURN ON INVESTMENT

EFFICIENCY	Greatly reduces hours of manual work and replaces with one click automated workflows, structured AI outputs
ACCELERATES	Decision cycles from days to minutes across networks, contracts, and IT operations
ELIMINATES	Costly prompt-engineering overhead - agents are pre-configured for your domain
PROTECTS	Revenue and reputation through proactive compliance and risk identification

G.H.O.S.T. FOR CONTRACTS & ACQUISITION INTELLIGENCE

Accelerate the full acquisition lifecycle - from requirements development to award closeout.

Federal acquisition teams are asked to award faster and document more with a shrinking contracting workforce. Market research, requirements packages, source selections, and post-award oversight all compete for the same analyst hours - stretching PALT and pushing files past audit readiness. G.H.O.S.T.'s Acquisition Agent supports the contracting workforce across every phase of the acquisition lifecycle - delivering structured, defensible outputs that keep every file audit-ready and protest-defensible.

CORE ACQUISITION CAPABILITIES

MARKET RESEARCH & INDUSTRY ANALYSIS

Draft RFIs and Sources Sought notices, consolidate vendor responses, and analyze SAM.gov, FPDS-NG, and USASpending data to identify capable industry sources. G.H.O.S.T. generates structured market research reports and small business set-aside analyses that stand up to review.

Market research reports in minutes, not hours or weeks

ACQUISITION PLANNING & PRE-AWARD DOCUMENTATION

Assemble acquisition plan packages and flag gaps before review. G.H.O.S.T. produces IGCE first drafts, D&F and J&A justifications, and consolidation and bundling determinations - each citation-backed to the FAR and agency supplements for a defensible pre-award file.

Acquisition package prep time cut by 65%

REQUIREMENTS & SOLICITATION DEVELOPMENT

Generate PWS and SOW first drafts from program office inputs, develop evaluation criteria mirrored to Sections L and M, and run solicitation readiness checks - including sensitive-data scrubs before industry release and consolidated Q&A and amendment support.

First-draft PWS/SOW in under 2 minutes

PROPOSAL EVALUATION & SOURCE SELECTION SUPPORT

Support technical evaluation boards with structured consensus report drafting and award decision documentation that is consistency-checked against the evaluation record. Prepare unsuccessful-offeror notices and debriefings on protest-defensible timelines.

Evaluation documentation time cut by 70%

CONTRACT ADMINISTRATION & PERFORMANCE OVERSIGHT

Track contractor deliverables against CDRL matrices, analyze monthly status reports, monitor burn rate and funding runway, and review invoice packages before payment approval. Modification reviews enforce SF-30 and scope discipline; CPARS narratives are drafted from the performance record.

Modification review time reduced by 60%

CONTRACT CLOSEOUT & AUDIT READINESS

Run transition-in and transition-out readiness checks, review closeout packages, and complete final documentation of the contract record. G.H.O.S.T. generates closeout checklists and audit-ready files - ensuring clean contract completion with no outstanding obligations.

Closeout documentation time cut by 55%

REGULATORY FRAMEWORKS & CONTRACT VEHICLES SUPPORTED

FAR / DFARS | GSA SCHEDULES | GWACS / IDIQS | SAM.GOV / FPDS | SBIR / STTR

Works inside your own network or fully disconnected.

AGENTS ACROSS THE ENTERPRISE

Security, contracts, infrastructure and knowledge, ready on day one.

Beyond the four featured agents, G.H.O.S.T. ships a roster built for the work businesses and agencies actually do. Each one is pre-configured, cites its sources and leaves the decision to a person.

Intelligence and security

TALON	Cyber threat intelligence that turns reporting into defensive action
SIFT	Digital forensics and incident response analysis
BASTION	Security assessment and authorization through the risk management framework
PULSE	Finding the blind spots in security telemetry

Acquisition and legal

PRYME	Pre-award acquisition, from market research to solicitation
ARIA	Post-award contract administration
CADENCE	Requirements, deliverables and program rhythm
JURIS	Legal evaluation and regulatory analysis

Networks and infrastructure

NODE	Network operations and troubleshooting across vendors
LYNX	Linux and cloud router engineering
SYSTOWER	Server operations and infrastructure analysis
TRIAD	Multi-cloud engineering and cost
DATAMIND	Datacenter facilities and critical infrastructure
CASCADE	What a proposed change will actually touch
FATHOM	Rebuilding legacy systems from the evidence left behind
LATTICE	Database, data pipeline and data-quality analysis

Knowledge and assurance

UNIVERSAL AI	The general assistant every user starts with
AMBASSADOR	Translation and linguistic intelligence
LORE	Capturing what experts know before they leave
VELLUM	Governing what a retrieval system may answer from
GAUNTLET	Offline testing to prove an AI system works before it ships



Deploy and get started

Operational in under 48 hours. Secure by design.

DEPLOYMENT MODELS

Your security posture and your mission decide where G.H.O.S.T. runs.

ON-PREMISES

On your own servers and hardware, behind your own firewall.

PRIVATE CLOUD

In your own cloud tenant, so the data boundary does not move.

HYBRID

Sensitive work stays local while approved work uses the cloud.

FULLY DISCONNECTED

No internet needed. An optional zero-connectivity, no-call-home design for closed networks.

WHY IT DEPLOYS QUICKLY

WORKS OFFLINE OR ON CLASSIFIED NETWORKS

No internet is required, and nothing leaves your network.

NO DATA EXFILTRATION RISK

It runs inside DoD containerized environments, and inference stays local.

READY ON DAY ONE

Agents arrive pre-trained and pre-configured, so there is no experimental prompting.

BUILT ON DOCTRINE

Agents are trained on field manuals, joint publications and military standards, so answers stay accurate and trusted.

PORTABLE

It runs on hardened laptops or tactical servers, for field and austere sites.

RACK-SIDE HELP, FULLY DISCONNECTED

The Data Center Operations agent runs on the G.H.O.S.T. Mobile laptop and diagnoses switch, router and firewall faults from command-line output.

GET STARTED

Get started today and achieve mission dominance.

G.H.O.S.T. is available today for classified and unclassified use. It is actively deployed, validated, and ready for broader service adoption.

TO REQUEST A DEMO OR EVALUATION LICENSE IN A SECURE ENVIRONMENT

SUPPORT	support@ghostplatform.ai
WEBSITE	www.ghostplatform.ai

UEI and CAGE Code available upon request.

LEADERSHIP

CEO	Joseph Anderson | janderson@cypher-llc.com
CTO	John Izquierdo | jizquierdo@cypher-llc.com

